

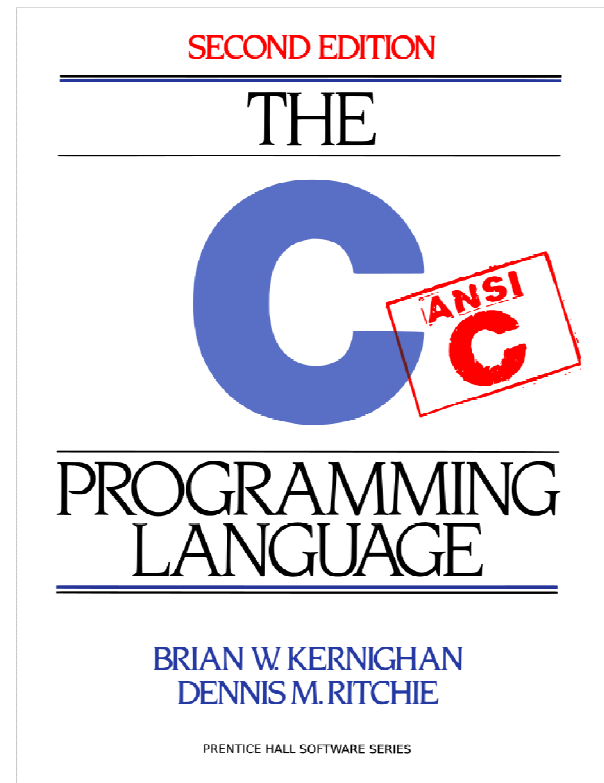
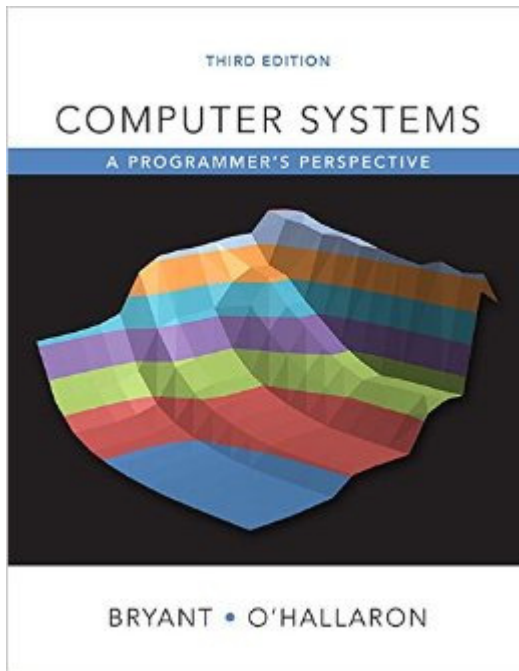
# Computer Systems Principles

C Structures



# Announcements

- Mid-term (Th, Feb 25) exam on paper, open book, close notes



# Announcements

- Quiz 4 and HW3 are out

# Learning Objectives

- To learn and apply C structures
- To understand a little about alignment
- To understand and apply C typedef

# C Structures

- **Essential**
  - For building up interesting data structures
- **Definition**
  - A **C structure** is a collection of one or more variables, typically of different types, grouped together under a single name for convenient handling
  - Kind of like a Java class with public instance variables and no methods

# C struct

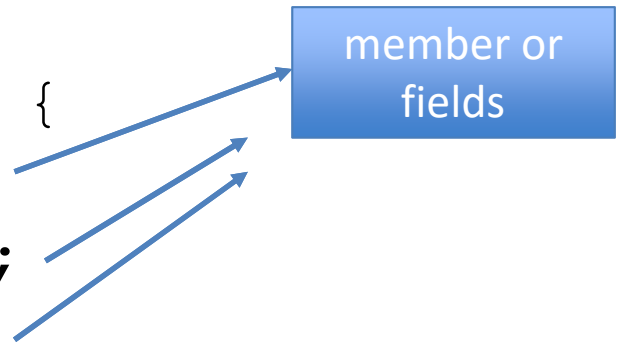
- **Defines a new type**
  - A new kind of data type that the compiler regards as a unit or aggregate of variables/types.
- **Example:**

```
struct Date {  
    int day;  
    int month;  
    int year;  
};
```

# C struct

- **Defines a new type**
  - A new kind of data type that the compiler regards as a unit or aggregate of variables/types.
- **Example:**

```
struct Date {  
    int day;  
    int month;  
    int year;  
};
```

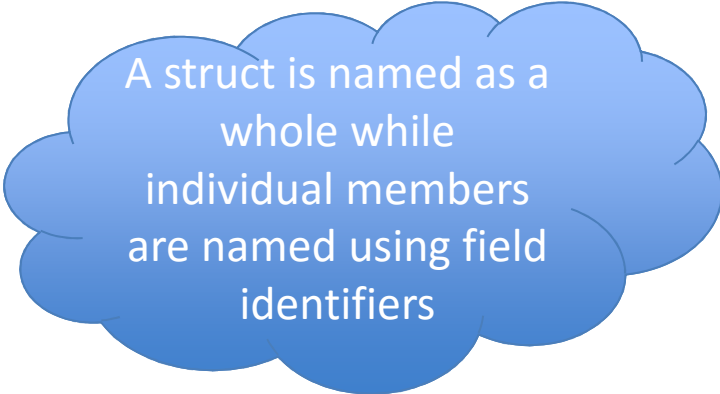


A diagram illustrating the components of a C struct. Three blue arrows originate from the members 'day', 'month', and 'year' within the 'struct Date' definition and point towards a blue rectangular box on the right. The box contains the text 'member or fields' in white, indicating that these are the individual components that make up the struct.

# C struct

- **Defines a new type**
  - A new kind of data type that the compiler regards as a unit or aggregate of variables/types.
- **Example:**

```
struct Date {  
    int day;  
    int month;  
    int year;  
};
```



A struct is named as a whole while individual members are named using field identifiers

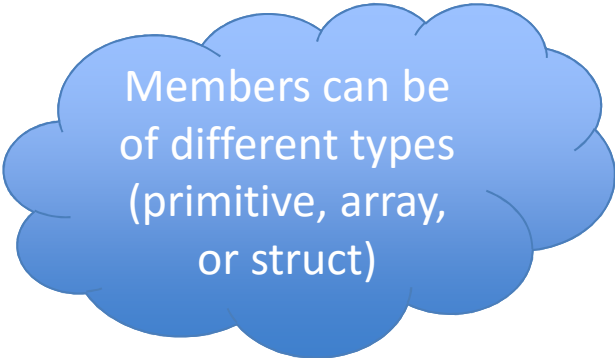


# More struct Examples

- **Examples:**

```
struct Date {  
    int day;  
    int month;  
    int year;  
};
```

```
struct Employee{  
    char ename[20];  
    int ssn;  
    float salary;  
    struct Date doj;  
};
```



Members can be  
of different types  
(primitive, array,  
or struct)

# More struct Examples

- **Examples:**

```
struct Date {  
    int day;  
    int month;  
    int year;  
};
```

```
struct Employee[3];
```

```
struct Employee{  
    char ename[20];  
    int ssn;  
    float salary;  
    struct Date doj;  
};
```

# Declaring a `struct` Variable

- Declaration of a variable of struct type:

**<struct type> <identifier list>;**

- **Example:**

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};  
struct StudentRecord student1;
```

# Declaring a `struct` Variable

- Declaration of a variable of struct type:

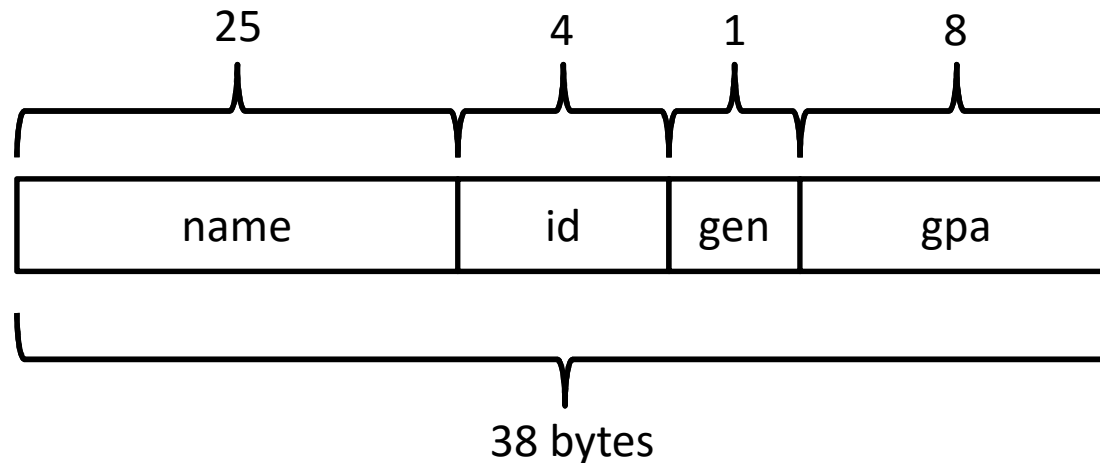
**<struct type> <identifier list>;**

- **Example:**

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};  
struct StudentRecord student1;  
... = student1.name  
student1.gpa = ...
```

# sizeof(struct StudentRecord)=?

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```



# student-01.c example

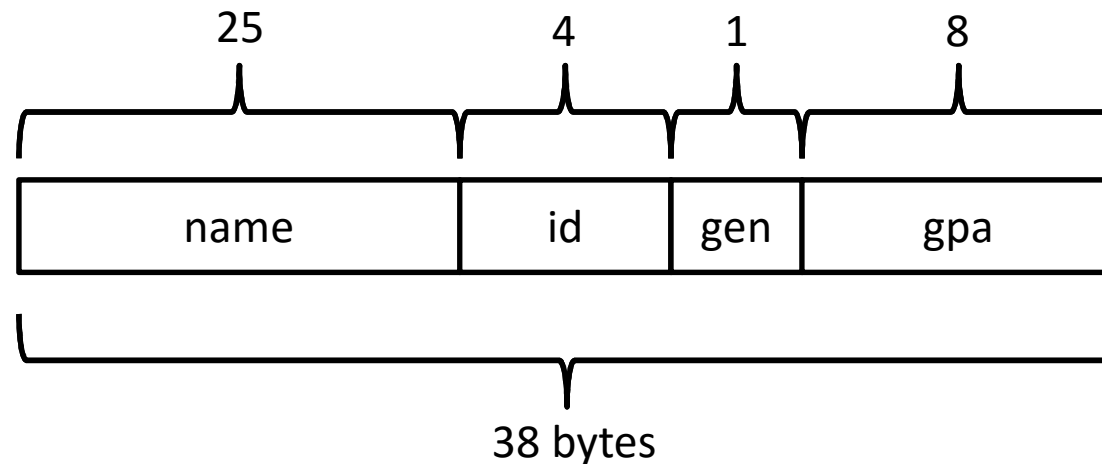
- **Let us compile this example**
  - What did you notice about the output of this program?
  - Is the size of this struct the same as we predicted?
  - Why is this or is this not the case?

# student-01.c example

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```



sizeof(struct  
StudentRecord) = 48



# Data Allocation and Alignment

- **Data Allocation**

- Each variable definition is allocated bytes in memory according the type of that variable
- e.g., char = 1 byte, int = 4 bytes, double = 8 bytes
- This is allocated in a special place in memory known as the **stack**.

- **Data Alignment**

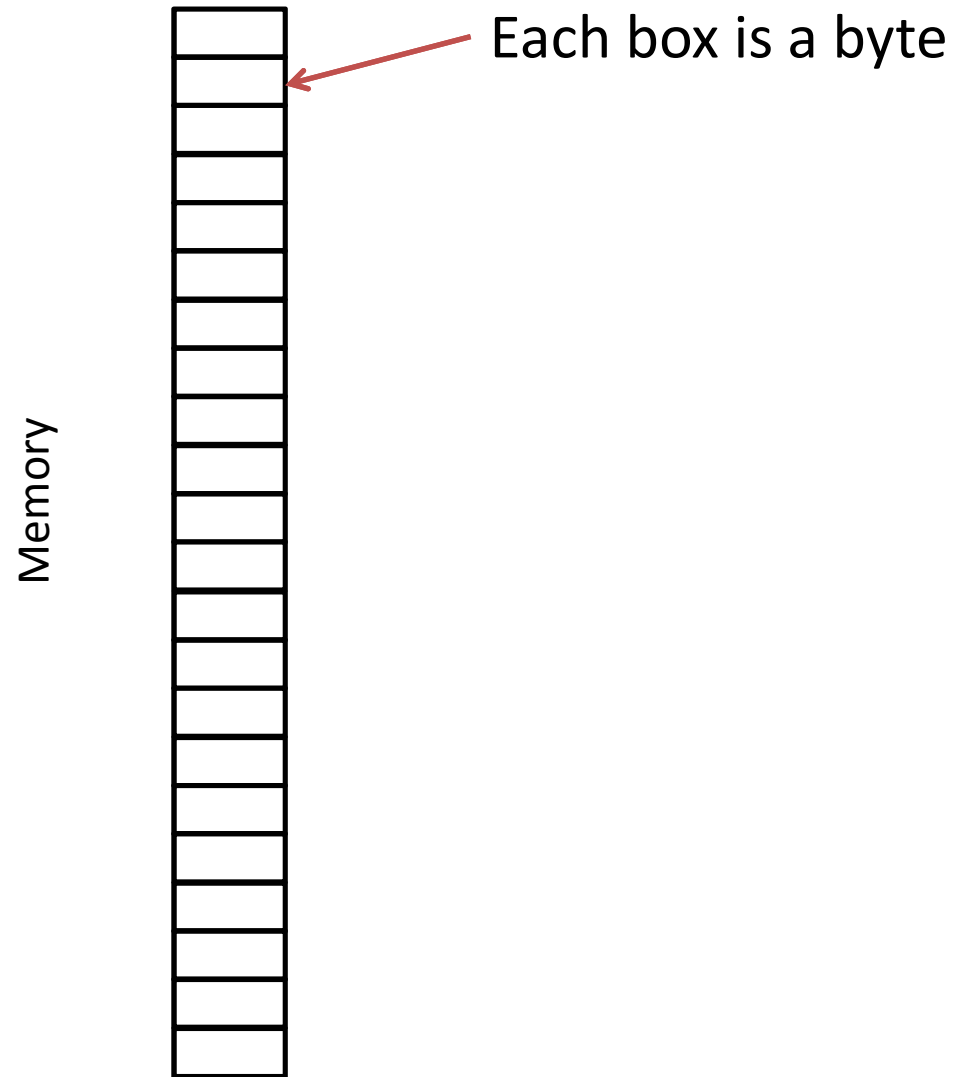
- Alignment helps the memory store data in a structured manner
  - e.g. 2-byte shorts must start on an even address
- Machines are more efficient if allocated data is accessed in a structured manner



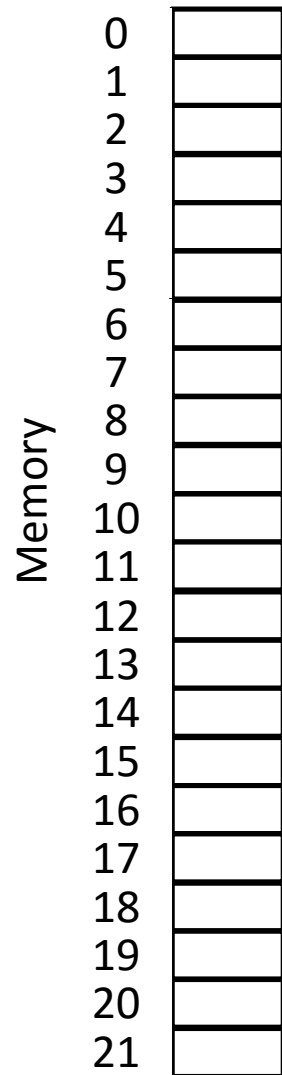
# Data Alignment

- All types except char doesn't normally start at an arbitrary address.
- A 1-byte char can start on any byte address
- A 2-byte short must start on an even address
- A 4-byte int/float must start on an address divisible by 4
- A 8-byte long/double must start on address divisible by 8

# Alignment in Memory



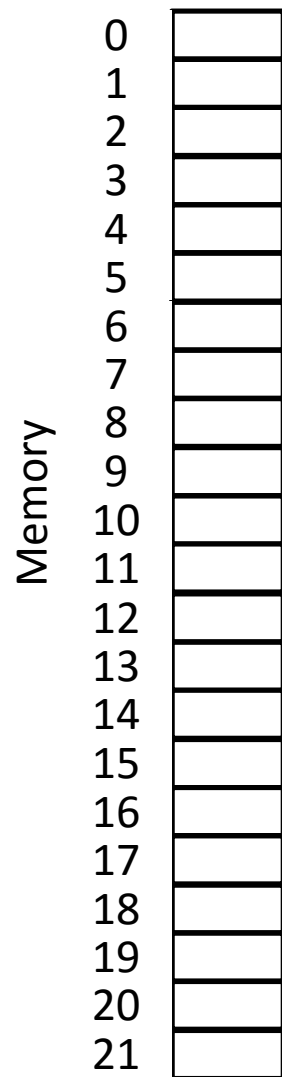
# Alignment in Memory



Each box is a byte  
**and** has a location.

Memory is very much like a a  
giant character array!

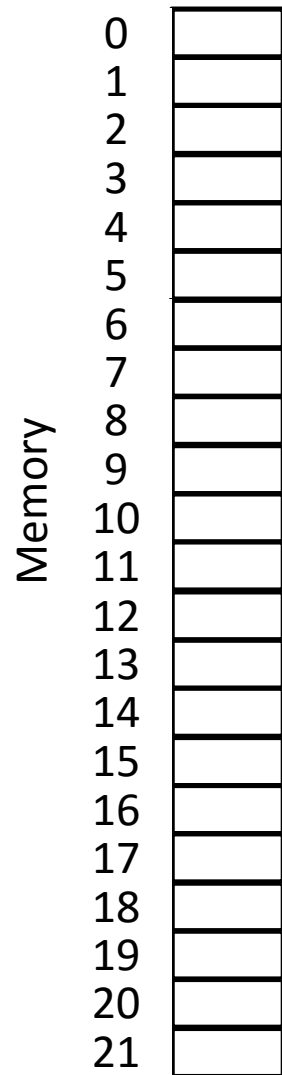
# Alignment in Memory



So, how do we allocate memory  
for the variable declarations  
below?

```
long a;  
char b;  
int c;
```

# Alignment in Memory

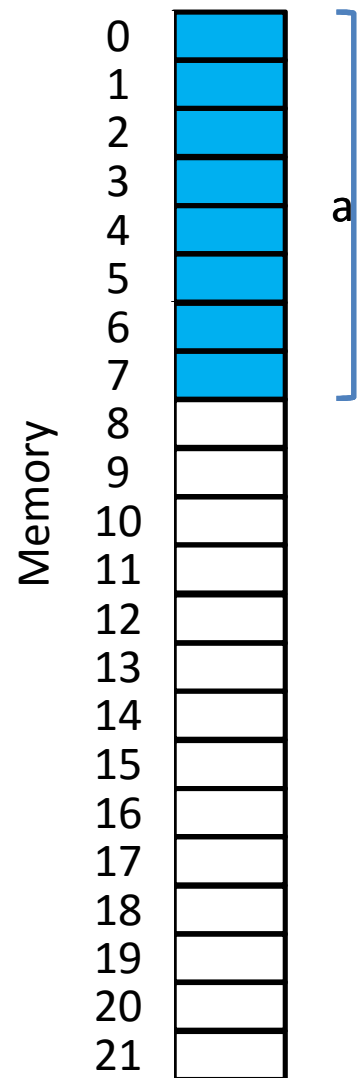


So, how do we allocate memory  
for the variable declarations  
below?

```
long a;  
char b;  
int c;
```

$8 + 1 + 4 = 13$  bytes ??

# Alignment in Memory

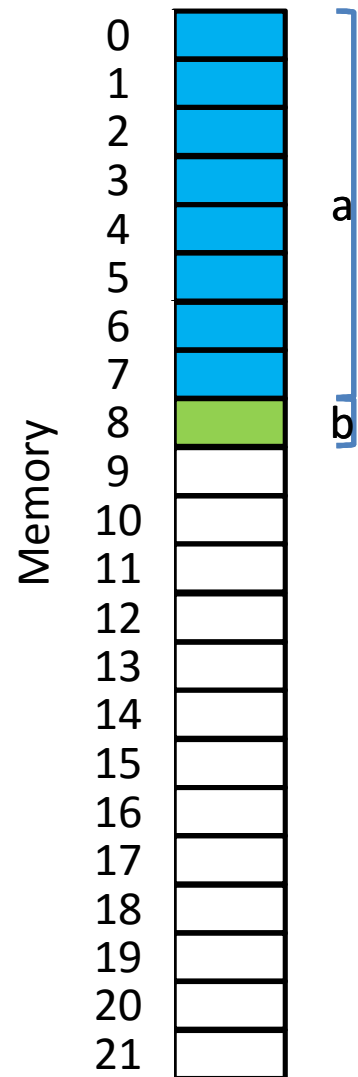


So, how do we allocate memory for the variable declarations below?

```
long a;  
char b;  
int c;
```

$8 + 1 + 4 = 13$  bytes ??

# Alignment in Memory

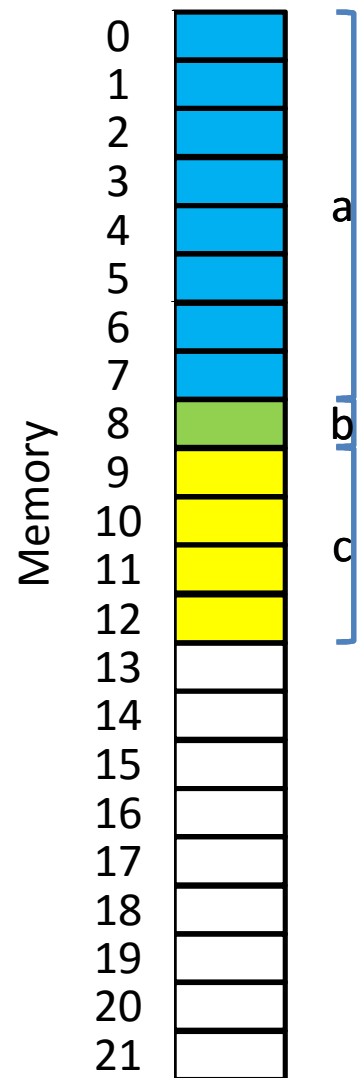


So, how do we allocate memory  
for the variable declarations  
below?

```
long a;  
char b;  
int c;
```

$8 + 1 + 4 = 13$  bytes ??

# Alignment in Memory



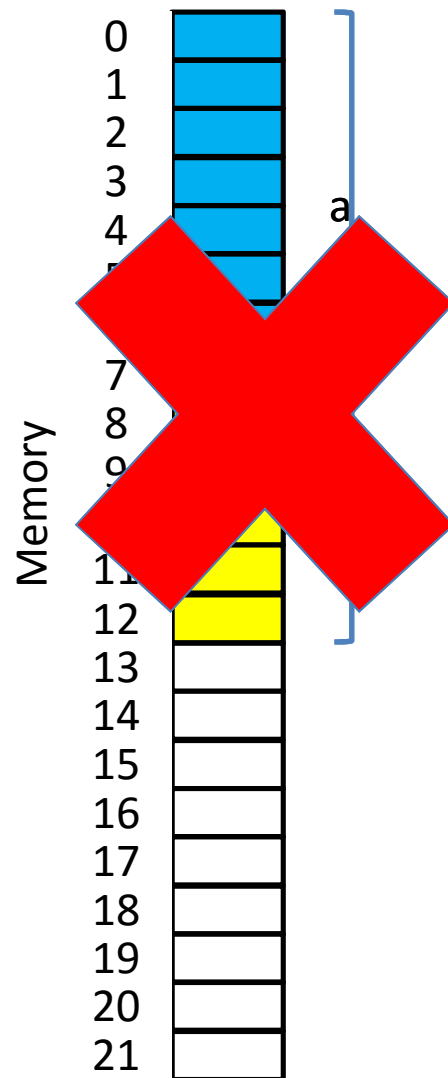
So, how do we allocate memory  
for the variable declarations  
below?

```
long a;  
char b;  
int c;
```

$8 + 1 + 4 = 13$  bytes ??



# Alignment in Memory

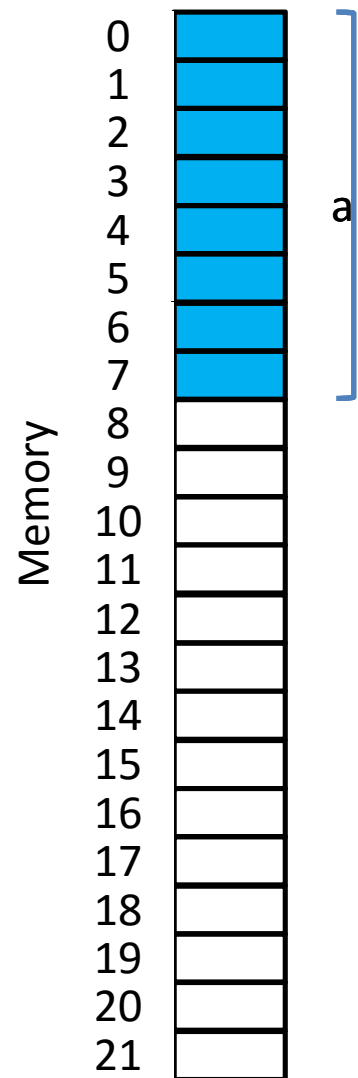


So, how do we allocate memory for the variable declarations below?

```
long a;  
char b;  
int c;
```

$8 + 1 + 4 = 13$  bytes ??

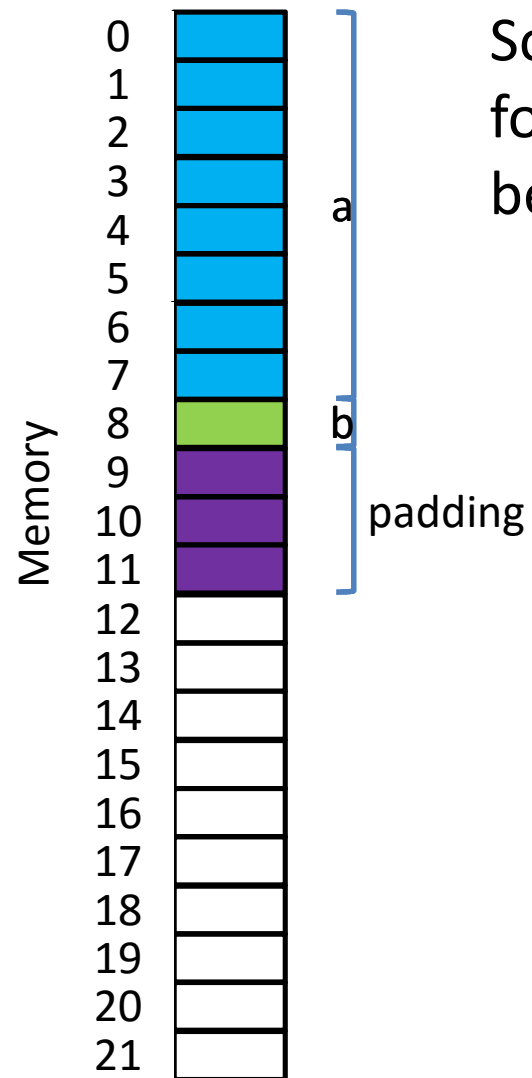
# Alignment in Memory



So, how do we allocate memory  
for the variable declarations  
below?

```
long a;  
char b;  
int c;
```

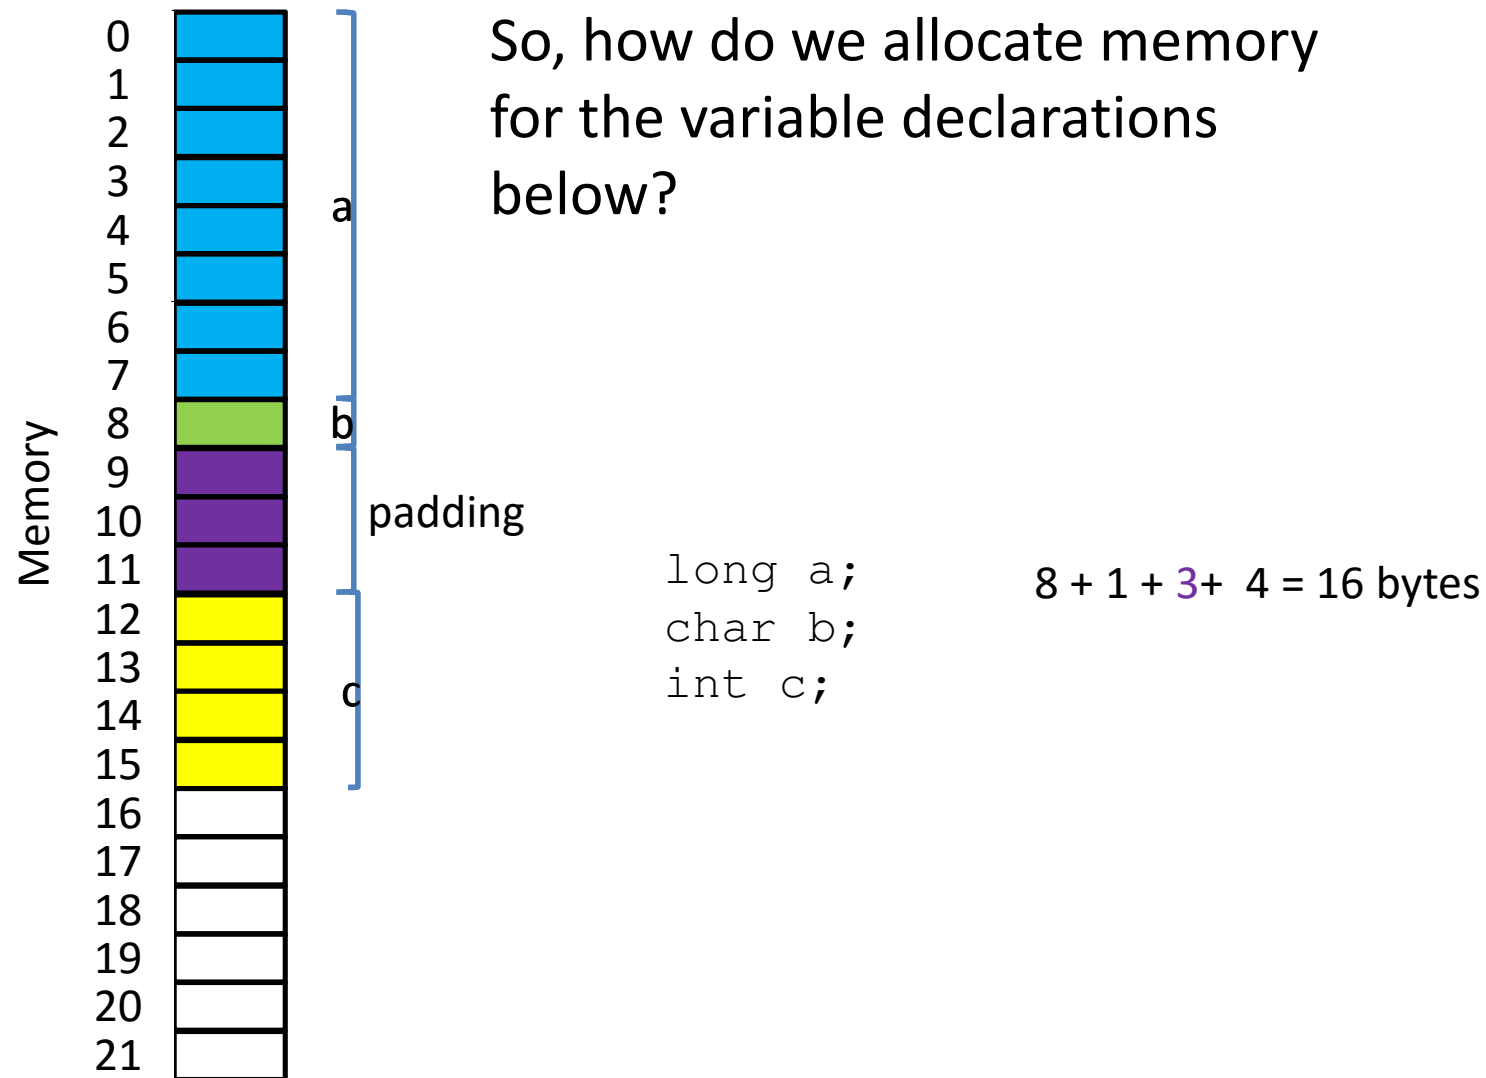
# Alignment in Memory



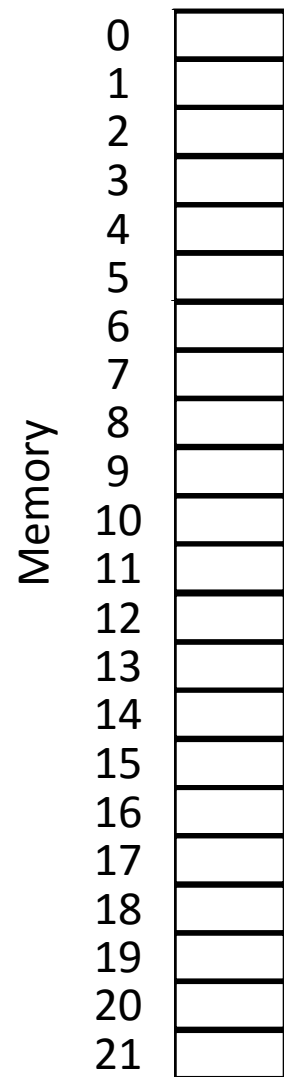
So, how do we allocate memory  
for the variable declarations  
below?

```
long a;  
char b;  
int c;
```

# Alignment in Memory



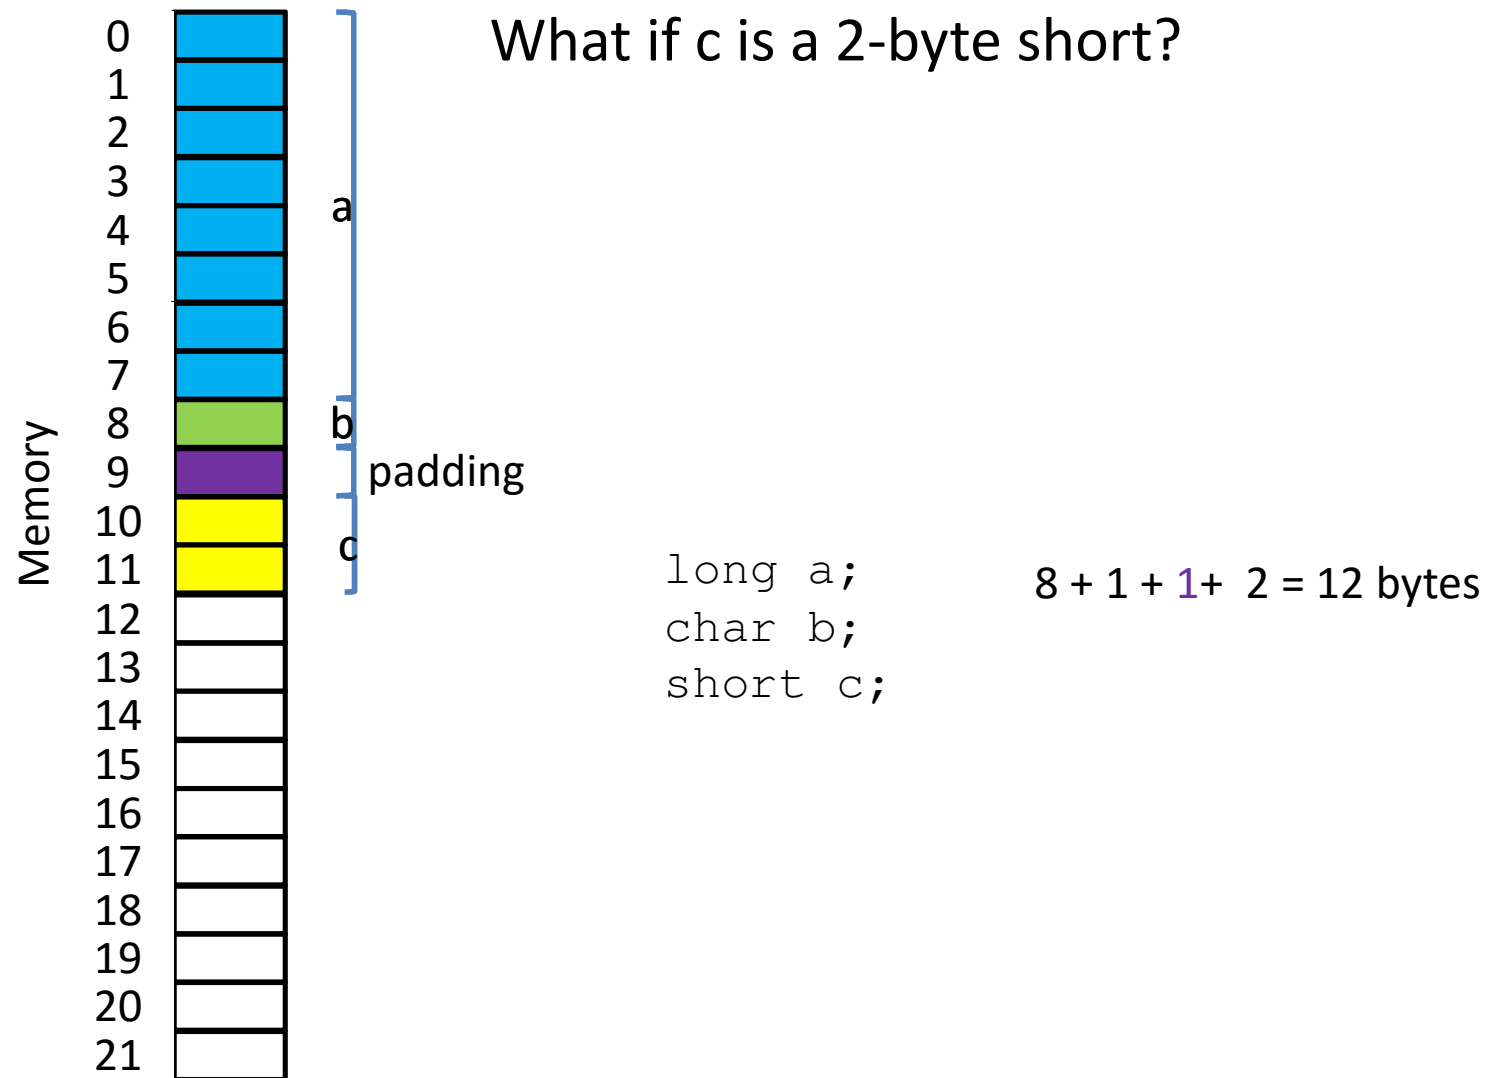
# Alignment in Memory



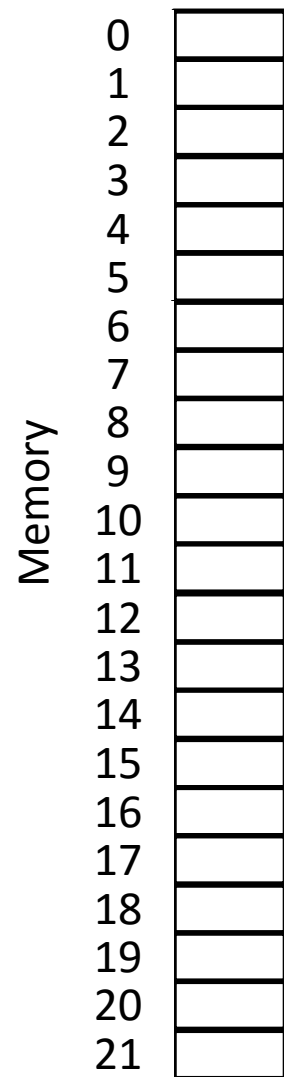
What if c is a 2-byte short?

```
long a;  
char b;  
short c;
```

# Alignment in Memory



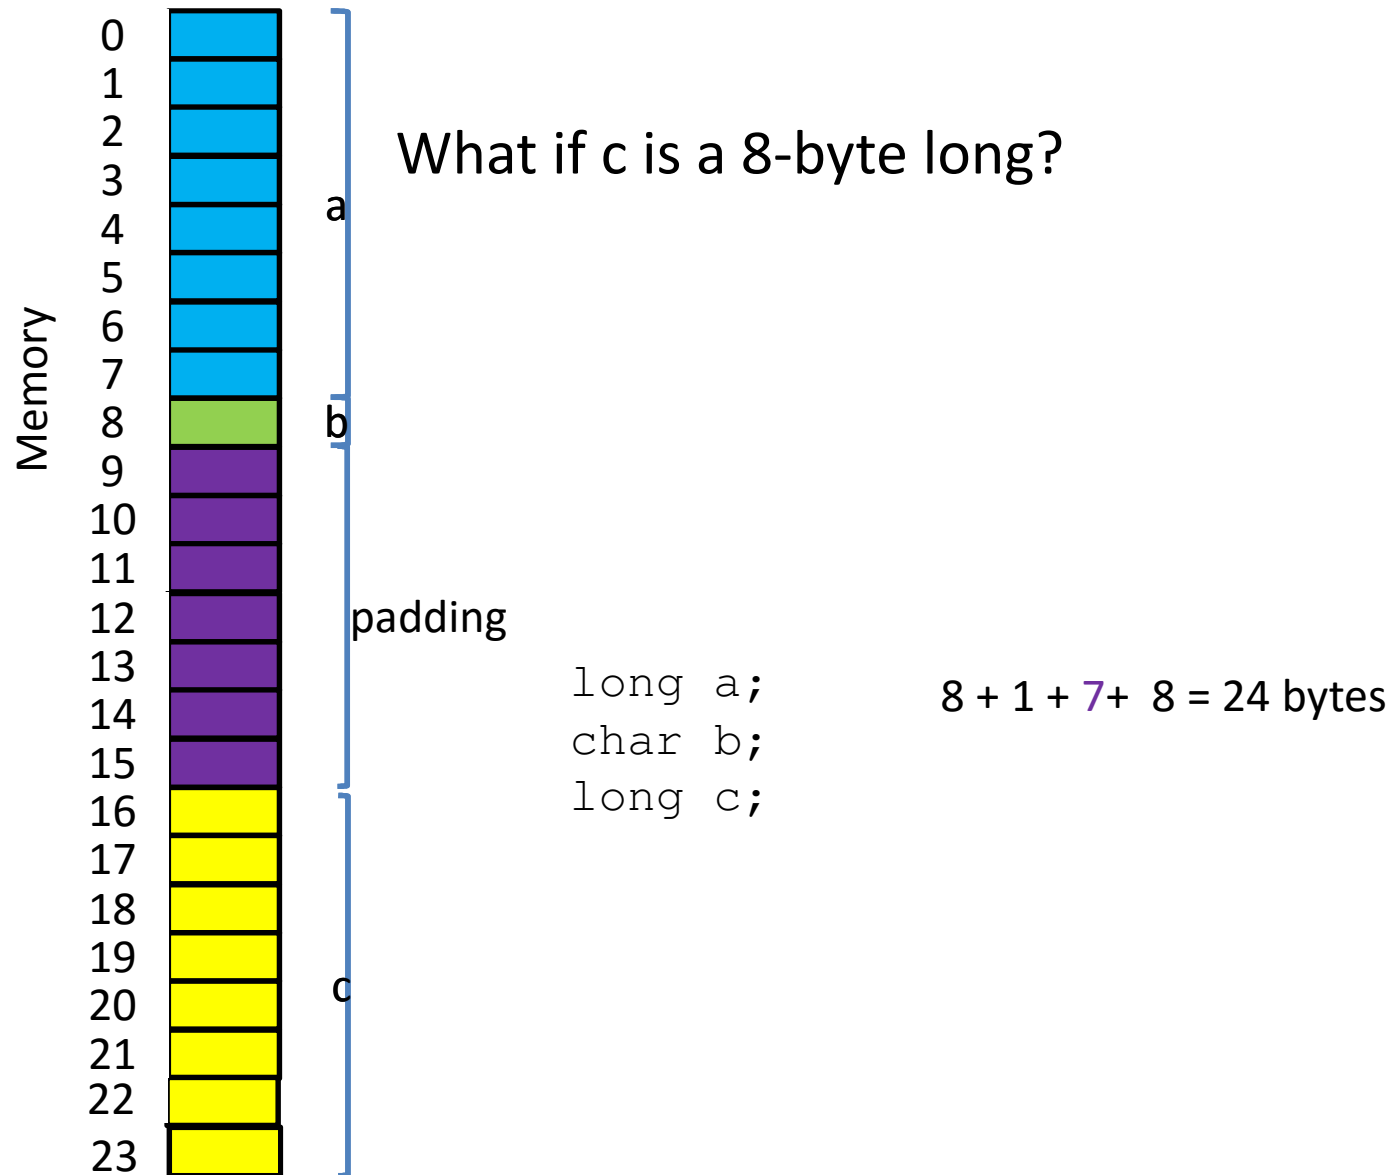
# Alignment in Memory



What if c is a 8-byte long?

```
long a;  
char b;  
long c;
```

# Alignment in Memory





# i-clicker question

What is the size of padding in bytes between a and c on a 64-bit machine?

```
char b;  
long a;  
int c;
```

- A. 0
- B. 1
- C. 2
- D. 3

# i-clicker question

What is the size of padding in bytes between b and a on a 64-bit machine??

```
char b;  
long a;  
int c;
```

- A. 0
- B. Determined by the address of b
- C. 7
- D. 3

# Reduce memory size

- We can reduce used memory size by reorder declaration statements

```
char b;  
long a;  
int c;
```



```
long a;  
int c;  
char b;
```

$1 + 8 + 4 + ? = 13 + ?$   
bytes

$8 + 4 + 1 = 13$   
bytes

# Alignment in Memory

- There is ***no*** internal padding within a primitive type array
- But **there can be** internal padding within a struct
- **There can be** internal padding within struct array.

# Structure alignment

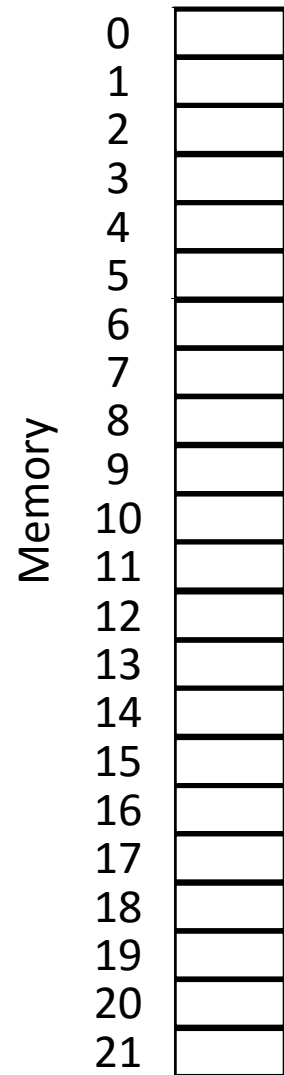
- Structure alignment requirements:
  - Within a struct, each member has different alignment requirements.
  - The structure as a whole has the alignment of its widest scalar member

# Structure alignment

- Structure alignment requirements:
  - Within a struct, each member has different alignment requirements.
  - The structure as a whole have the alignment of its widest scalar member
- Find memory layout
  - Step 1: Align each member as if these member has been separately declared
  - Step 2: Pad at the end so the structure has the alignment of its widest scalar member

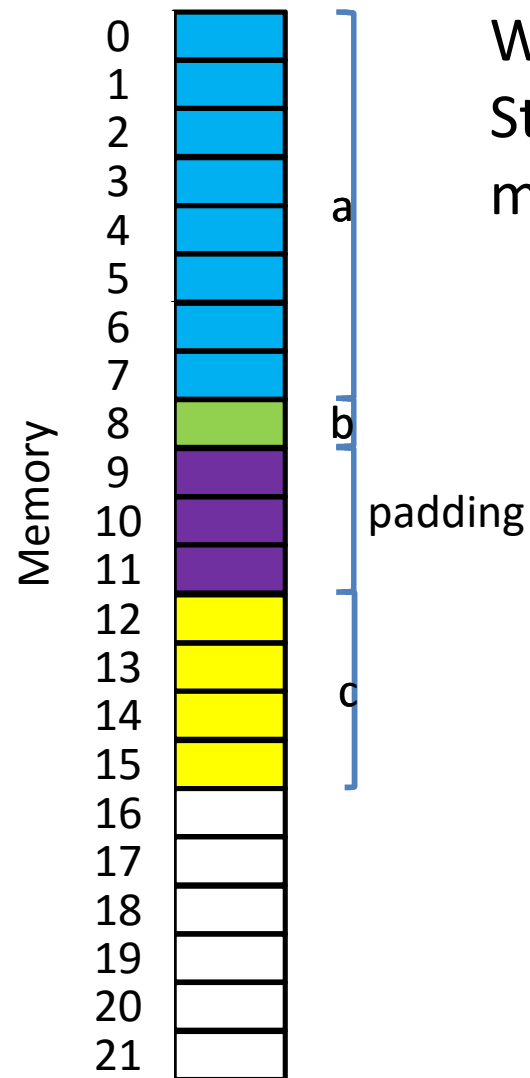
# Structure alignment

What's the memory layout of struct foo?



```
struct foo {  
    long a;  
    char b;  
    int c;  
};
```

# Structure alignment



What's the memory layout of struct foo?  
Step 1: Align each member as if these member has been separately declared

```
struct foo {  
    long a;  
    char b;  
    int c;  
};
```



# Structure alignment



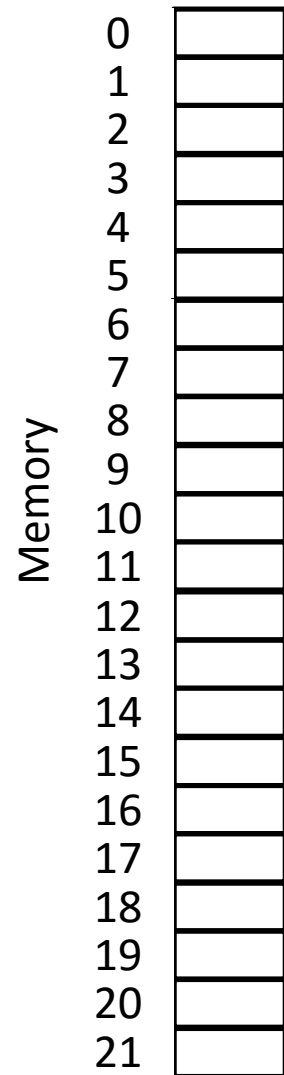
What's the memory layout of struct foo?

Step 1: Align each member as if these member has been separately declared

Step 2: Pad at the end so the structure has the alignment of its widest scalar member

```
struct foo {  
    long a;  
    char b;  
    int c;  
};
```

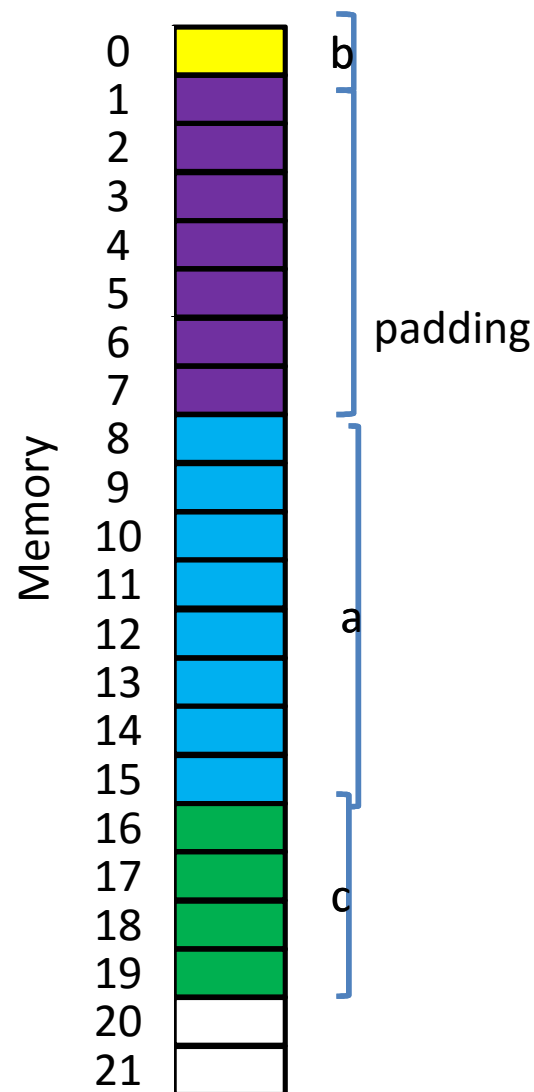
# Structure alignment



What's the memory layout of struct bar?

```
struct bar{  
    char b;  
    long a;  
    int c;  
};
```

# Structure alignment

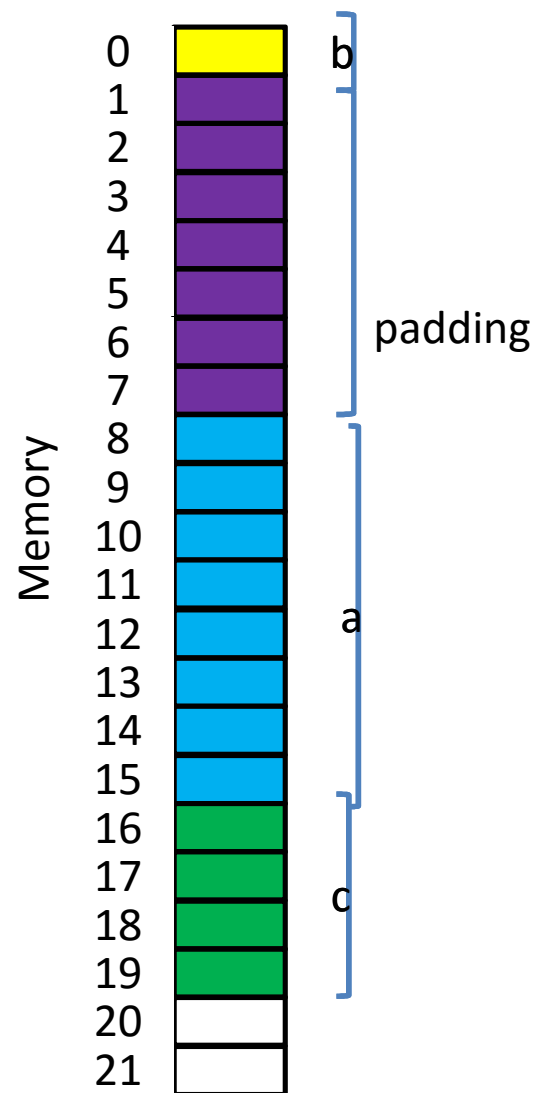


What's the memory layout of struct bar?

Step 1: Align each member as if these member has been separately declared

```
struct bar{  
    char b;  
    long a;  
    int c;  
};
```

# Structure alignment



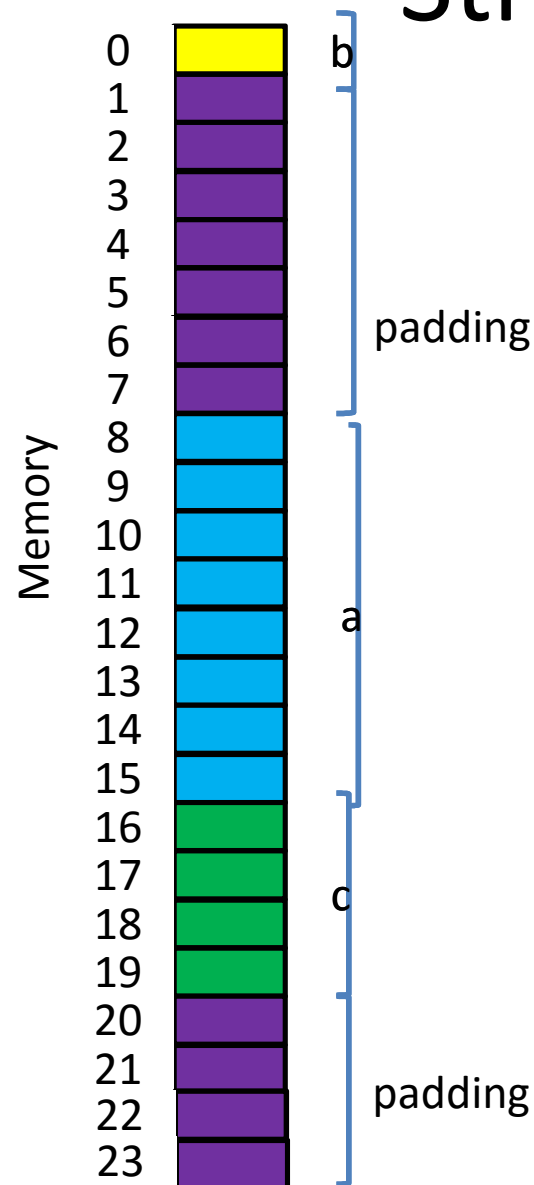
What's the memory layout of struct bar?

Step 1: Align each member as if these member has been separately declared

Step 2: Pad at the end so the structure has the alignment of its widest scalar member

```
struct bar{  
    char b;  
    long a;  
    int c;  
};
```

# Structure alignment



What's the memory layout of struct bar?

Step 1: Align each member as if these member has been separately declared

Step 2: Pad at the end so the structure has the alignment of its widest scalar member

```
struct bar{  
    char b;  
    long a;  
    int c;  
};
```

# i-clicker question

What is the size of the following structure on a 64-bit machine?

```
struct baz {  
    long a;  
    char c;  
};
```

- A. 9
- B. 24
- C. 20
- D. 16

# i-clicker question

What is the size of the following structure on a 64-bit machine?

```
struct qux {  
    char a;  
    char b;  
    char c;  
    char d;  
    char e;  
};
```

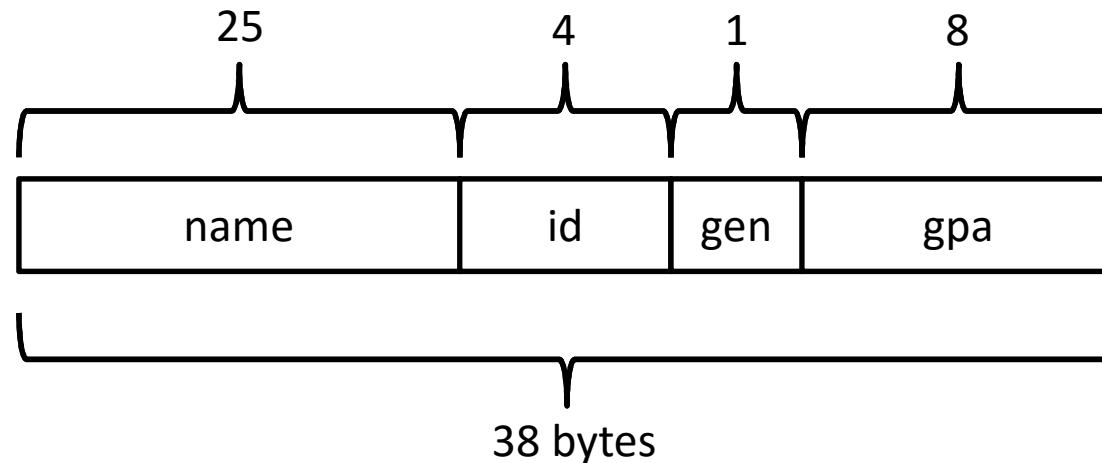
- A. 10
- B. 20
- C. 5
- D. 8

# student-01.c example

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```

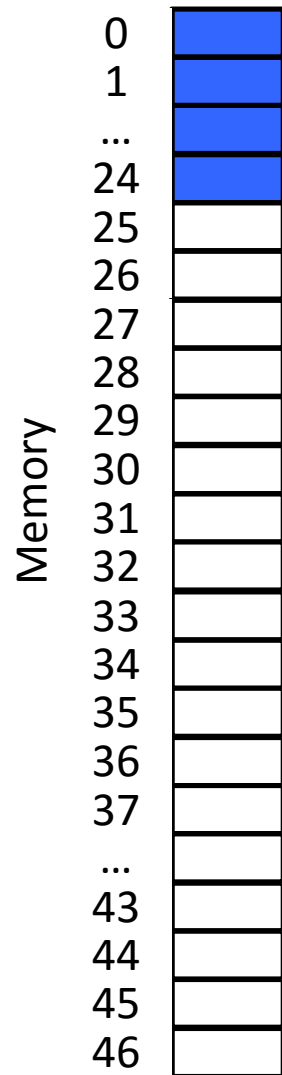


sizeof(struct  
StudentRecord) = 48





# Alignment in Memory

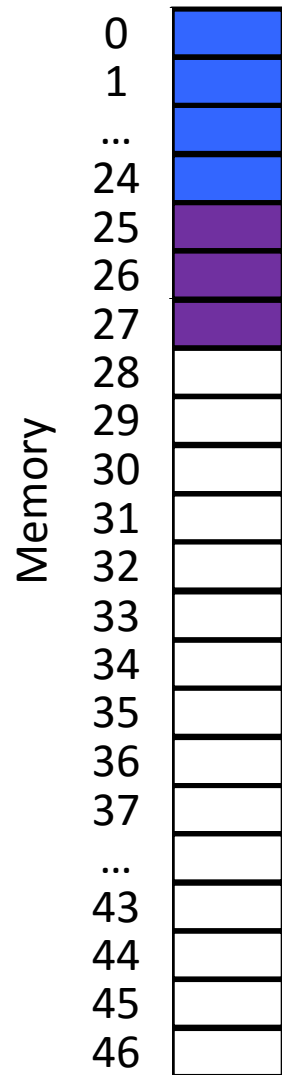


What's the memory layout of struct `StudentRecord`?  
Step 1: Align each member as if these member has been separately declared

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```

```
struct StudentRecord student1;
```

# Alignment in Memory

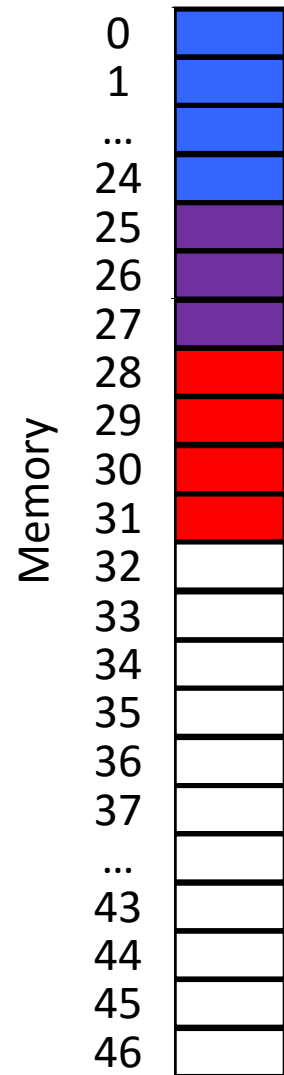


What's the memory layout of struct `StudentRecord`?  
Step 1: Align each member as if these member has been separately declared

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```

```
struct StudentRecord student1;
```

# Alignment in Memory

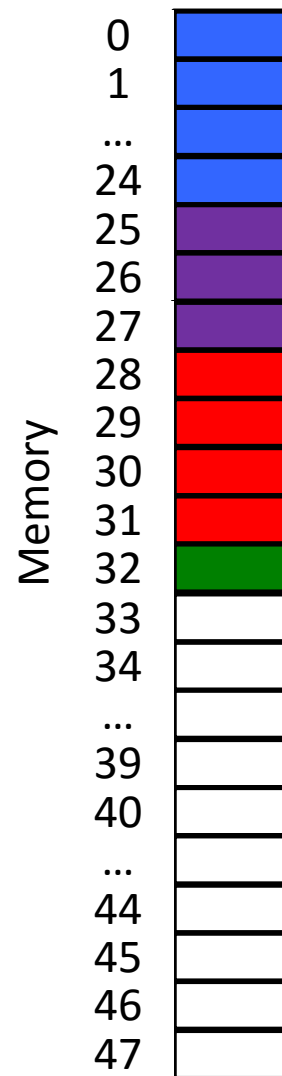


What's the memory layout of struct `StudentRecord`?  
Step 1: Align each member as if these member has been separately declared

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```

```
struct StudentRecord student1;
```

# Alignment in Memory

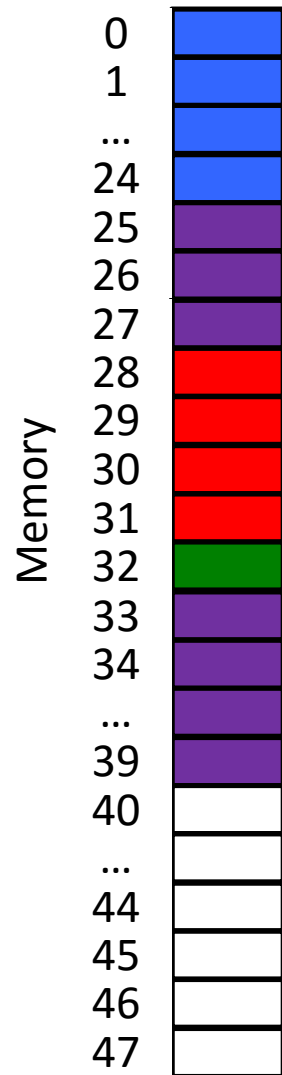


What's the memory layout of struct `StudentRecord`?  
Step 1: Align each member as if these member has been separately declared

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```

```
struct StudentRecord student1;
```

# Alignment in Memory

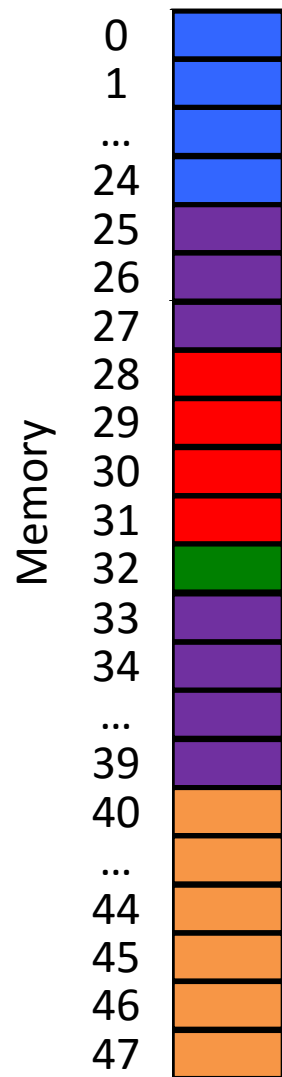


What's the memory layout of struct `StudentRecord`?  
Step 1: Align each member as if these member has been separately declared

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```

```
struct StudentRecord student1;
```

# Alignment in Memory

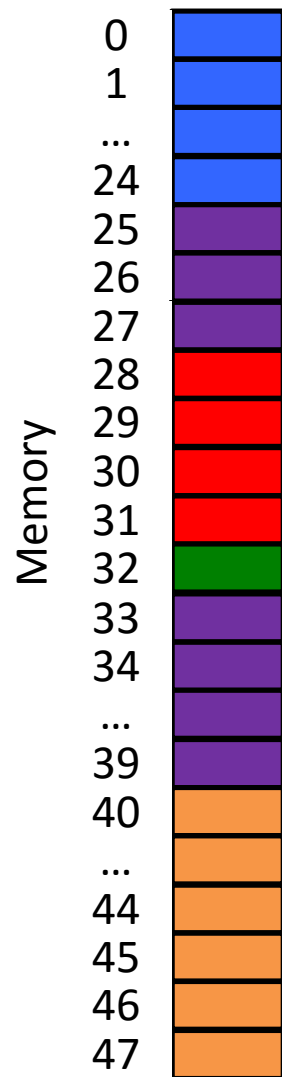


What's the memory layout of struct `StudentRecord`?  
Step 1: Align each member as if these member has been separately declared

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```

```
struct StudentRecord student1;
```

# Alignment in Memory



What's the memory layout of struct `StudentRecord`?  
Step 1: Align each member as if these member has been separately declared

Step 2: Pad at the end so the structure has the alignment of its widest scalar member

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```

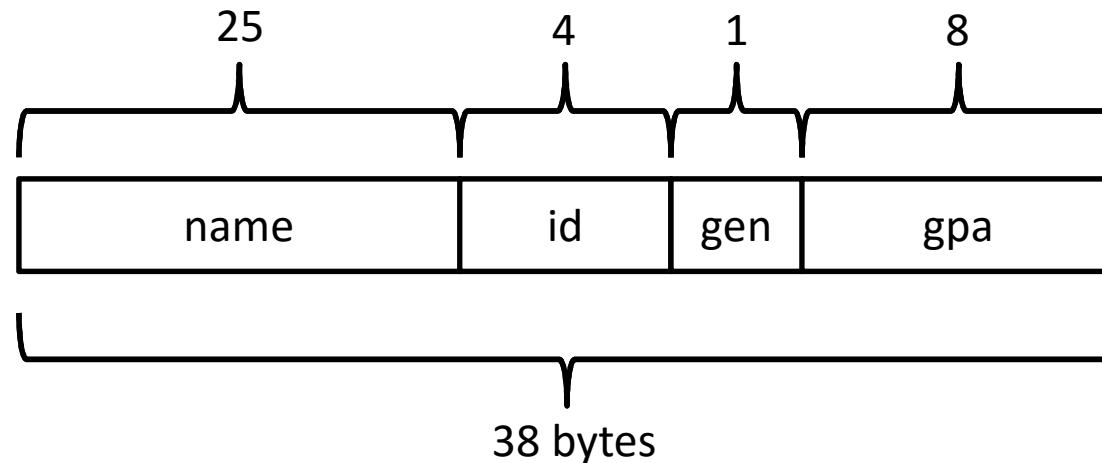
```
struct StudentRecord student1;
```

# student-01.c example

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```



sizeof(struct  
StudentRecord) = 48





# student-01\_offsetof.c

- Explore the memory layout of a struct by yourself:
  - `offsetof (type,member)`

# Structure Initialization

- **There are four ways to initialize a struct**
  - Positional initialization
  - Named initialization
  - Copy initialization
  - Initialize individual fields

# Positional Initialization

**Positional initialization** allows you to provide the values for each of the fields based on the position of each structure member:

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};  
  
struct StudentRecord student1 = {  
    "John Doe", 1234567, 'M', 3.95  
};
```

Not a recommended way!

# Positional Initialization

**Positional initialization** allows you to provide the values for each of the fields based on the position of each structure member:

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};  
  
struct StudentRecord student1 = {  
    "John Doe", 1234567, 'M', 3.95  
};
```

Not a recommended way!

# Named Initialization

**Named initialization** allows you to provide the values for each of the fields based on the name of each structure member:

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```

```
struct StudentRecord student1 = {  
    .id      = 1234567,  
    .gpa     = 3.95,  
    .gender  = 'M',  
    .name= "Harry Potter"  
};
```

# Copy Initialization

**Copy initialization** allows you to initialize a structure by assigning an existing structure:

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```

```
struct StudentRecord student1 = {  
    .id          = 1234567,  
    .gpa         = 3.95,  
    .gender      = 'M',  
    .name = "Harry Potter"  
};
```

```
struct StudentRecord student2 = student1;
```

# Field Initialization

**Field initialization** allows you to initialize a structure by assigning to its fields:

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```

```
struct StudentRecord student1;
```

```
student1.id          = 1234567;  
student1.gender      = 'M';  
student1.gpa         = 3.95;
```

# Field Initialization

**Field initialization** allows you to initialize a structure by assigning to its fields:

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```

```
struct StudentRecord student1;
```

```
student1.id          = 1234567;  
student1.gender      = 'M';  
student1.gpa         = 3.95;  
student1.name        = "Harry Potter";
```

**What about this one?**



# student-05.c example

- **Let us compile this example**
  - What problems do we encounter with this example?
  - Why can't we assign a string to a character array?
    - Arrays are not modifiable values, that is, you can't reassign them to "point" to different locations in memory.

# Field Initialization

**Field initialization** allows you to initialize a structure by assigning to its fields:

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```

```
struct StudentRecord student1;
```

```
student1.id          = 1234567;  
student1.gender      = 'M';  
student1.gpa         = 3.95;  
student1.name        = "Harry Potter";
```

**So, how do we fix this?**

# strncpy

- **Copying Strings**

- #include <string.h>

- A library for manipulating C strings

- To assign a new string value to a C string (e.g., character array) you must use the *strncpy* function to **copy** the bytes into the array.

# Field Initialization

**Field initialization** allows you to initialize a structure by assigning to its fields:

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```

```
struct StudentRecord student1;
```

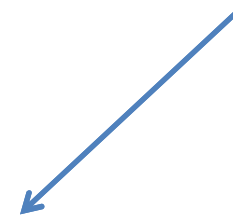
```
student1.id          = 1234567;
```

```
student1.gender      = 'M';
```

```
student1.gpa         = 3.95;
```

```
strncpy(student1.name, "Harry Potter", 25);
```

Size of the *destination*

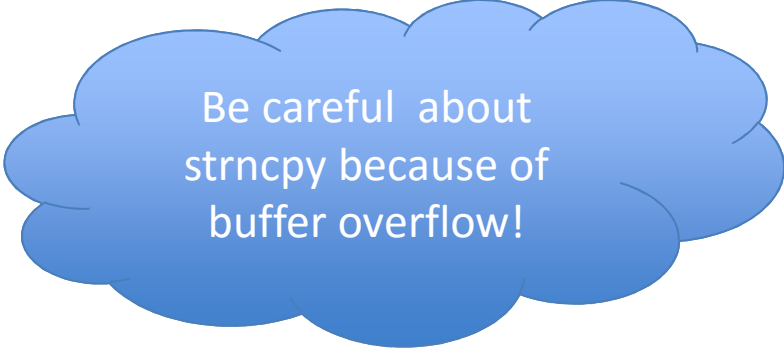


**We use the strncpy function!**

# Field Initialization

**Field initialization** allows you to initialize a structure by assigning to its fields:

```
struct StudentRecord {  
    char name[25];  
    int id;  
    char gender;  
    double gpa;  
};
```



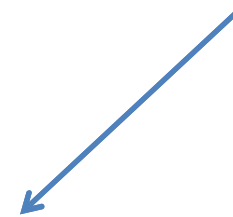
Be careful about  
strncpy because of  
buffer overflow!

```
struct StudentRecord student1;
```

```
student1.id          = 1234567;  
student1.gender      = 'M';  
student1.gpa         = 3.95;
```

```
strncpy(student1.name, "Harry Potter", 25);
```

Size of the *destination*



**We use the strncpy function!**

# strlen

- **Calculate the length of a string**
  - #include <string.h>  
A library for manipulating C strings
  - Return an unsigned integer
  - Example:
    - strlen("Revenant") == 8

# Buffer overflow example

```
void f(char[] bar)
{
    char c[12];
    strncpy(c, bar, strlen(bar));
}
```