

Computer Systems Principles

Linking and Loading

Learning Objectives

- Learn about the linker and its uses
- Learn to identify types of linking
- Learn about the creation and contents of an Executable and Linkable (ELF) file
- Learn how symbol resolution is performed inside the linker
- Learn how linking is performed in Static libraries
- Learn how linking is performed in Dynamic libraries

Why Linkers?

- Modularity
 - Program can be written as a collection of smaller source files, rather than one monolithic mass.
 - Can build libraries of common functions
 - e.g., Math library, standard C library, etc.

Why Linkers?

- Efficiency
 - Time: **Separate compilation**
 - Change one source file, compile, and then re-link.
 - No need to recompile other source files.
 - Space: **Libraries**
 - Common functions can be aggregated into a single file
 - Yet executable files and running memory images contain code only for the functions they actually use.

Example C Program

swap.c

main.c

```
int buf[2] = {1, 2};

int main()
{
    swap();
    return 0;
}
```

```
extern int buf[];

int *bufp0 = &buf[0];
static int *bufp1;

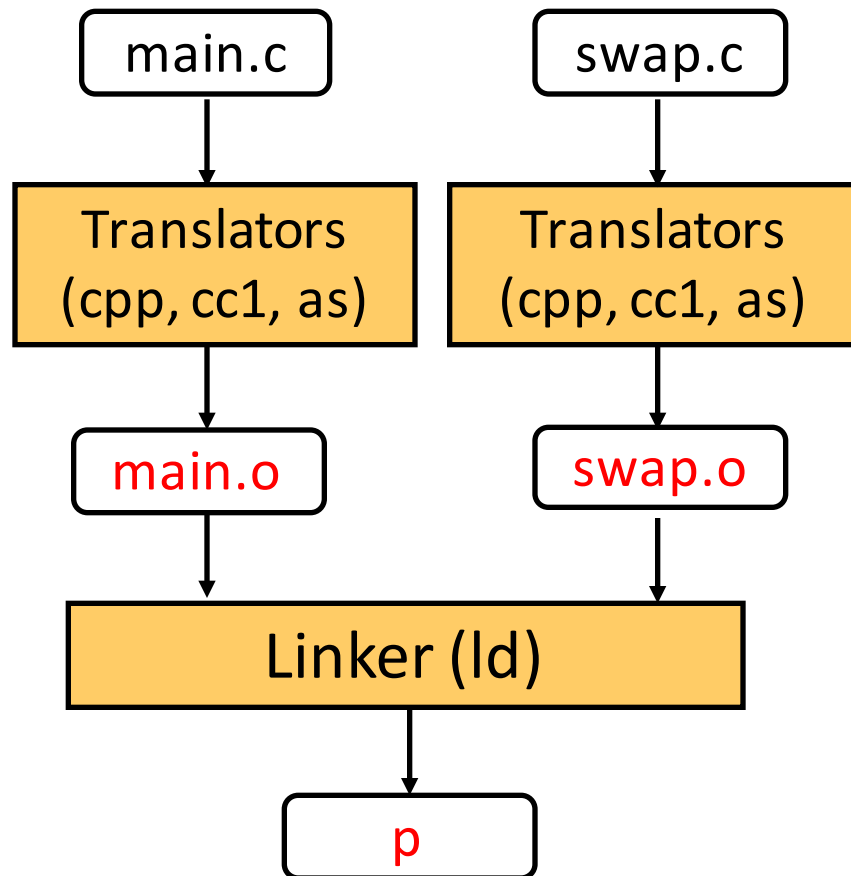
void swap()
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

Static Linking

Programs are translated and linked using a **compiler driver**:

- > gcc main.c swap.c -o p
- > ./p



Source files

Separately compiled
relocatable object files

Fully linked **executable object file** (contains code and data for all functions defined in `main.c` and `swap.c`)

What Do Linkers Do?

- Symbol Resolution [step 1]
 - Programs define and reference symbols (**global variables and functions**):
 - `void swap() {...}` // define symbol swap
 - `swap();` // reference symbol swap
 - `int *xp = &x;` // define symbol xp, reference x
 - Symbol definitions are stored in object file (by assembler) in a part called the **symbol table**.
 - The symbol table is an **array of structs**
 - **Each struct entry** includes name, size, and location of symbol.
 - Linker associates each symbol reference with **exactly one** symbol definition.

What Do Linkers Do?

- Relocation [step2]
 - **Merges** separate code and data sections into single sections
 - **Relocates symbols** from their relative locations in the **.o** files to their final **absolute memory locations** in the executable.
 - **Updates** all references to these symbols to reflect their new positions.

Let's look at these two steps in more detail....

Three Kinds of Object Files (Modules)

- Relocatable object file (.o file)
 - Contains code and data in a form that can be combined with other relocatable object files to form executable object file.
 - Each .o file is produced from exactly one source .c file
- Executable object file (a.out file)
 - Contains code and data in a form that can be copied directly into memory and then executed.
- Shared object file (.so file)
 - Special type of relocatable object file that can be loaded into memory and linked dynamically, at either load time or run time.
 - Called *Dynamic Link Libraries* (DLLs) by Windows

Executable and Linkable Format (ELF)

- Standard binary format for object files
- One unified format for
 - Relocatable object files (.o),
 - Executable object files (a.out)
 - Shared object files (.so)
- Generic name: ELF binaries

ELF Object File Format

- ELF header
 - Word size, byte ordering, file type (.o, exec, .so), machine type, etc.
- Segment header table
 - Page size, virtual addresses of memory segments (sections), segment sizes.

ELF header
Segment header table (required for executables)
.text section
.rodata section
.data section
.bss section
.symtab section
.rel.txt section
.rel.data section
.debug section
Section header table

ELF Object File Format

- **.text** section
 - Code
- **.rodata** section
 - Read only data: jump tables, ...
- **.data** section
 - Initialized global variables
- **.bss** section
 - Uninitialized global variables
 - “Block Started by Symbol”
 - “Better Save Space”
 - Has section header but **occupies no space** in the ELF file (space used at run time)

ELF header	0
Segment header table (required for executables)	
.text section	
.rodata section	
.data section	
.bss section	
.symtab section	
.rel.txt section	
.rel.data section	
.debug section	
Section header table	

ELF Object File Format

- **.symtab** section
 - Symbol table
 - Procedure and static variable names
 - Section names and locations
- **.rel.text** section
 - Relocation info for **.text** section
 - Addresses of instructions that will need to be modified in the executable
 - Directions for modifying them
- **.rel.data** section
 - Relocation info for **.data** section
 - Addresses of pointer data that will need to be modified in the merged executable
- **.debug** section
 - Info for symbolic debugging (**gcc -g**)
- Section header table
 - Offsets and sizes of each section

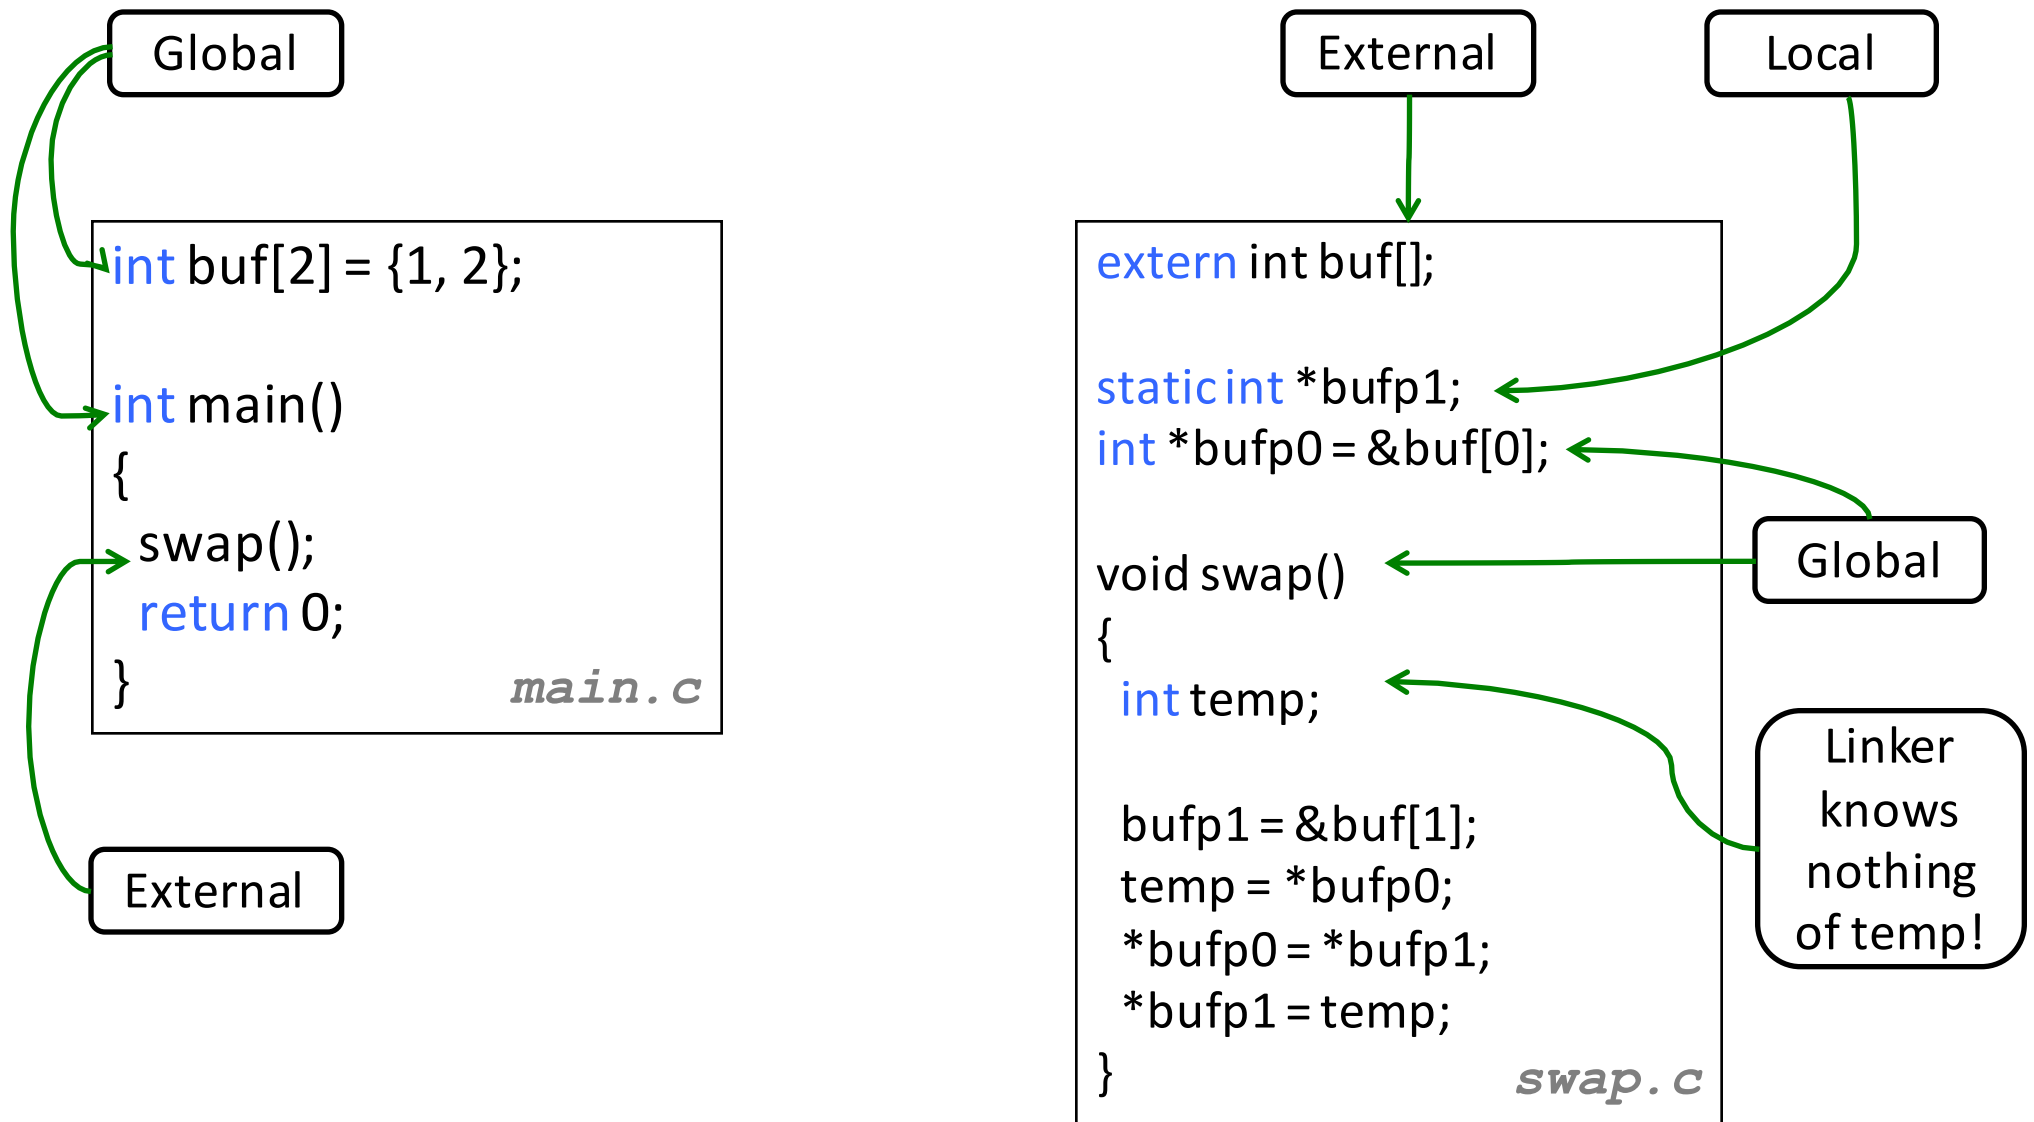
ELF header	0
Segment header table (required for executables)	
.text section	
.rodata section	
.data section	
.bss section	
.symtab section	
.rel.txt section	
.rel.data section	
.debug section	
Section header table	

Linker Symbols

For relocatable object module m:

- Global symbols
 - Symbols **defined by module m** that can be referenced by other modules.
 - E.g.: **non-static C functions** and **non-static global variables**.
- External symbols
 - Global symbols that are referenced by module m **but defined by some other module**.
- Local symbols
 - Symbols that are defined and referenced **exclusively** by module m.
 - E.g.: C functions and global variables defined with the **static** attribute.
 - **Local linker symbols are not the same as local program variables!**

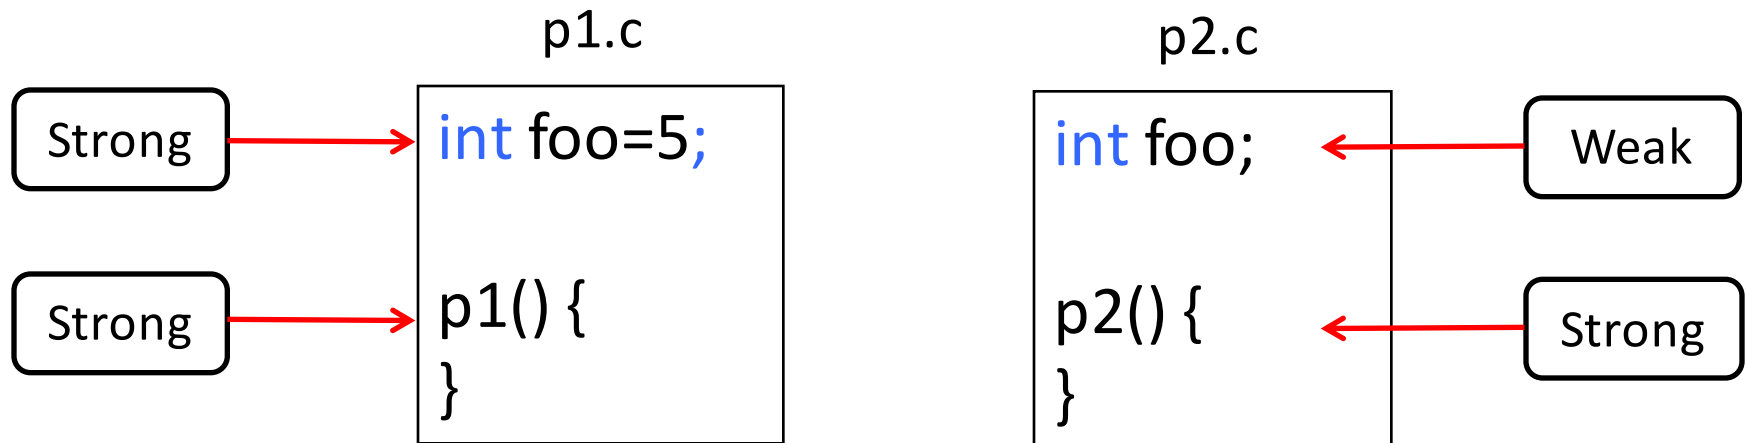
Symbol Resolution



Question: How do linkers resolve duplicate names?

Strong and Weak Symbols

- Program symbols are either strong or weak
 - **Strong**: procedures and initialized **global**
 - **Weak**: uninitialized **global**



Linker's Symbol Rules

- Rule 1: Multiple strong symbols are not allowed
 - Each item can be **defined only once**
 - Otherwise: Linker error
- Rule 2: Given a strong symbol and multiple weak symbols, **choose the strong symbol**
 - References to the weak symbols resolve to the strong symbol
- Rule 3: If there are multiple weak symbols, pick an **arbitrary weak symbol**
 - Can override this with `gcc -fno-common`

Linker Puzzles

```
int x;  
p1() {}
```

```
p1() {}
```

Link time error: **two strong symbols** (p1)

```
int x;  
p1() {}
```

```
int x;  
p2() {}
```

References to x will refer to the same **uninitialized int**.

```
int x;  
int y;  
p1() {}
```

```
double x;  
p2() {}
```

Writes to x in p2 **might overwrite y!**

```
int x;  
int y;  
p1() {}
```

```
double x=44;  
p2() {}
```

Writes to x in p2 will **overwrite y!**

```
int x=7;  
p1() {}
```

```
int x;  
p2() {}
```

References to x will refer to the same initialized variable.

two identical weak structs, compiled by different compilers with different alignment rules.

iClicker question

Consider alpha.c and beta.c: What will print?

```
int x;  
int main () {  
    x = 5;  
    beta();  
    printf("%d\n", x);  
}
```

```
static int x;  
void beta () {  
    x = 10;  
    return;  
}
```

- A) 5
- B) 10
- C) Something else

iClicker question

Consider alpha.c and beta.c: What will print?

```
int x;  
int main () {  
    x = 5;  
    beta();  
    printf("%d\n", x);  
}
```

```
static int x;  
void beta () {  
    x = 10;  
    return;  
}
```

A) 5

B) 10

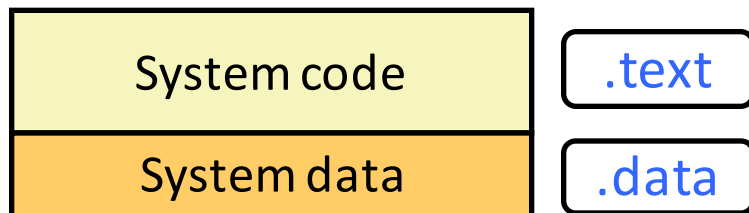
C) Something else

Global Variables

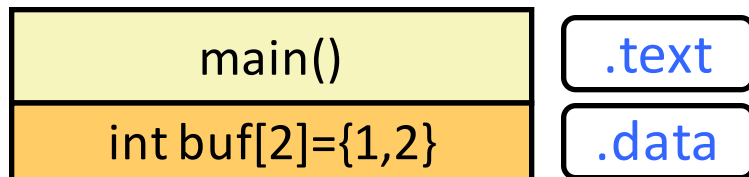
- Avoid them if possible
- Otherwise
 - Use `static` when possible
 - `Initialize` when a global variable is defined
 - Use `extern` for an external global variable

Step 2: Relocation

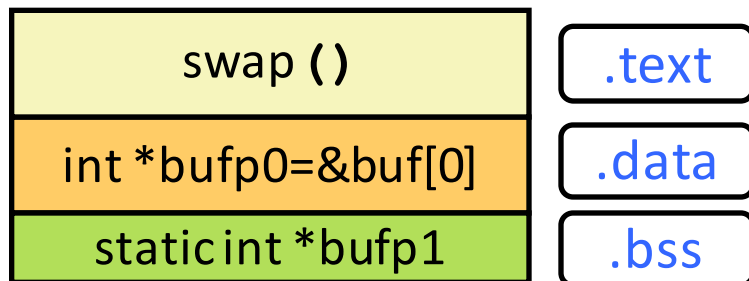
Relocatable Object Files



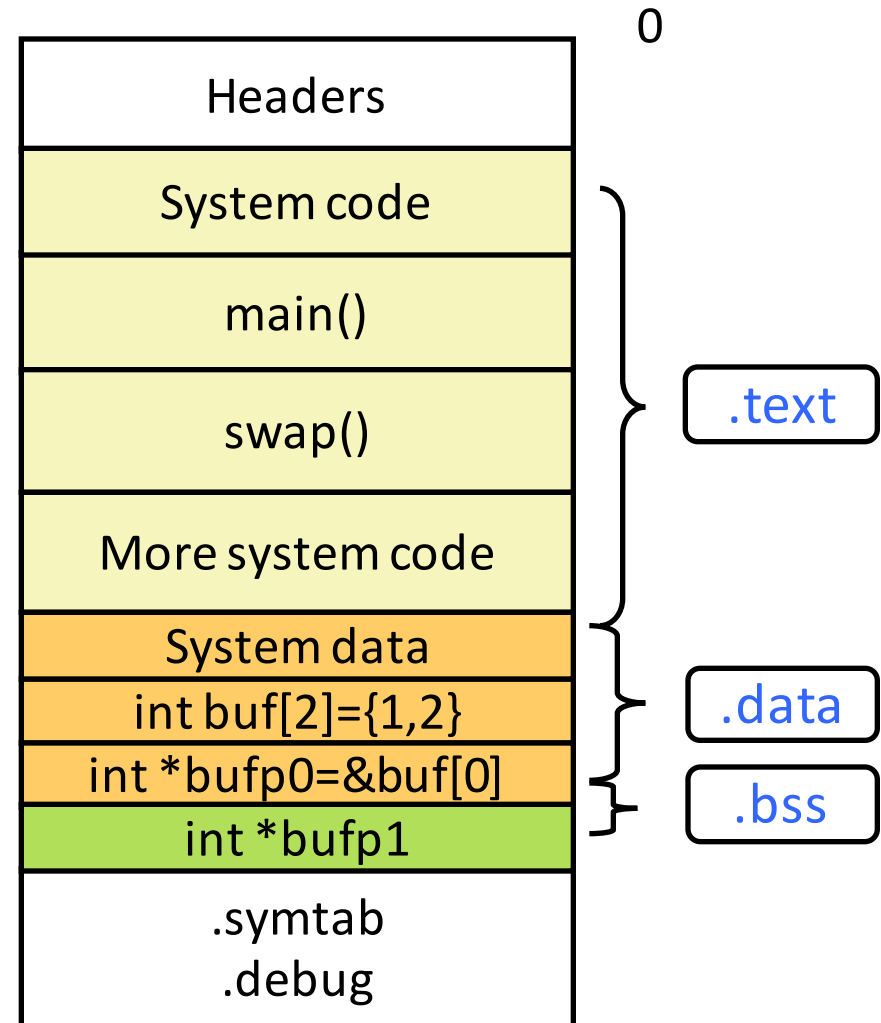
main.o



swap.o



Executable Object File



Even though private to swap,
requires allocation in `.bss`

Relocation Info (main)

main.c

```
int buf[2] =  
{1, 2};
```

```
int main()  
{  
    swap();  
    return 0;  
}
```

0:	8d 4c 24 04	lea	0x4(%esp),%ecx
4:	83 e4 f0	and	\$0xffffffff0,%esp
7:	ff 71 fc	pushl	0xffffffffc(%ecx)
a:	55	push	%ebp
b:	89 e5	mov	%esp,%ebp
d:	51	push	%ecx
e:	83 ec 04	sub	\$0x4,%esp
11:	e8 fc ff ff	call	12 <main+0x12>
		12: R_386_PC32	swap
16:	b8 00 00 00 00	mov	\$0x0,%eax
1b:	83 c4 04	add	\$0x4,%esp
1e:	59	pop	%ecx
1f:	5d	pop	%ebp
20:	8d 61 fc	lea	0xffffffffc(%ecx),%esp
23:	c3	ret	

main.o

-4

Source: `objdump-r-d main.o`

Disassembly of section `.data`: Source: `objdump-j .data -d main.o`

Executable Before/After Relocation (.text)

00000000 <main>:

```
...
e:      83 ec 04      sub    $0x4, %esp
11:     e8 fc ff ff ff  call   12 <main+0x12>
                        12: R_386_PC32 swap
16:     b8 00 00 00 00  mov    $0x0, %eax
...
```

Link time :
0x8048398
+ (-4)
- 0x8048386
= 0xe

08048374 <main>:

```
8048374:      8d 4c 24 04      lea    0x4(%esp),%ecx
8048378:      83 e4 f0          and    $0xffffffff0,%esp
804837b:      ff 71 fc          pushl  0xffffffffc(%ecx)
804837e:      55                push   %ebp
804837f:      89 e5              mov    %esp,%ebp
8048381:      51                push   %ecx
8048382:      83 ec 04          sub    $0x4,%esp
8048385:      e8 0e 00 00 00    call   8048398 <swap>
804838a:      b8 00 00 00 00    mov    $0x0,%eax
804838f:      83 c4 04          add    $0x4,%esp
8048392:      59                pop    %ecx
8048393:      5d                pop    %ebp
8048394:      8d 61 fc          lea    0xffffffffc(%ecx),%esp
8048397:      c3                ret
```

Runtime:
0x804838a +
0xe
=
0x8048398

Relocation Info (swap, .text)

swap.c

```
extern int buf[];

static int *bufp1;
int *bufp0 = &buf[0];

void swap()
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

swap.o

```
00000000 <swap>:
0: 55                                push    %ebp
1: 89 e5                            mov     %esp,%ebp
3: 53                                push    %ebx
4: c7 05 00 00 00 00 04 movl     $0x4,0x0
b: 00 00 00

                                6: R_386_32 .bss
                                a: R_386_32 buf
e: 8b 0d 00 00 00 00 mov     0x0,%ecx
                                10: R_386_32 bufp0
14: 8b 19                            mov     (%ecx),%ebx
16: ba 04 00 00 00 mov     $0x4,%edx
                                17: R_386_32 buf
1b: 8b 02                            mov     (%edx),%eax
1d: 89 01                            mov     %eax,(%ecx)
1f: 89 1a                            mov     %ebx,(%edx)
21: 5b                                pop     %ebx
22: 5d                                pop     %ebp
23: c3                                ret
```

Relocation Info (swap, .data)

swap.c

```
extern int buf[];

static int *bufp1;
int *bufp0 = &buf[0];

void swap()
{
    int temp;

    bufp1 = &buf[1];
    temp = *bufp0;
    *bufp0 = *bufp1;
    *bufp1 = temp;
}
```

Disassembly of section .data:

```
00000000 <bufp0>:
    0:  00 00 00 00
               0: R_386_32 buf
```

00000000 <swap>:

...

4: c7 05 00 00 00 00 04 movl \$0x4,0x0

b: 00 00 00

6: R_386_32 .bss

a: R_386_32 buf

e: 8b 0d 00 00 00 00 mov 0x0,%ecx

10: R_386_32 bufp0

14: 8b 19 mov (%ecx),%ebx

16: ba 04 00 00 00 mov \$0x4,%edx

17: R_386_32 buf

Before relocation

After relocation

08048398 <swap>:

8048398: 55 push %ebp

8048399: 89 e5 mov %esp,%ebp

804839b: 53 push %ebx

804839c: c7 05 14 96 04 08 04 movl \$0x8049604,0x8049614

80483a3: 96 04 08

80483a6: 8b 0d 08 96 04 08 mov 0x8049608,%ecx

80483ac: 8b 19 mov (%ecx),%ebx

80483ae: ba 04 96 04 08 mov \$0x8049604,%edx

80483b3: 8b 02 mov (%edx),%eax

80483b5: 89 01 mov %eax,(%ecx)

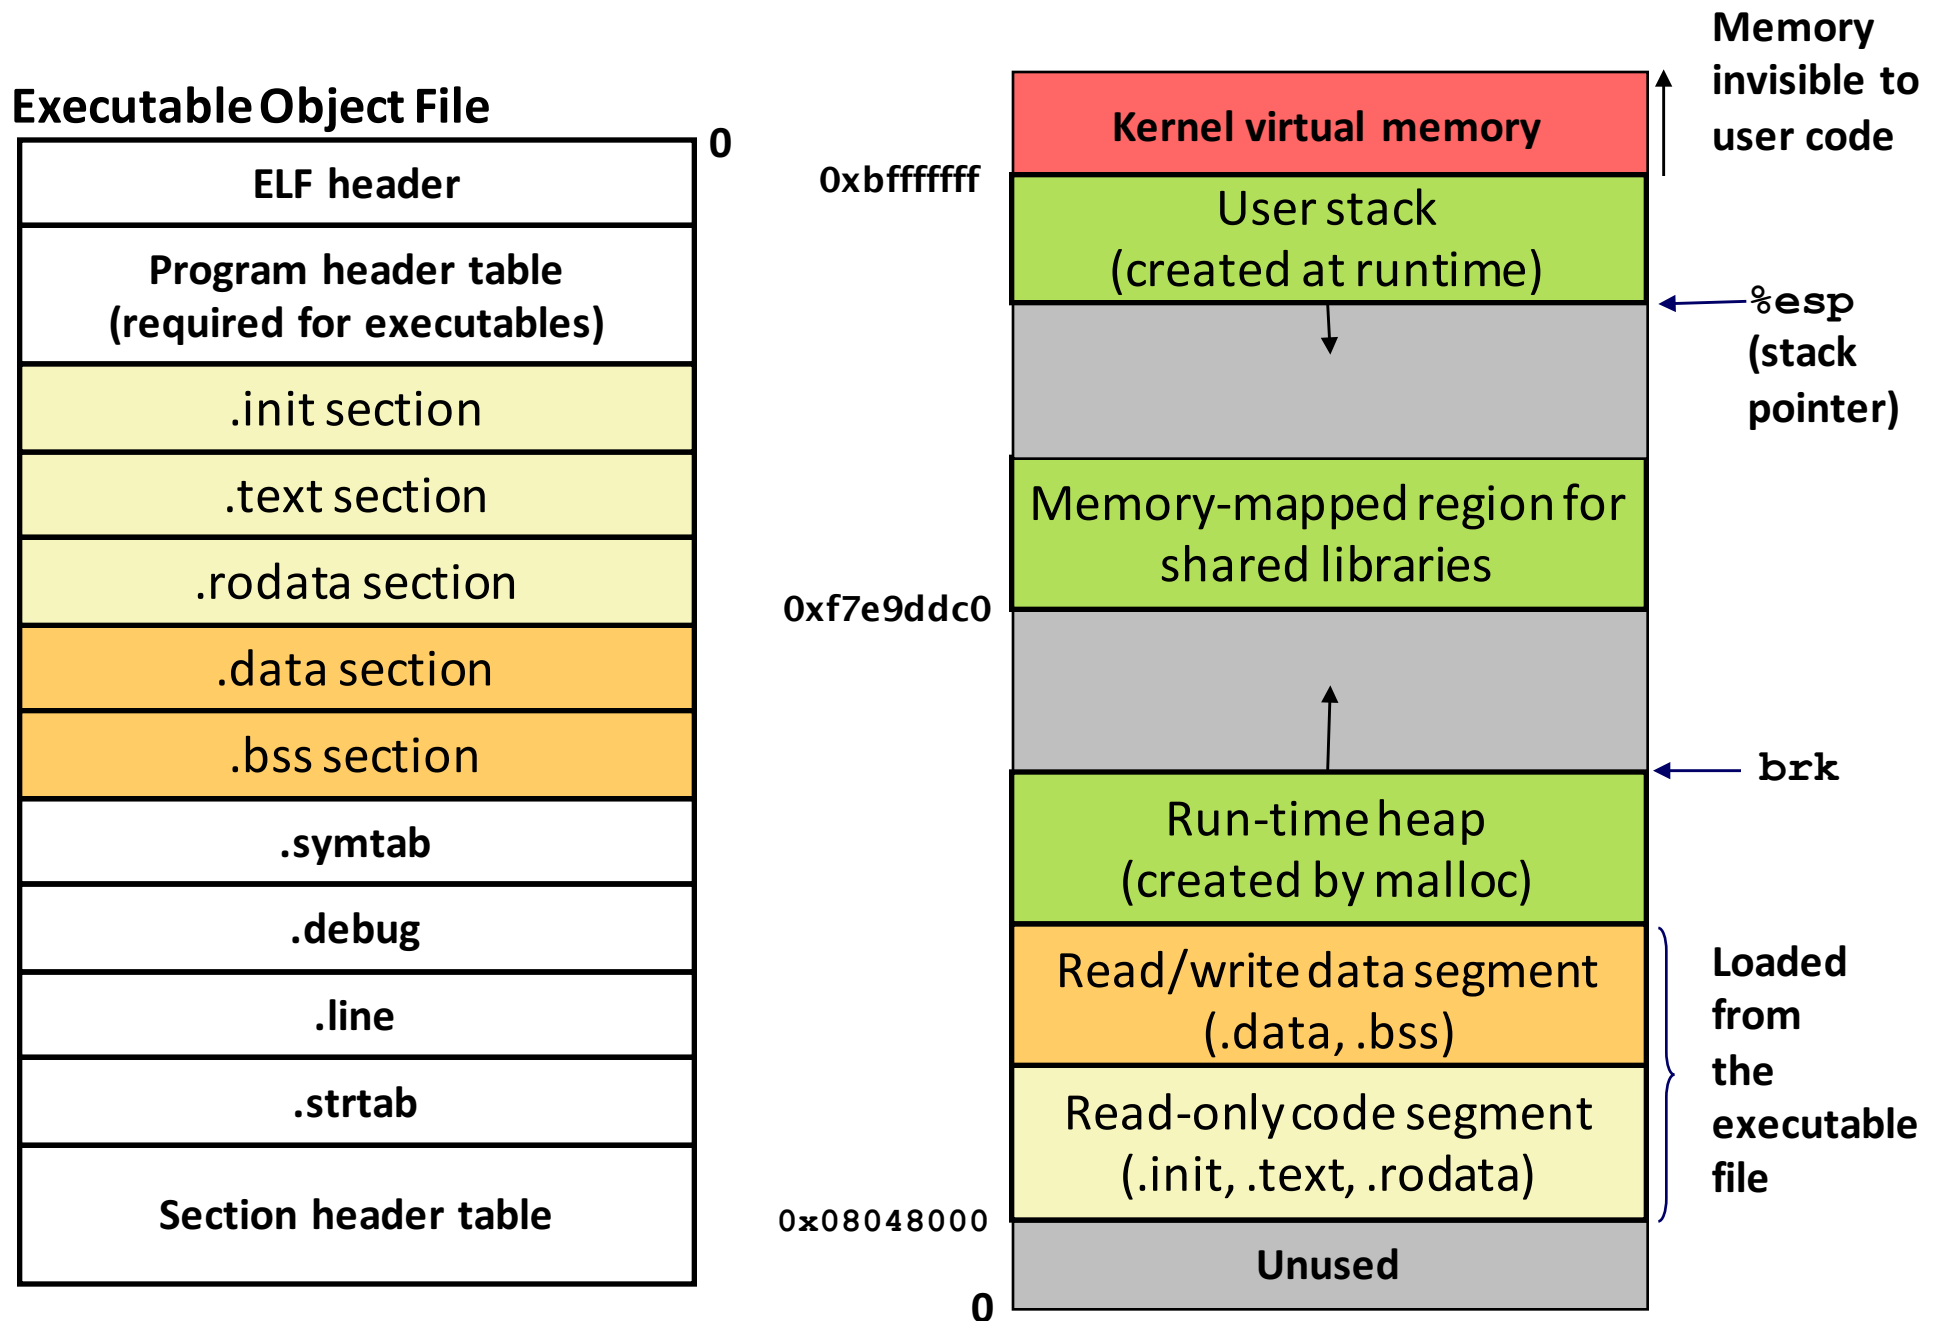
80483b7: 89 1a mov %ebx,(%edx)

80483b9: 5b pop %ebx

80483ba: 5d pop %ebp

80483bb: c3 ret

Loading Executable Object Files



Packaging Commonly Used Functions

- How to package functions commonly used by programmers?
 - Math, I/O, memory management, string manipulation, etc.
- Awkward, given the linker framework so far:
 - Option 1: Put each function in a separate source file
 - Programmers explicitly link appropriate binaries into their programs
 - More efficient, but burdensome on the programmer
 - Option 2: Put all functions into a single source file
 - Programmers link big object file into their programs
 - Space and time inefficient

Commonly Used Libraries

libc.a (the C standard library)

- 8 MB archive of 1392 object files.
- I/O, memory allocation, signal handling, string handling, data and time, random numbers, integer math

libm.a (the C math library)

- 1 MB archive of 401 object files.
- floating point math (sin, cos, tan, log, exp, sqrt, ...)

```
% ar -t /usr/lib/libc.a | sort
```

```
...  
fork.o  
...  
fprintf.o  
fpu_control.o  
fputc.o  
freopen.o  
fscanf.o  
fseek.o  
fstab.o  
...
```

```
% ar -t /usr/lib/libm.a | sort
```

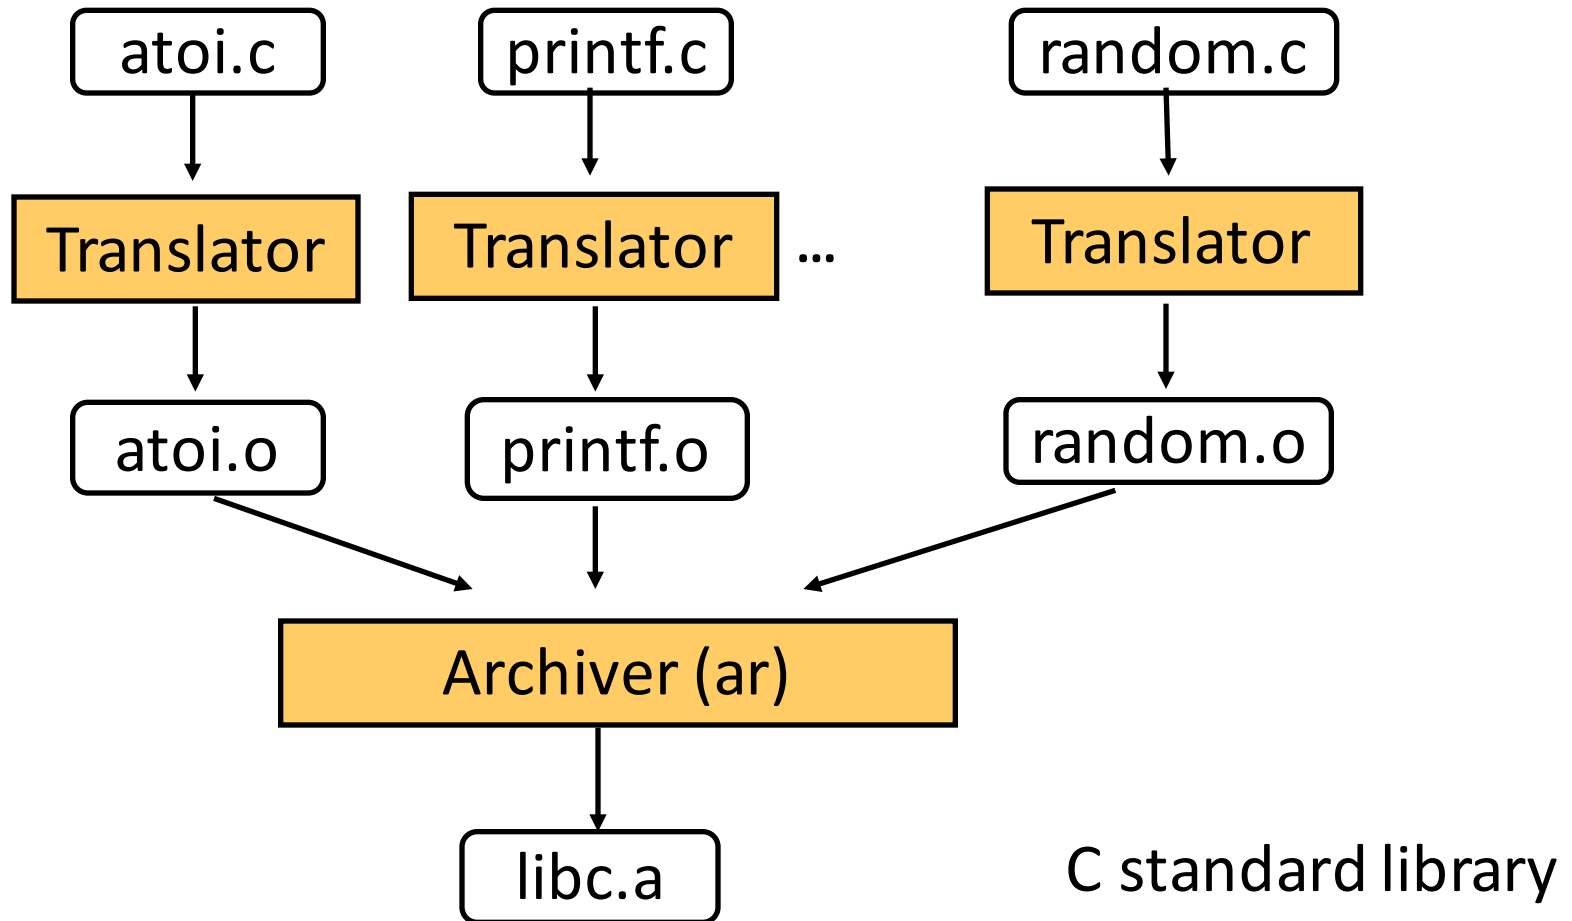
```
...  
e_acos.o  
e_acosf.o  
e_acosh.o  
e_acoshf.o  
e_acoshl.o  
e_acosl.o  
e_asin.o  
e_asinf.o  
e_asinl.o  
...
```

Static Libraries

Static libraries (.a archive files)

- Concatenate related **relocatable object files** into a single file with an index (called **an archive**).
- Enhance linker so that it tries to resolve **unresolved external references** by looking for the symbols in one or more archives.
- If an archive member file resolves reference, **link it into the executable**.

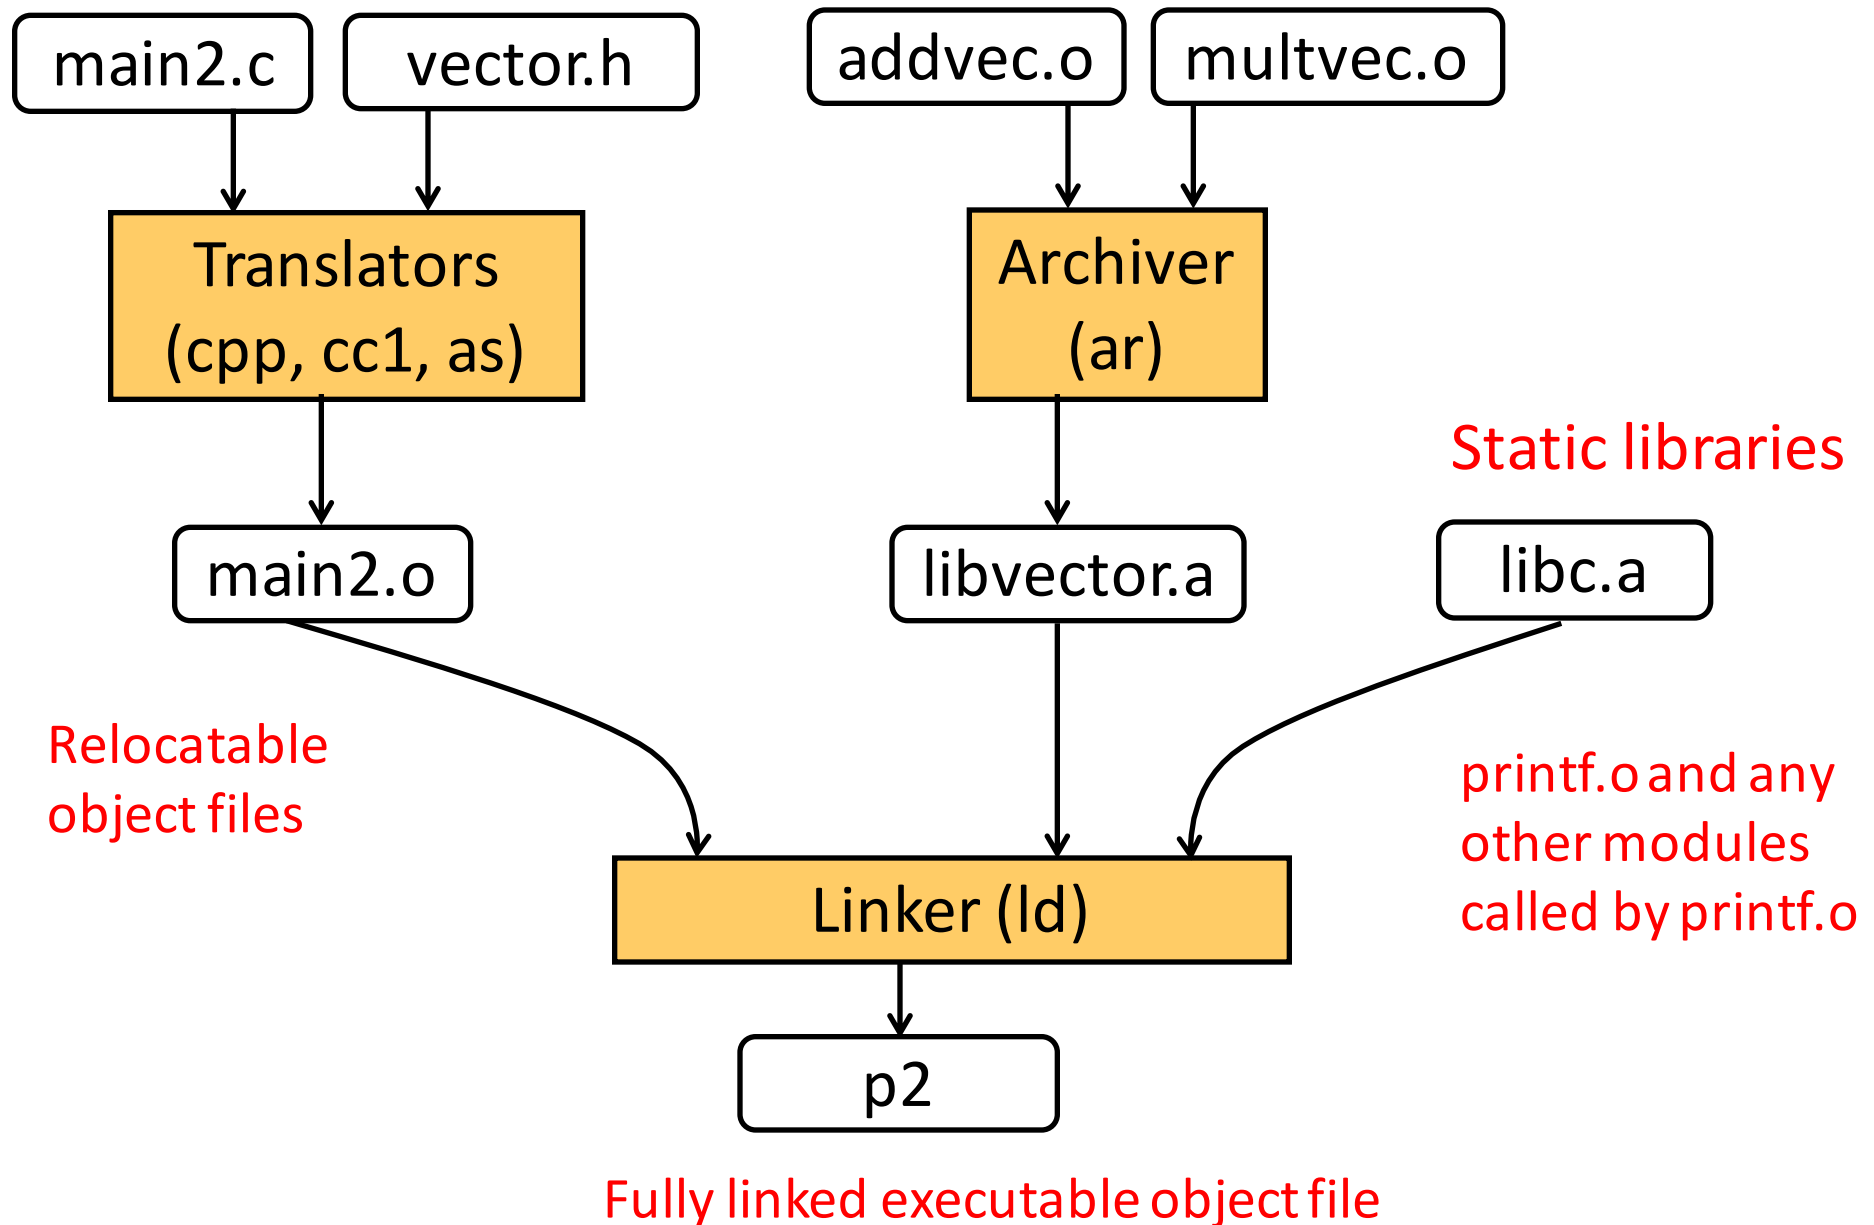
Creating Static Libraries



> `ar rcs libc.a atoi.o printf.o ...random.o`

- Archiver allows incremental updates
- Recompile function that changes and replace .o file in archive.

Linking with Static Libraries



Using Static Libraries

- Linker's algorithm for resolving external references:
 - Scan .o files and .a files in the **command line order**.
 - During the scan, keep a list of **the current unresolved** references.
 - As each new .o or .a file, obj, is encountered, try to resolve each unresolved reference in the list against the symbols defined in obj.
 - If any entries in the unresolved list at **end of scan, then error**.
- Problem:
 - **Command line order matters!**
 - Moral: put libraries at the end of the command line.

```
gcc -L. libtest.o -lmine  
gcc -L. -lmine libtest.o  
libtest.o: In function `main':  
libtest.o(.text+0x4): undefined reference to `libfun'
```

Shared Libraries

- **Static libraries** have the following disadvantages:
 - **Duplication** in the stored executables (every function needs libc)
 - Duplication in the running executables
 - Minor bug fixes of system libraries require each application to **explicitly relink**
- Modern solution: **Shared Libraries**
 - Object files that contain code and data that are loaded and linked into an application **dynamically, at either load-time or run-time**
 - Also called: dynamic link libraries, DLLs, .so files

Shared Libraries

- **Dynamic linking** can occur when executable is first loaded and run (**load-time linking**).
 - Common case for Linux, handled automatically by the dynamic linker (ld-linux.so).
 - Standard C library (libc.so) usually dynamically linked.
- Dynamic linking can also occur after program has begun (**run-time linking**).
 - In Linux, this is done by calls to the **dlopen()** interface .
 - Distributing software.
 - High-performance web servers.
 - Runtime library inter-positioning.
- Shared library routines can be shared by multiple processes.
 - More on this when we learn about virtual memory

Dynamic Linking at Load-time

main2.c vector.h

> gcc -shared -o libvector.so \
addvec.c multvec.c

Translators
(cpp, cc1, as)

libc.so
libvector.so

Relocatable
object file

main2.o

Relocation and
symbol table
info

Linker (ld)

Partially linked
executable
object file

p2

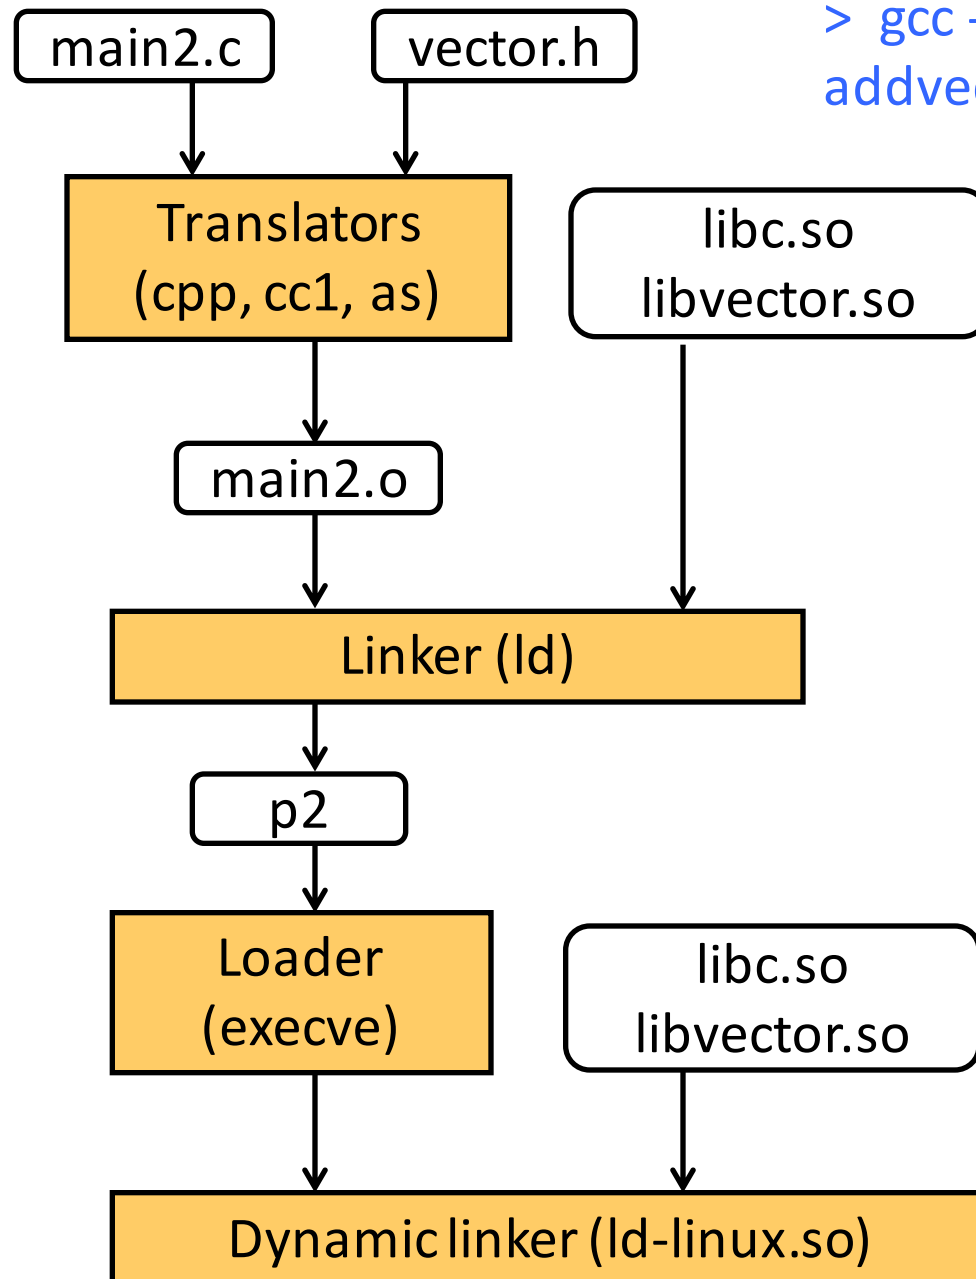
Loader
(execve)

libc.so
libvector.so

Code and data

Fully linked
executable
in memory

Dynamic linker (ld-linux.so)



Dynamic Linking at Run-time

```
#include <stdio.h>
#include <dlfcn.h>
int x[2] = {1, 2};
int y[2] = {3, 4};
int z[2];

int main() {
    void *handle;
    void (*addvec)(int *, int *, int *, int);
    char *error;

    // Dynamically load the shared lib that contains addvec()
    handle = dlopen("./libvector.so", RTLD_LAZY);
    if (!handle) {
        fprintf(stderr, "%s\n", dlerror());
        exit(1);
    }
```

Dynamic Linking at Run-time

```
...
// Get a pointer to the addvec() function we just loaded
addvec = dlsym(handle, "addvec");
if ((error = dlerror()) != NULL) {
    fprintf(stderr, "%s\n", error);
    exit(1);
}
// Now we can call addvec() just like any other function
addvec(x, y, z, 2);
printf("z = [%d %d]\n", z[0], z[1]);
// unload the shared library
if (dlclose(handle) < 0) {
    fprintf(stderr, "%s\n", dlerror());
    exit(1);
}
return 0;
}
```

Summary

- Linking is a technique that allows programs to be constructed from **multiple object files**.
- Linking can happen at different times in a program's lifetime:
 - **Compile time** (when a program is compiled)
 - **Load time** (when a program is loaded into memory)
 - **Run time** (while a program is executing)
- Understanding linking can help you avoid nasty errors and make you a better programmer.

Adapted from slides by Bryant and O'Hallaron