

# Computer System Principles

Network Programming




# Announcement


- We will drop the lowest assignment grade, except that we won't drop the grade for the last assignment if that grade is not a passing grade (70%).

# **NETWORK TOPOLOGY**

# Network Topology



# Client-server architecture



## server:

- Generally an “always-on” host
- (e.g. Amazon, telephone)
- Generally at a *permanent* IP address
  - data centers for scaling

## clients:

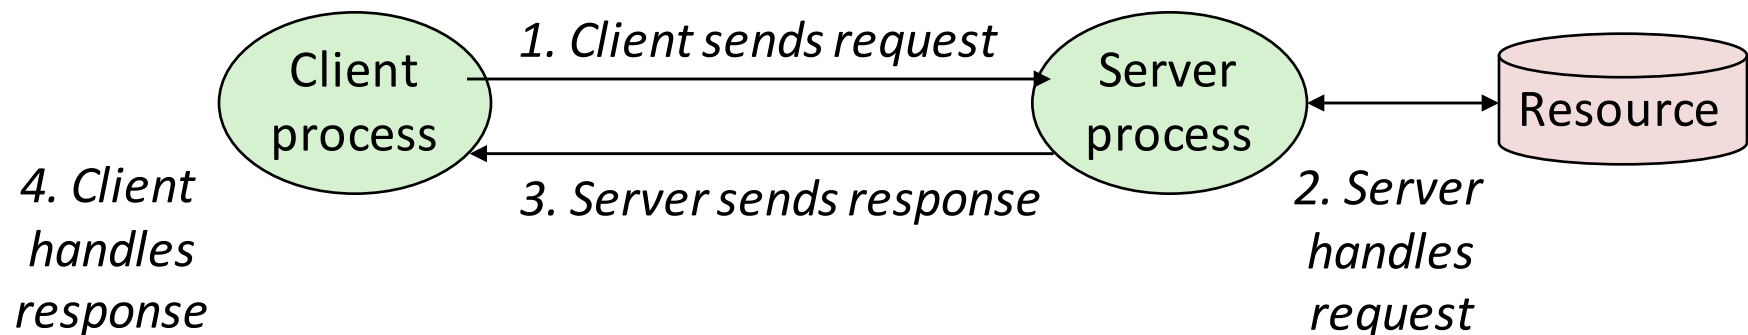
- communicate with server
- may be intermittently connected
- may have dynamic IP addresses
- do not typically communicate directly with each other

# Client-Server Model

- An application process is assigned a process identifier number (process ID) - likely to be different each time that process is started.
- Process IDs differ between OS platforms - **they are not uniform.**
- A server process can have **multiple connections to multiple clients** at a time - simple connection identifiers at one end are not unique.

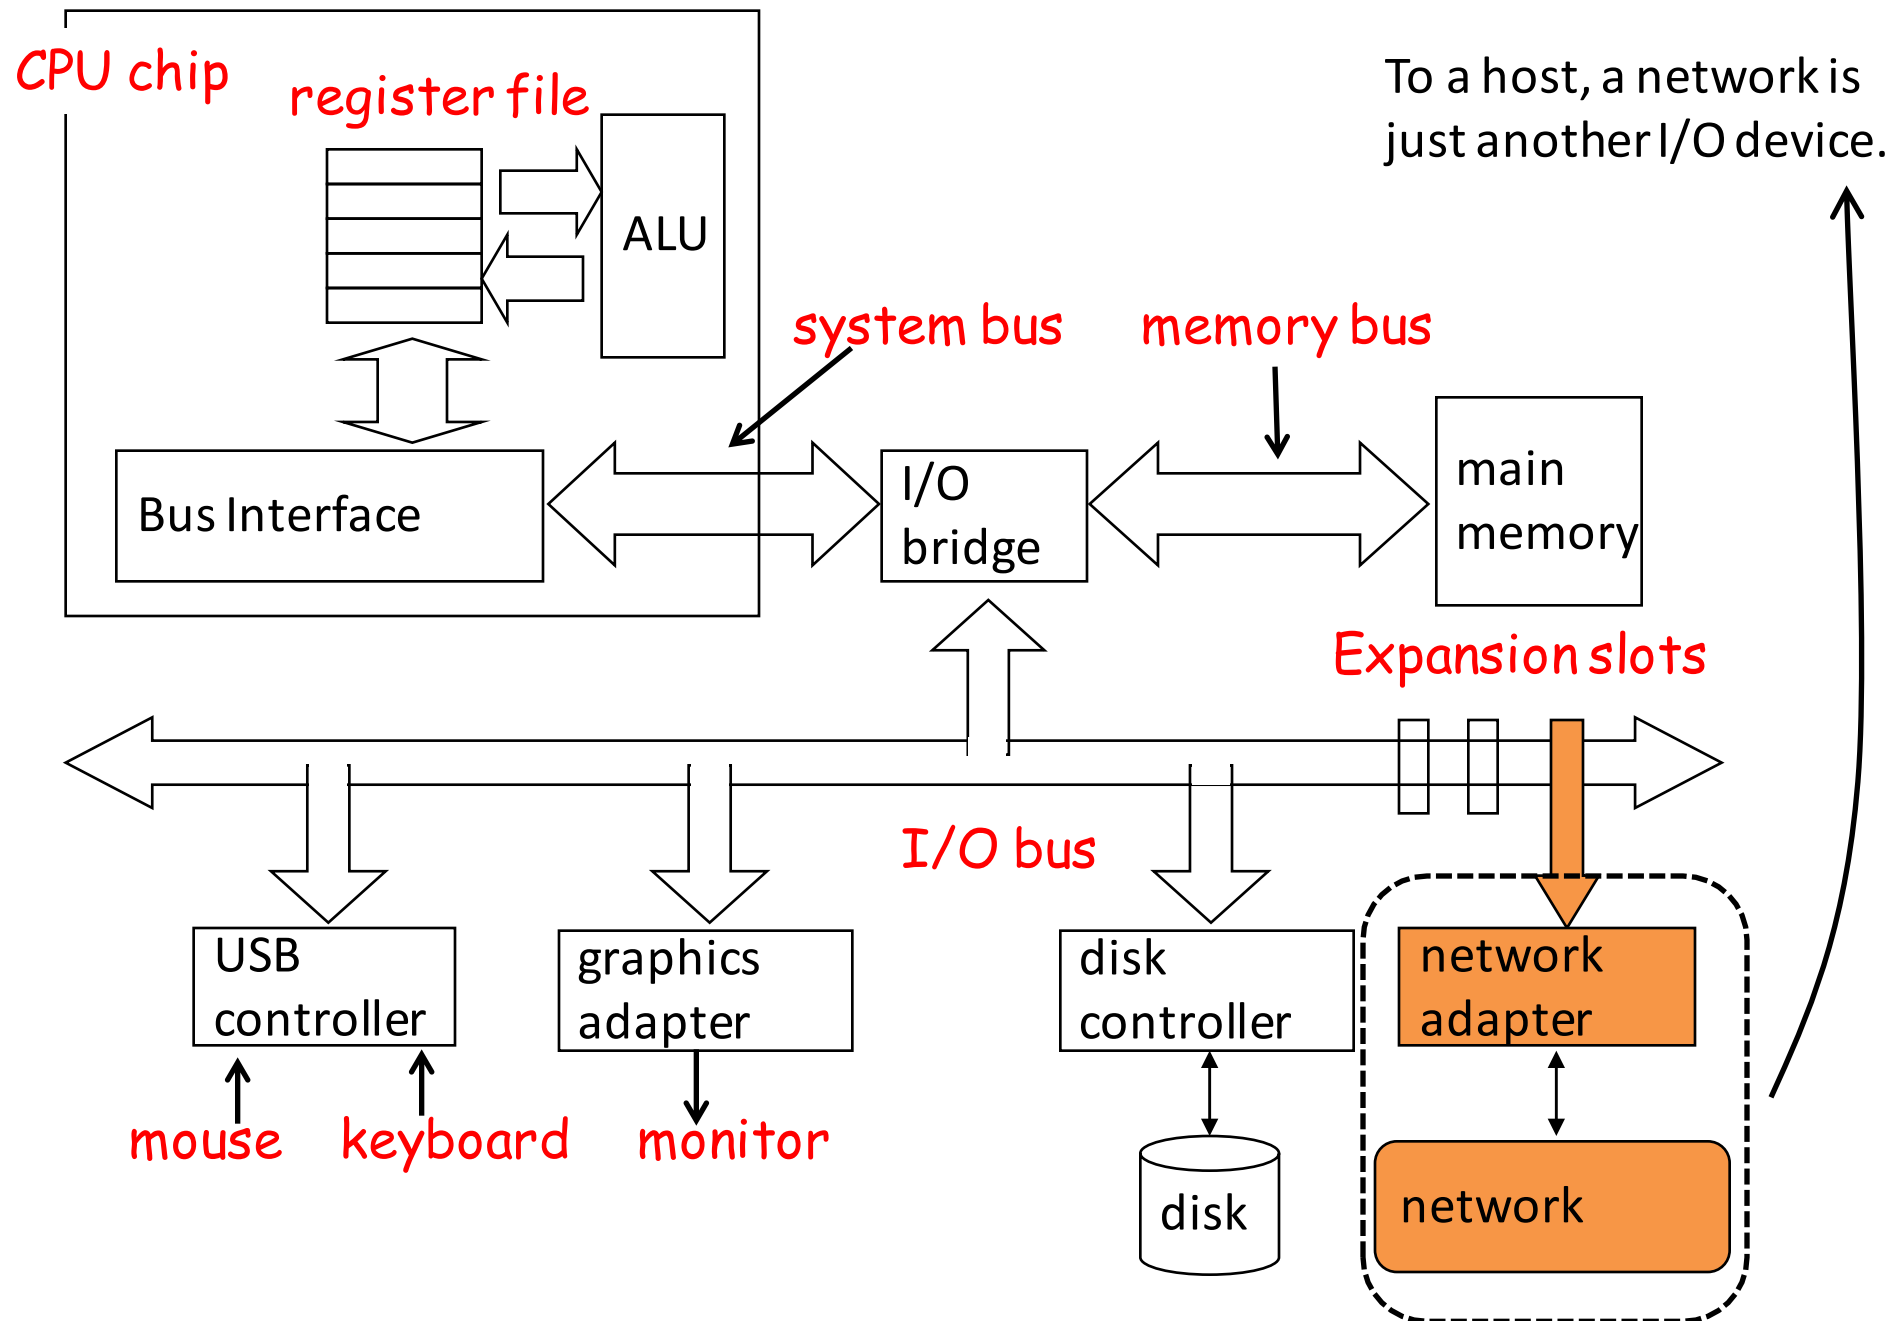
# A Client-Server Transaction

- Most network applications are based on the client-server model:
  - A **server** process and one or more **client** processes
  - Server manages some **resource**
  - Server provides **service** by manipulating resource for clients
  - Server activated by request from client (vending machine analogy)



*Note: clients and servers are processes running on hosts  
(can be the same or different hosts)*

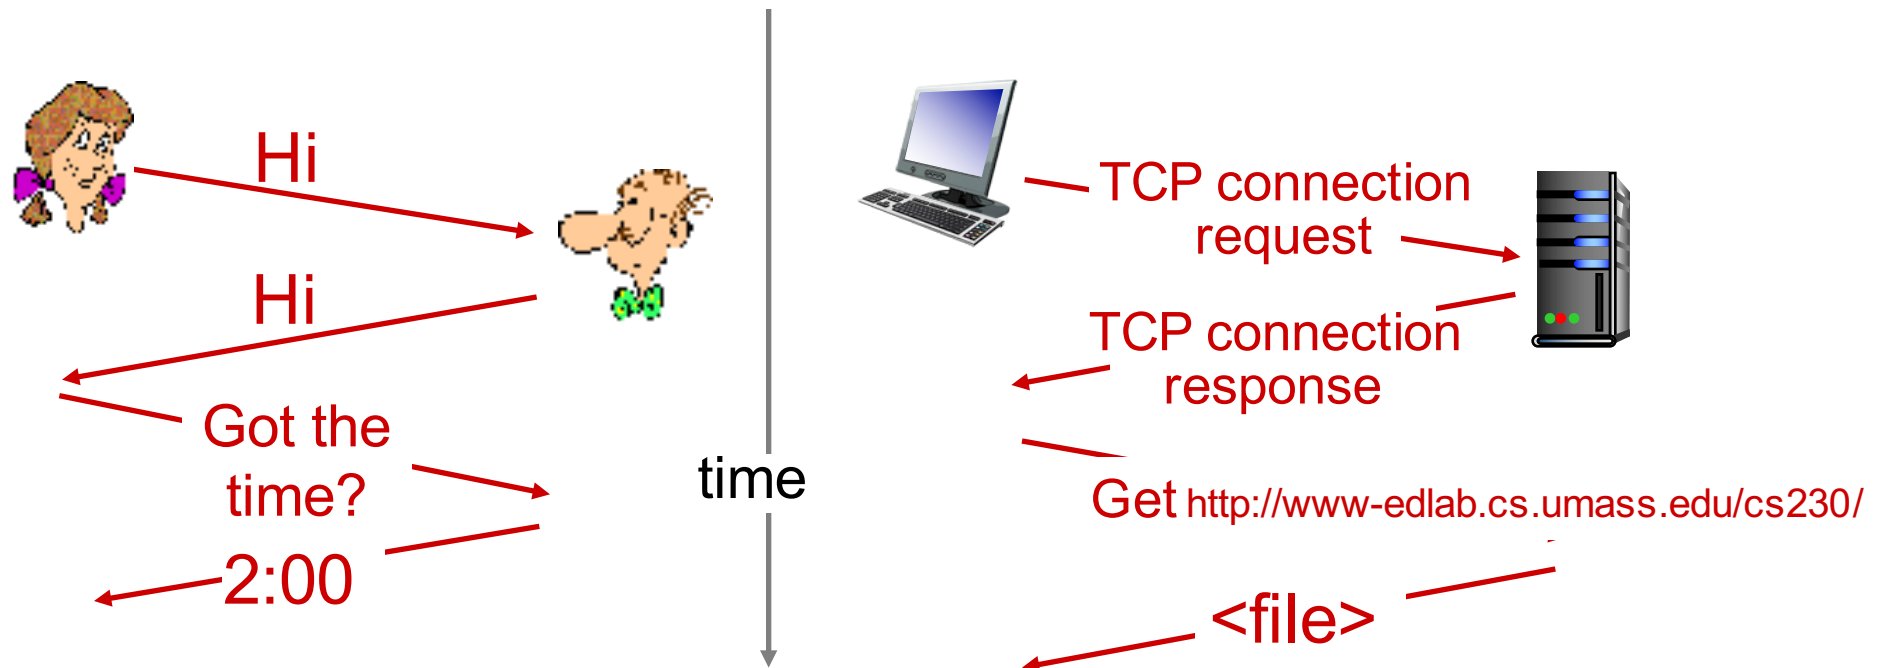
# Network Hardware






# Protocols

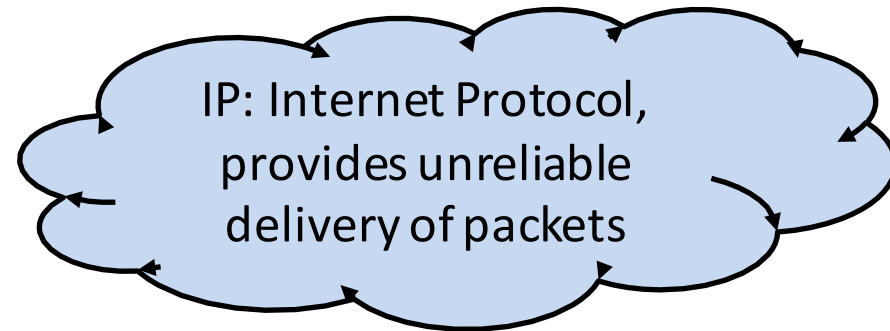
a human protocol and a computer network protocol:




*protocols* define *format*, *order* of *msgs sent and received* among network entities, and *actions taken* on msg transmission, receipt


Data Link  
Layer





Network  
Layer: IP  
Data Link  
Layer

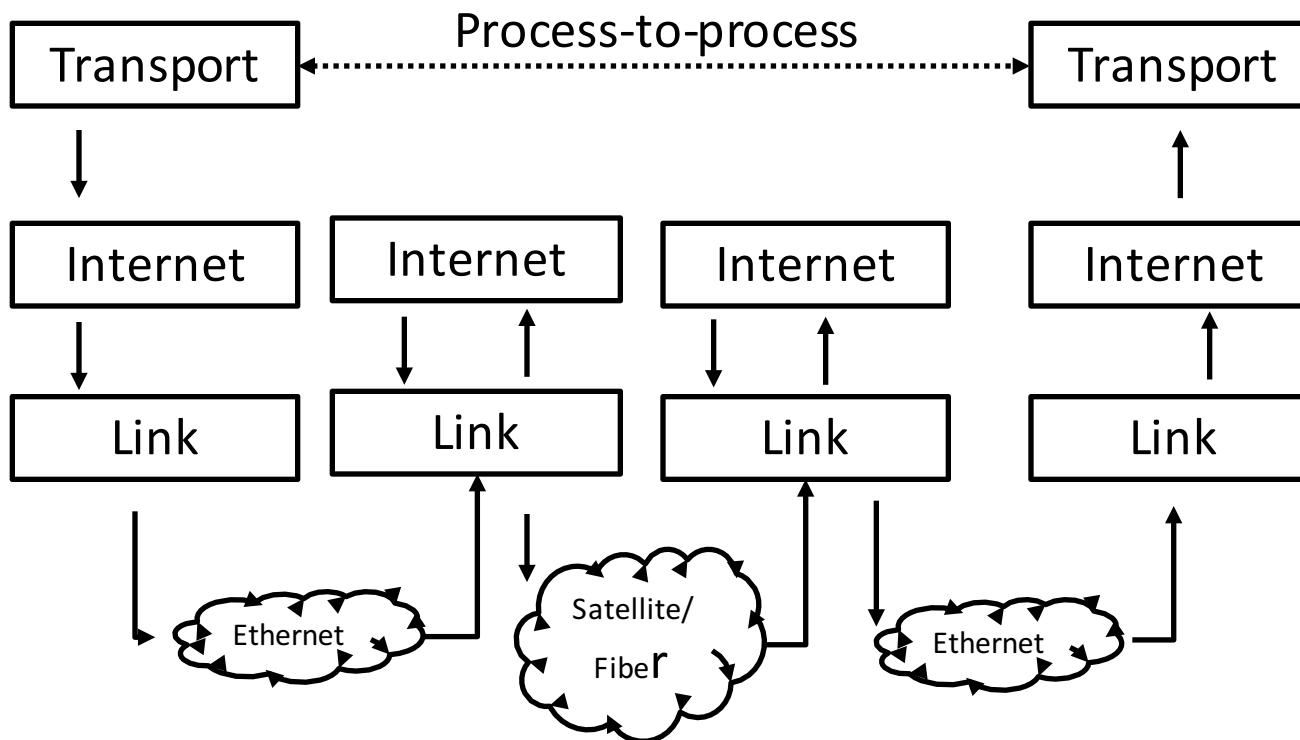


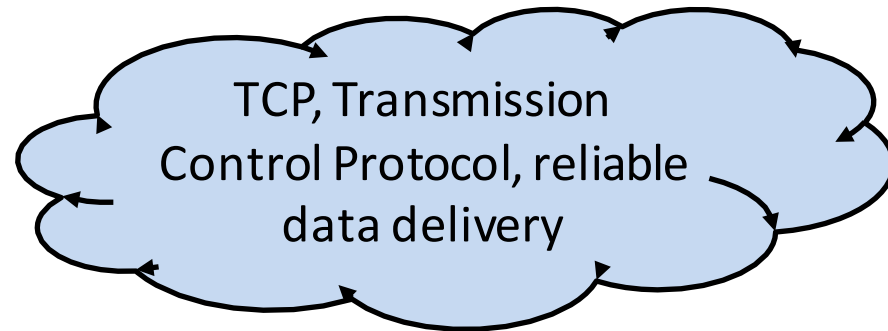


Transport Layer:  
TCP/UDP

Network Layer: IP

Data Link Layer

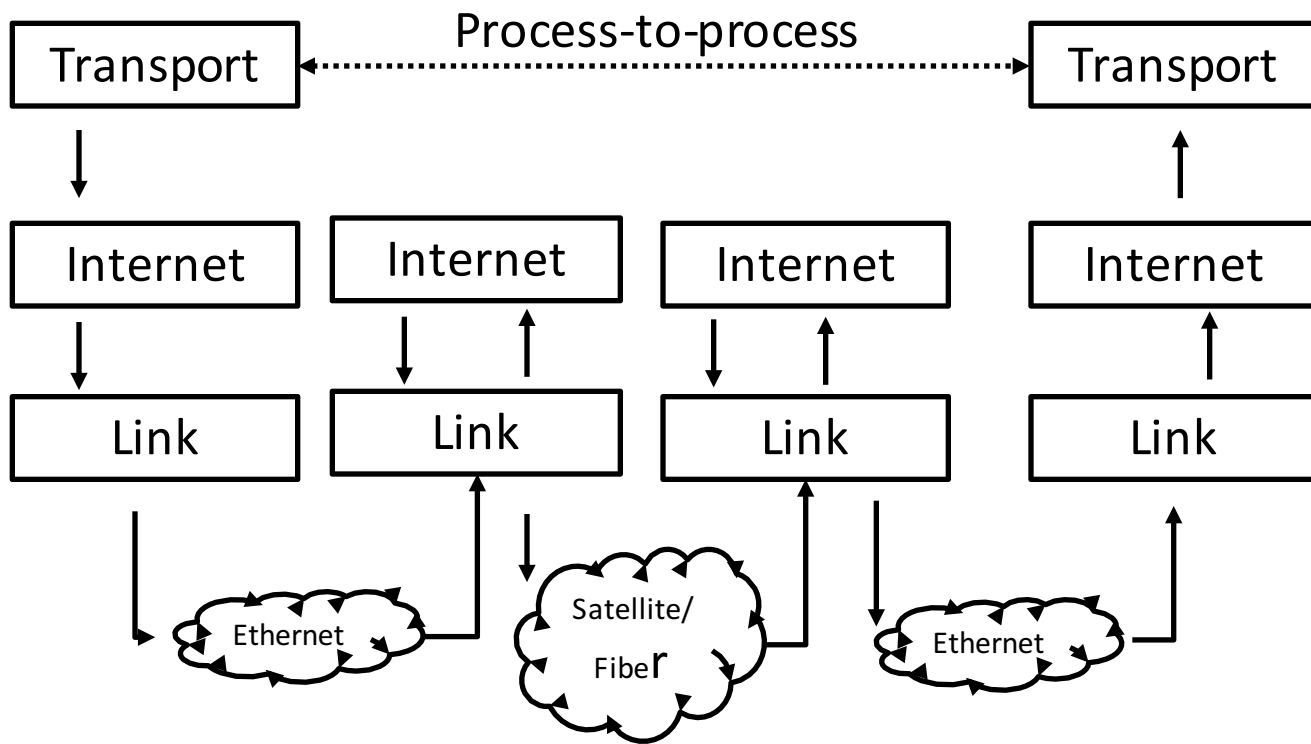




Transport Layer:  
TCP/UDP

Network Layer: IP

Data Link Layer

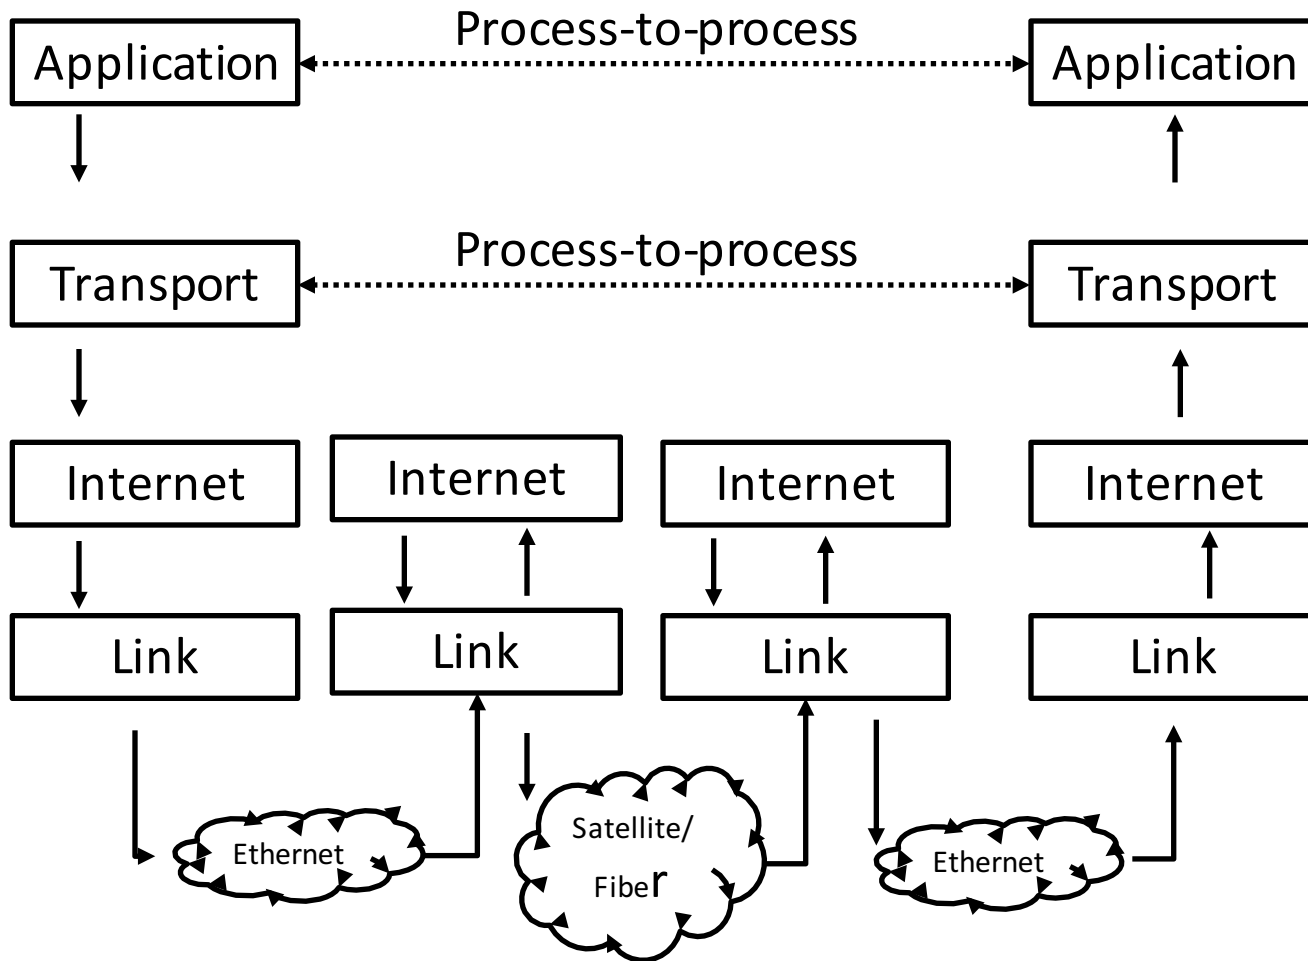


Application Layer:  
HTTP/FTP  
...

Transport Layer:  
TCP/UDP

Network Layer:  
IP

Data Link Layer

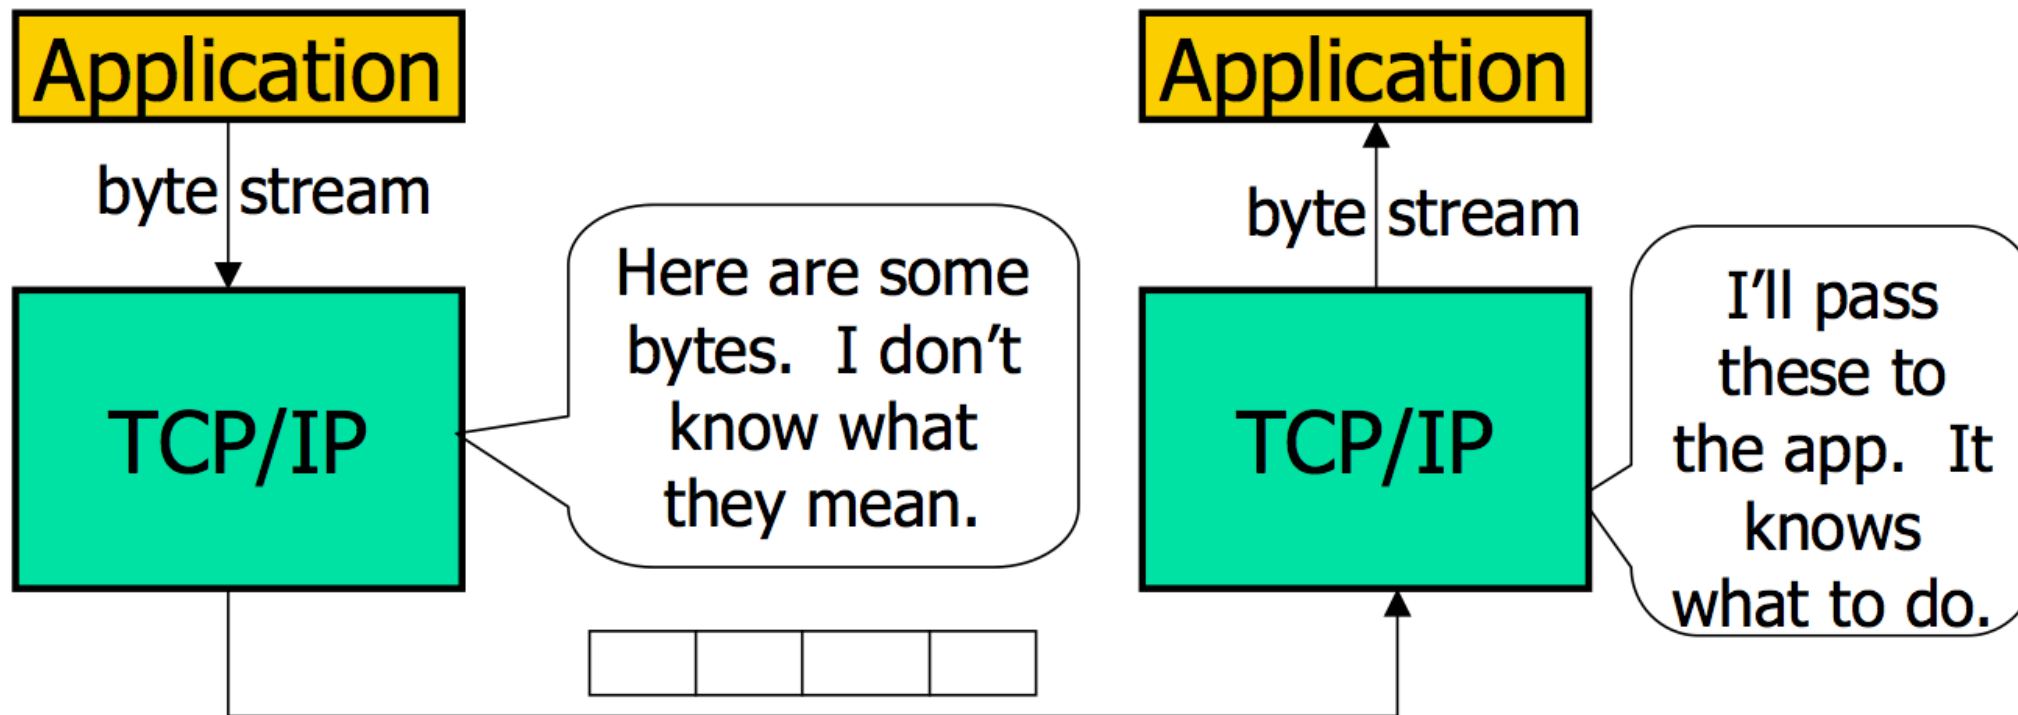


# **TCP OVERVIEW**

# TCP

The TCP data unit- segment

- does not recognize messages,
- sends a block of bytes from the byte stream between sender and receiver.





# TCP

- The primary purpose of TCP is to provide a reliable logical circuit or connection service between pairs of processes.
  - does not assume reliability from the lower-level protocols (such as IP), so TCP must guarantee this itself
    - Heavy network utilization and resulting congestion
    - Faulty network hardware or connectors


# TCP

- The primary purpose of TCP is to provide a reliable logical circuit or connection service between pairs of processes.
  - does not assume reliability from the lower-level protocols (such as IP), so TCP must guarantee this itself
    - Heavy network utilization and resulting congestion
    - Faulty network hardware or connectors
  - When a packet is successfully delivered to its destination, the destination should send an acknowledgement


# TCP

- The primary purpose of TCP is to provide a reliable logical circuit or connection service between pairs of processes.
  - does not assume reliability from the lower-level protocols (such as IP), so TCP must guarantee this itself
    - Heavy network utilization and resulting congestion
    - Faulty network hardware or connectors
  - When a packet is successfully delivered to its destination, the destination should send an acknowledgement
  - otherwise, the source system retransmits the packet
    - the destination never receive the packet
    - the ack is lost


# TCP Three-Way Handshake




# TCP Three-Way Handshake



# TCP Three-Way Handshake



# TCP Three-Way Handshake



**SOCKET**



# What is a socket?

- Clients and servers communicate by sending streams of bytes over **connections**:
  - Point-to-point, full-duplex (2-way communication), and reliable.

# What is a socket?

- Clients and servers communicate by sending streams of bytes over **connections**:
  - Point-to-point, full-duplex (2-way communication), and reliable.
- A **socket** is an endpoint of a connection
  - Socket address is an IPaddress:port pair

# What is a socket?

- Clients and servers communicate by sending streams of bytes over **connections**:
  - Point-to-point, full-duplex (2-way communication), and reliable.
- A **socket** is an endpoint of a connection
  - Socket address is an IP address:port pair
- A **port** is a 16-bit integer that identifies a process

# What is a socket?

- Clients and servers communicate by sending streams of bytes over **connections**:
  - Point-to-point, full-duplex (2-way communication), and reliable.
- A **socket** is an endpoint of a connection
  - Socket address is an IPaddress:port pair
- A **port** is a 16-bit integer that identifies a process:
  - **Ephemeral port**: Assigned automatically on client when client makes a connection request
  - **Well-known port**: Associated with some service provided by a server (e.g., port 80 is associated with Web servers)

# What is a socket?

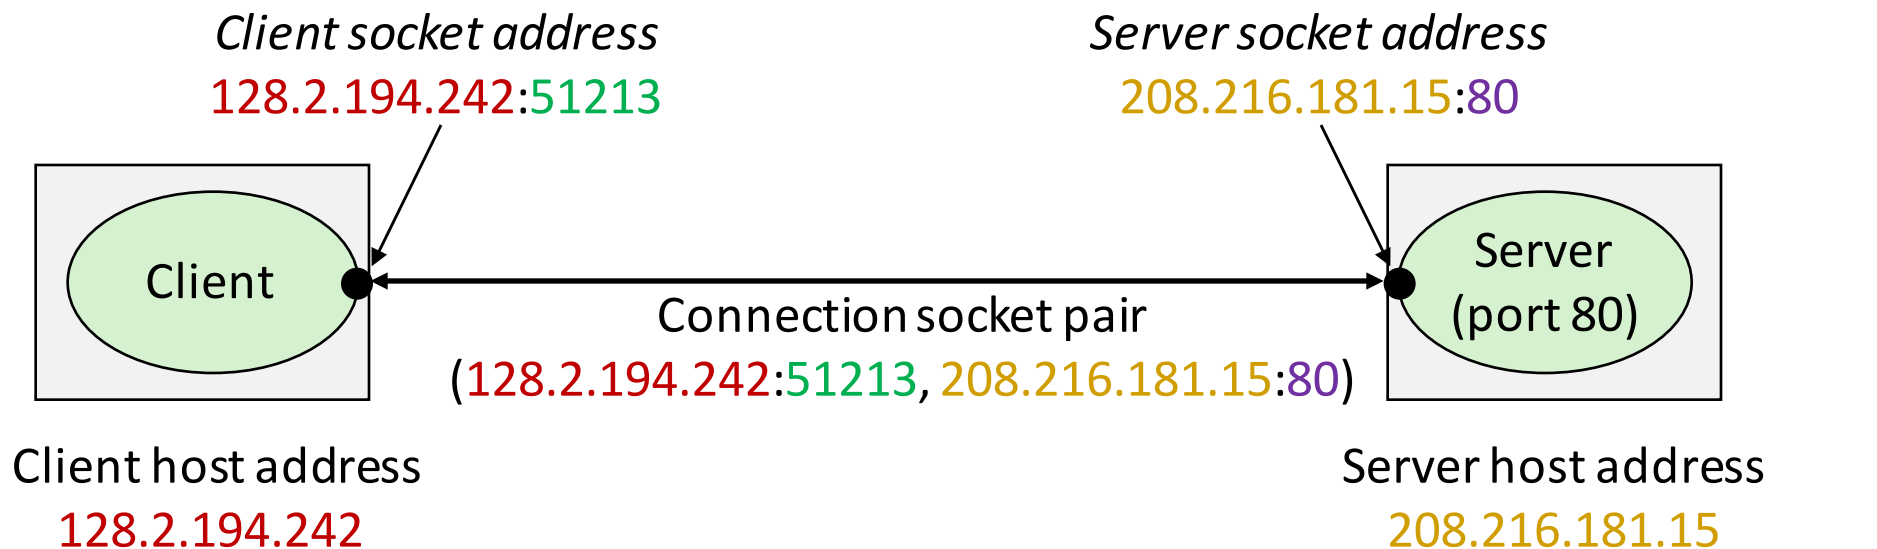
- Clients and servers communicate by sending streams of bytes over **connections**:
  - Point-to-point, full-duplex (2-way communication), and reliable.
- A **socket** is an endpoint of a connection
  - Socket address is an IPaddress:port pair
- A **port** is a 16-bit integer that identifies a process:
  - **Ephemeral port**: Assigned automatically on client when client makes a connection request
  - **Well-known port**: Associated with some service provided by a server (e.g., port 80 is associated with Web servers)
- A connection is uniquely identified by the socket addresses of its endpoints (**socket pair**) (and its protocol, e.g., TCP).
  - (cliaddr:cliport, servaddr:servport)

# TCP Sockets -Ports

- TCP uses **well-known** and **ephemeral ports**
  - Well Known Ports: range 0-1023
    - assigned to the server side of an application.
  - Registered Ports: range 1024-49151
    - publicly defined; convenience for the Internet community; avoids vendor conflicts
  - Dynamic and/or Private Ports: range 49152-65535
- Server
- can be used freely by any client or server

# Anatomy of a Connection

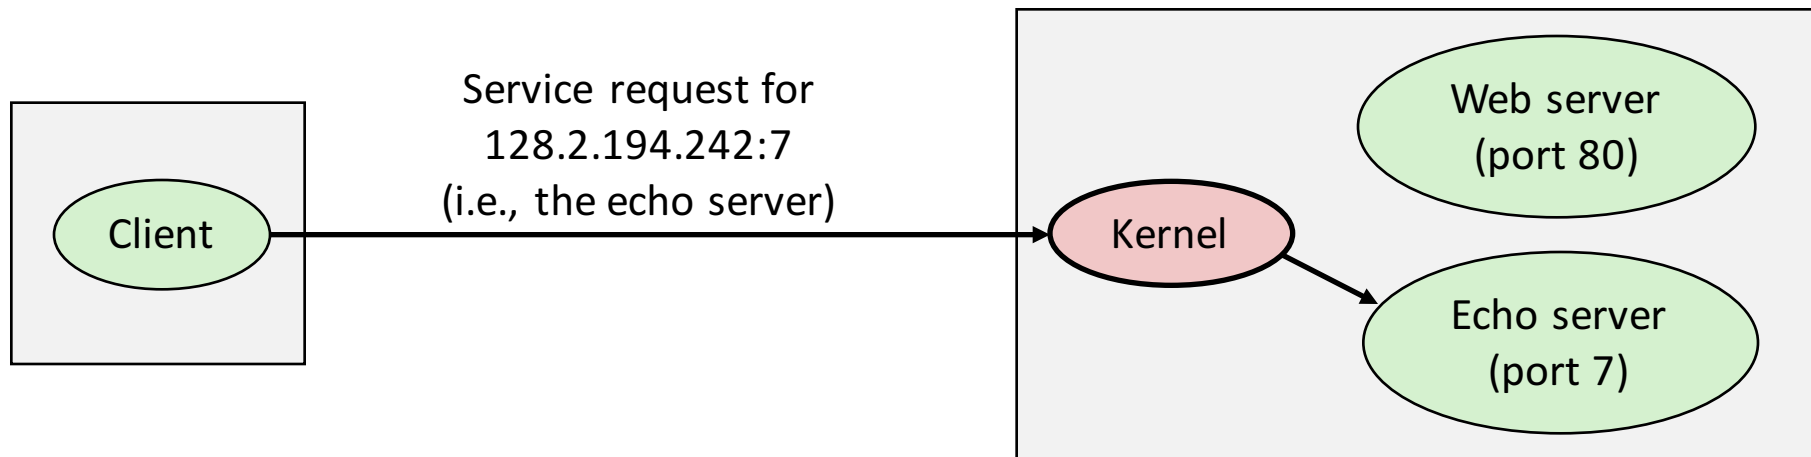
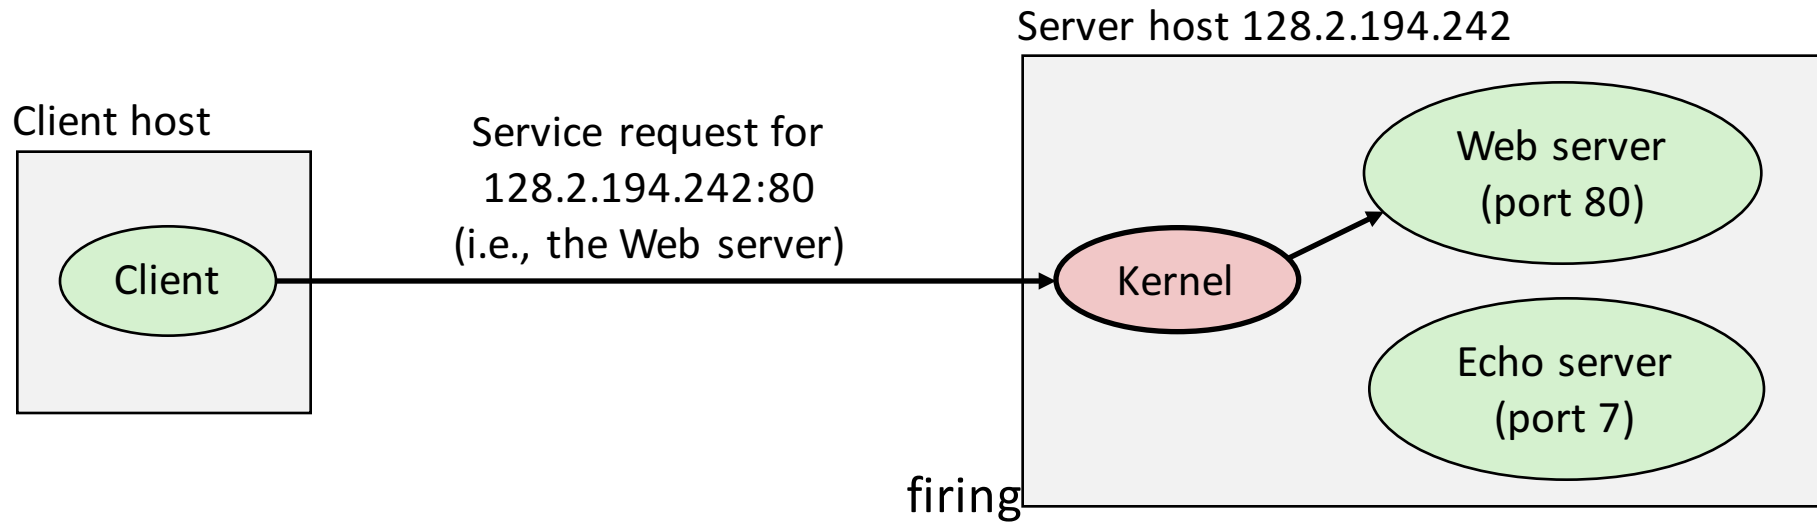
- A connection is uniquely identified by the socket addresses of its endpoints (*socket pair*)
  - (cliaddr:cliport, servaddr:servport)



51213 is an ephemeral port  
allocated by the kernel

80 is a well-known port  
associated with Web servers

# Using Ports to Identify Services






# Sockets Interface

- Set of system-level functions used in conjunction with Unix I/O to build network applications.
- Created in the early 80's as part of the original Berkeley distribution of Unix that contained an early version of the Internet protocols.
- Available on all modern systems
  - Unix variants, Windows, OS X, IOS, Android, ARM

# Socket is like file I/O


- To the kernel, a socket is an endpoint of communication
- To an application, a socket is a file descriptor that lets the application read/write from/to the network
- **Remember:** All Unix I/O devices, including networks, are modeled as “files”
  - Clients and servers communicate with each other by reading from and writing to socket descriptors



- The main distinction between regular file I/O and socket I/O is how the application “opens” the socket descriptors

# Socket v.s. pipe

- Fundamental difference between a socket and a pipe: sockets serve to interconnect two processes that **execute on arbitrary machines**. (Pipes work only on the same machine.)
  - Two processes can execute on the same host, or
  - on networked hosts across the world from each other.The set up and operations on the sockets are the same!



- TCP sockets implement a high level abstraction:
  - It gives the programmer a byte-stream communication channel across networked hosts that is reliable and order-preserving.

# iClicker question

Suppose there is a connection from:

192.168.0.17:53268 to 128.119.40.1:80

Which of these statements is true?

- A) 192.168.0.17 is the client
- B) 53268 is an ephemeral port
- C) 128.119.40.1 is the server
- D) 80 is a well-known port (for HTTP)
- E) All of the above

# iClicker question solution

Suppose there is a connection from:

192.168.0.17:53268 to 128.119.40.1:80

Which of these statements is true?

- A) 192.168.0.17 is the client
- B) 53268 is an ephemeral port
- C) 128.119.40.1 is the server
- D) 80 is a well-known port (for HTTP)
- E) All of the above

# **SOCKET ADDRESSING**

# Ipv4 address: a dotted quad

4 bytes

\_\_\_\_\_

/ \

128.119.240.84

\\_/\_

1 byte

# Storing an IP address in struct in\_addr

```
typedef uint32_t in_addr_t;  
struct in_addr  
{  
    in_addr_t s_addr;  
};
```



# Storing an IP address in struct in\_addr

```
int main(){  
    struct in_addr addr;  
    unsigned char * ip;  
  
    ip = (unsigned char *) &(addr.s_addr);  
  
    ip[0]=128;  
    ip[1]=119;  
    ip[2]=240;  
    ip[3]=84;  
  
    printf("Hello %s\n", inet_ntoa(addr));  
}
```



# Storing an IP address in struct in\_addr

```
int main(){
    struct in_addr addr;
    unsigned char * ip;

    ip = (unsigned char *) &(addr.s_addr);

    ip[0]=128;
    ip[1]=119;
    ip[2]=240;
    ip[3]=84;

    printf("Hello %s\n", inet_ntoa(addr));
}
```



# Storing an IP address in struct in\_addr

```
int main(){
    struct in_addr addr;
    inet_aton("128.119.240.84", &(addr));
    printf("Hello %s\n", inet_ntoa(addr));
}
```

# Socket Address Structures

```
struct sockaddr {  
    uint16_t  sa_family;    /* Protocol family */  
    char      sa_data[14];  /* Address data.  */  
};
```

sa\_family




Family Specific

# Socket Address Structures

- Generic socket address:
  - For address arguments to **connect**, **bind**, and **accept**
  - Necessary only because C did not have generic (**void \***) pointers when the sockets interface was designed

```
struct sockaddr {  
    uint16_t    sa_family;    /* Protocol family */  
    char        sa_data[14];  /* Address data.   */  
};
```

sa\_family

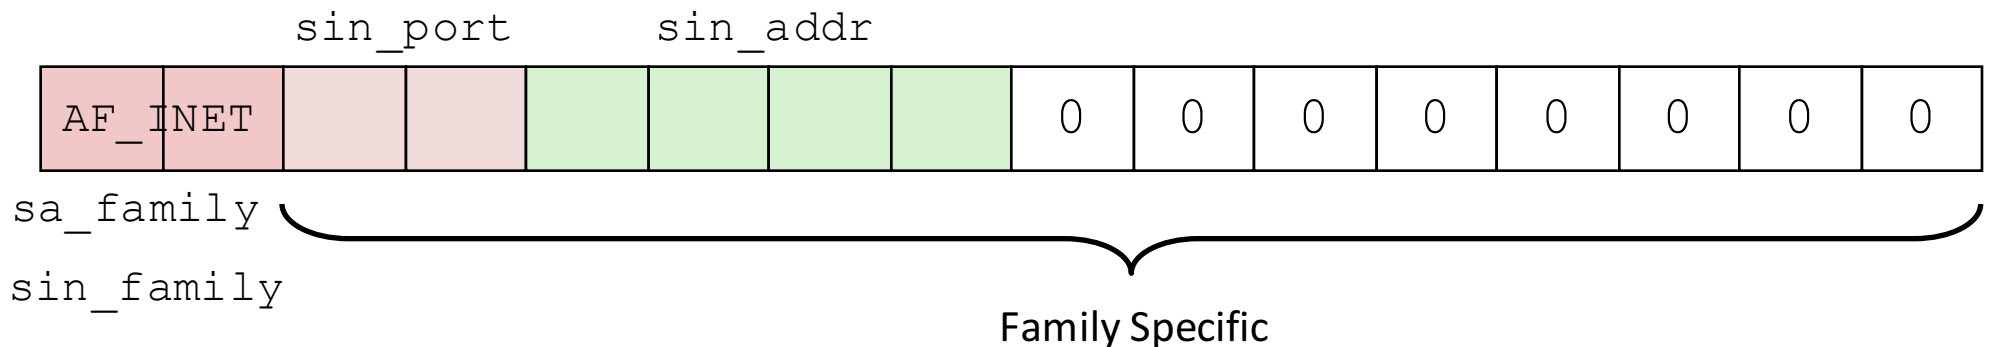


Family Specific

# Socket Address Structures

- Internet-specific socket address


```
struct sockaddr_in {  
    uint16_t      sin_family; /* Protocol family (always AF_INET) */  
    uint16_t      sin_port;  /* Port num in network byte order */  
    struct in_addr sin_addr;  /* IP addr in network byte order */  
    unsigned char  sin_zero[8]; /* Pad to sizeof(struct sockaddr) */  
};
```



# Socket Address Structures

- Internet-specific socket address
  - sockaddr\_in is like a subtype of sockaddr

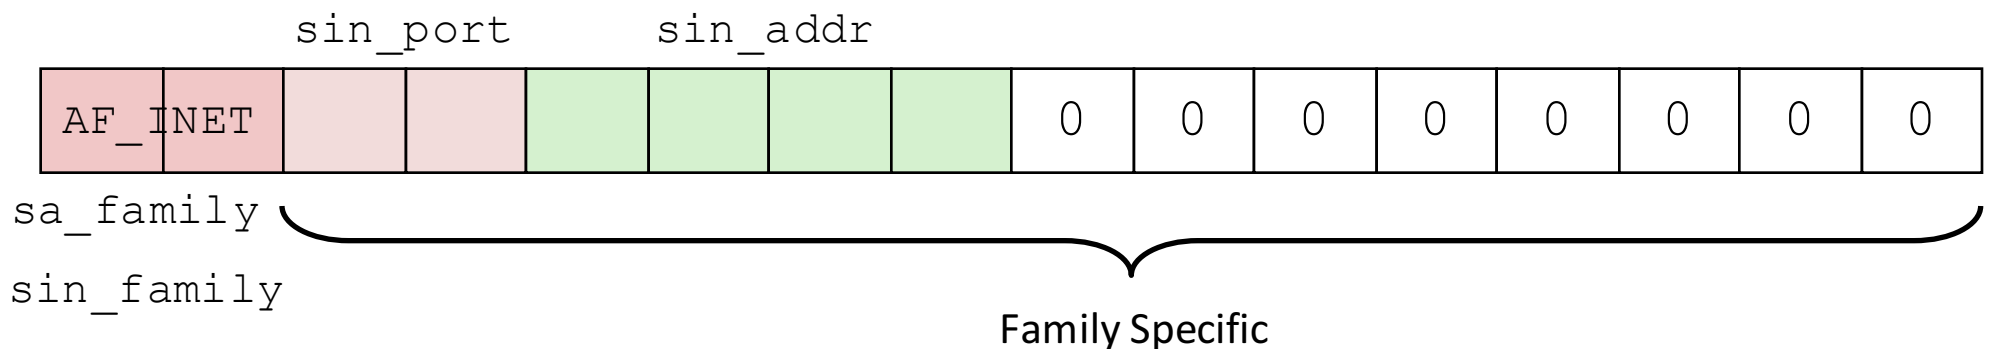
```
struct sockaddr_in {  
    uint16_t      sin_family; /* Protocol family (always AF_INET) */  
    uint16_t      sin_port;   /* Port num in network byte order */  
    struct in_addr sin_addr;   /* IP addr in network byte order */  
    unsigned char  sin_zero[8]; /* Pad to sizeof(struct sockaddr) */  
};
```



# Socket Address Structures

- Internet-specific socket address:
  - `sockaddr_in` is like a subtype of `sockaddr`
  - Must cast `(struct sockaddr_in *)` to `(struct sockaddr *)` for functions that take socket address arguments.

```
struct sockaddr_in {  
    uint16_t      sin_family; /* Protocol family (always AF_INET) */  
    uint16_t      sin_port;   /* Port num in network byte order */  
    struct in_addr sin_addr;   /* IP addr in network byte order */  
    unsigned char  sin_zero[8]; /* Pad to sizeof(struct sockaddr) */  
};
```





# Storing IP Addresses in sockaddr\_in

```
int main(){  
    struct sockaddr saddr;  
    struct sockaddr_in * saddr_in = (struct sockaddr_in *)  
&saddr;  
    inet_aton("128.119.240.84", &(saddr_in->sin_addr));  
    printf("Hello %s\n", inet_ntoa(saddr_in->sin_addr));  
}
```

# Socket Address Structures

```
struct sockaddr_in {  
    uint16_t      sin_family; /* Protocol family (always AF_INET) */  
    uint16_t      sin_port;   /* Port num in network byte order */  
    struct in_addr sin_addr;   /* IP addr in network byte order */  
    unsigned char  sin_zero[8]; /* Pad to sizeof(struct sockaddr) */  
};
```

# Is this correct in little endian systems?

```
struct sockaddr_in {  
    uint16_t      sin_family; /* Protocol family (always AF_INET) */  
    uint16_t      sin_port;   /* Port num in network byte order */  
    struct in_addr sin_addr;   /* IP addr in network byte order */  
    unsigned char  sin_zero[8]; /* Pad to sizeof(struct sockaddr) */  
};
```

```
struct sockaddr saddr;
```

```
struct sockaddr_in * saddr_in = (struct sockaddr_in *) &saddr;
```

```
saddr->sin_port = 80;
```

# Is this correct in little endian systems?

```
struct sockaddr_in {  
    uint16_t      sin_family; /* Protocol family (always AF_INET) */  
    uint16_t      sin_port;   /* Port num in network byte order */  
    struct in_addr sin_addr;   /* IP addr in network byte order */  
    unsigned char  sin_zero[8]; /* Pad to sizeof(struct sockaddr) */  
};
```

struct sockaddr saddr;


struct sockaddr\_in \* saddr\_in = (struct sockaddr\_in \*) &saddr;

saddr->sin\_port = 80;

increasing address  $\longrightarrow$

---

bits:    0   0   0   0   1   0   1   0



# Is this correct in little endian systems?


```
struct sockaddr_in {  
    uint16_t      sin_family; /* Protocol family (always AF_INET) */  
    uint16_t      sin_port;   /* Port num in network byte order */  
    struct in_addr sin_addr;   /* IP addr in network byte order */  
    unsigned char  sin_zero[8]; /* Pad to sizeof(struct sockaddr) */  
};
```

struct sockaddr saddr;

struct sockaddr\_in \* saddr\_in = (struct sockaddr\_in \*) &saddr;

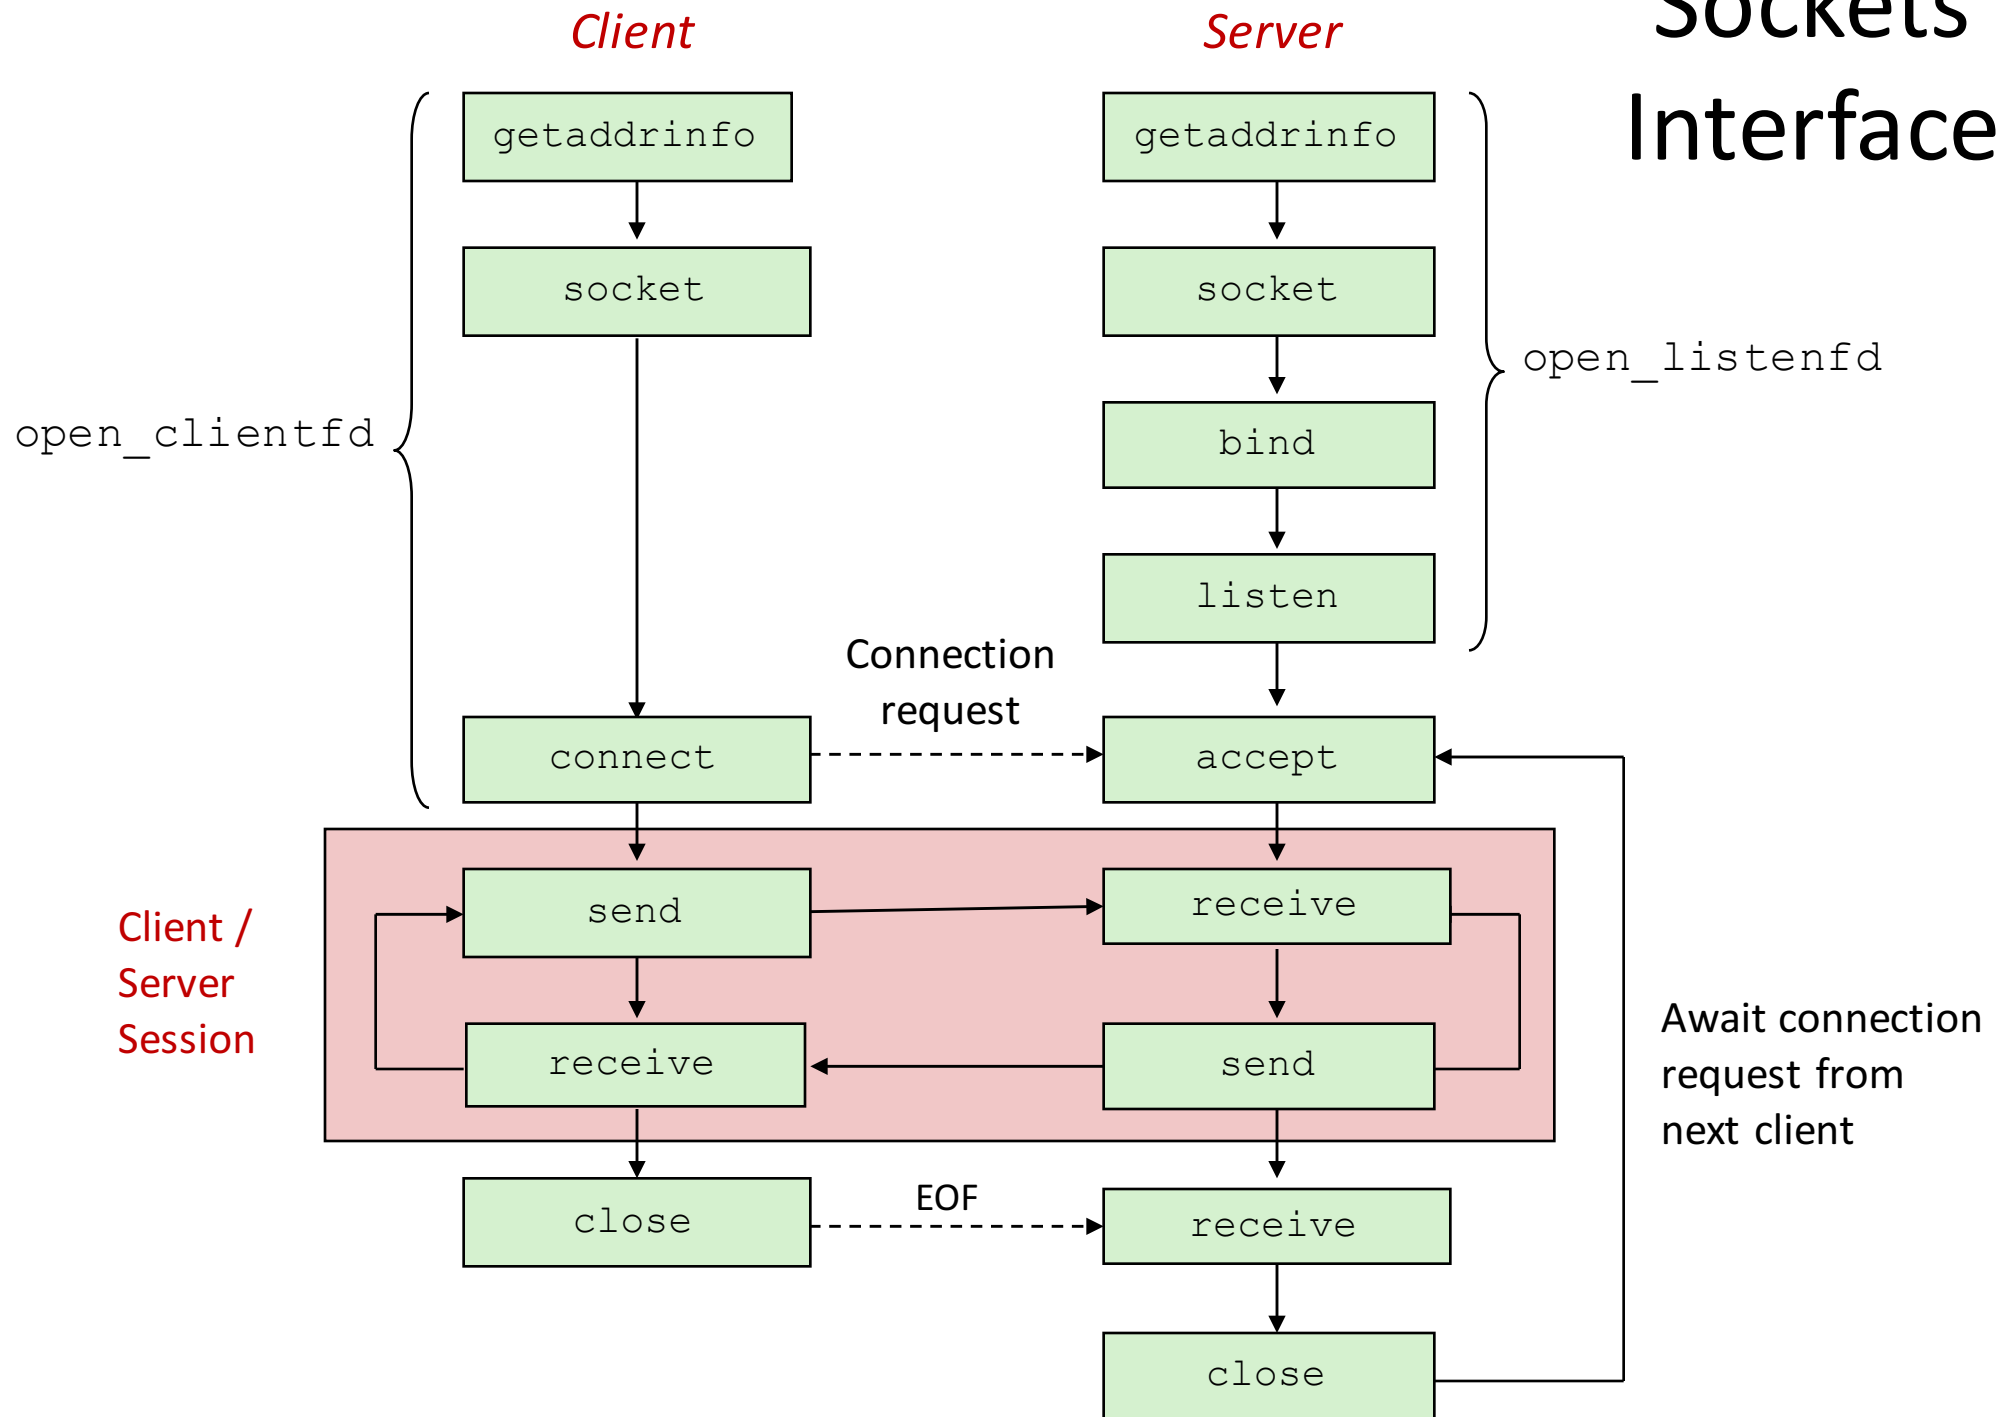
saddr->sin\_port = htons(80);

increasing address →  
-----  
bits: 0 0 0 0 1 0 1 0

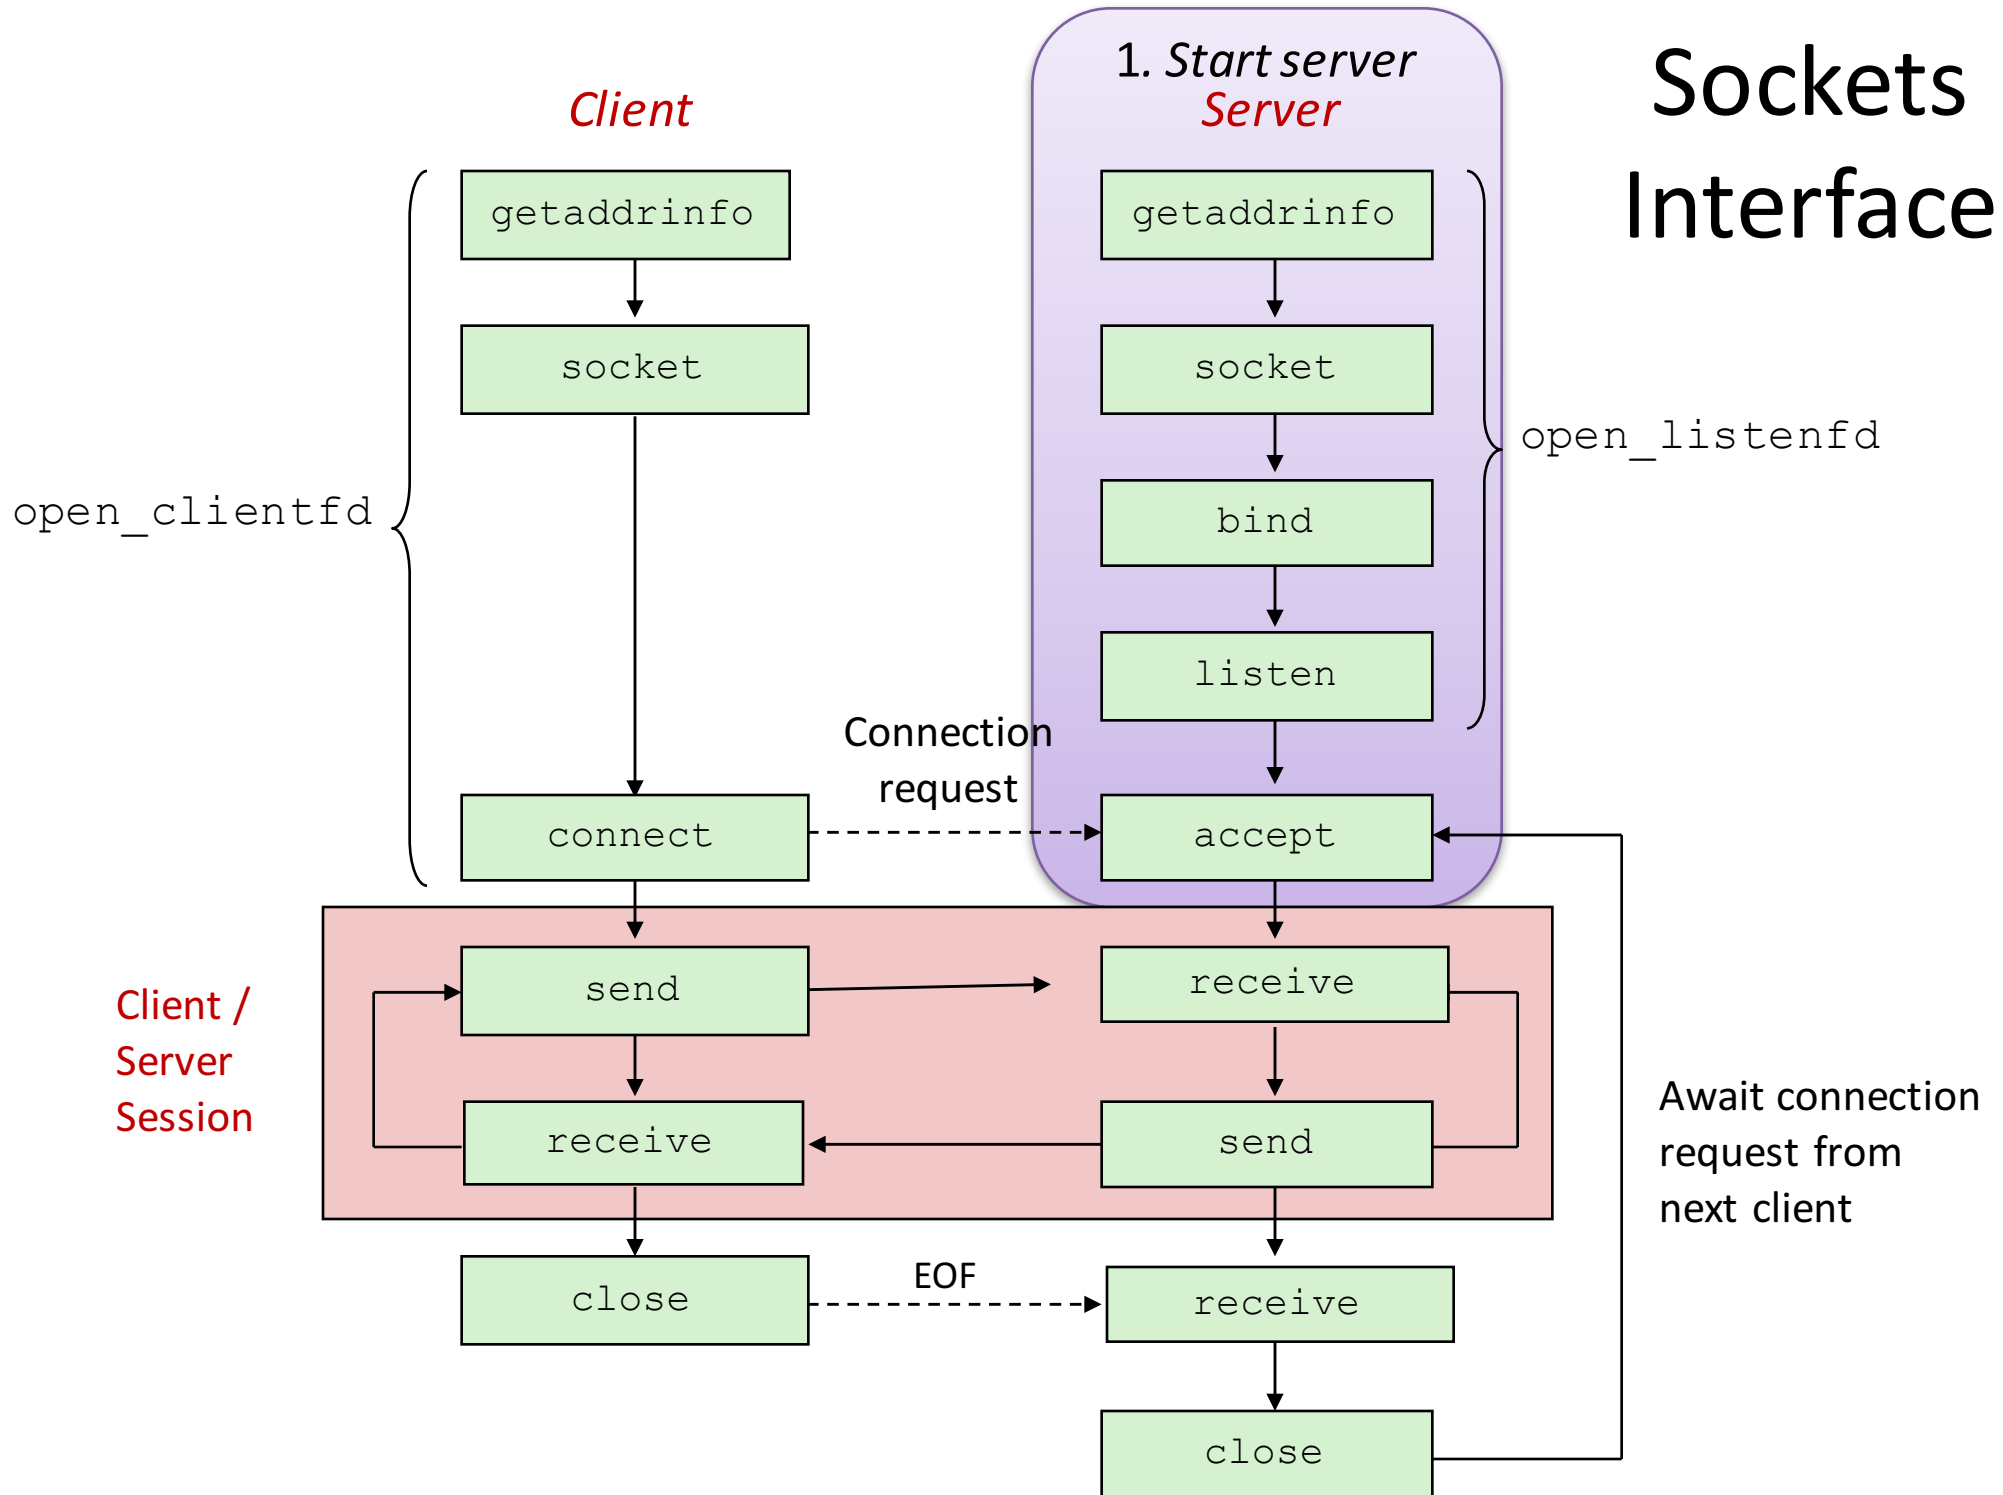


**SOCKET API**

# Sockets Interface



# Sockets Interface





The diagram illustrates the Sockets Interface, showing the flow of data and control between a client and a server. It is divided into two main sections: **2. Start client** (Client) and **1. Start server** (Server).

**Client Side (2. Start client):**

- The process starts with `getaddrinfo`, followed by `socket`, and then `connect`.
- A bracket labeled `fd` groups the `getaddrinfo` and `socket` steps.
- The `connect` step sends a **Connection request** to the server's `accept` step.
- After a successful connection, the client enters a loop of `send` and `receive` operations.
- The process ends with `close`.

**Server Side (1. Start server):**

- The process starts with `getaddrinfo`, followed by `socket`, `bind`, `listen`, and then `accept`.
- A bracket labeled `open_listenfd` groups the `socket`, `bind`, and `listen` steps.
- The `accept` step receives the **Connection request** from the client.
- After a successful connection, the server enters a loop of `receive` and `send` operations.
- The process ends with `close`.

**Communication and Control:**

- The `send` and `receive` operations on both sides are connected by arrows, indicating the flow of data.
- The `close` operation on the client side sends an **EOF** (End of File) signal to the `receive` operation on the server side.
- The `close` operation on the server side sends a signal back to the client's `close` operation, labeled **Await connection request from next client**.

The diagram illustrates the Sockets Interface, showing the flow of data and control between a client and a server. It is divided into two main sections: **2. Start client** (Client) and **1. Start server** (Server).

**Client Side (2. Start client):**

- The process starts with `getaddrinfo`, followed by `socket`, and then `connect`.
- A bracket labeled `fd` groups the `getaddrinfo` and `socket` steps.
- The `connect` step sends a **Connection request** to the server's `accept` step.
- After a successful connection, the client enters a loop of `send` and `receive` operations.
- The process ends with `close`.

**Server Side (1. Start server):**

- The process starts with `getaddrinfo`, followed by `socket`, `bind`, `listen`, and then `accept`.
- A bracket labeled `open_listenfd` groups the `socket`, `bind`, and `listen` steps.
- The `accept` step receives the **Connection request** from the client.
- After a successful connection, the server enters a loop of `receive` and `send` operations.
- The process ends with `close`.

**Communication and Control:**

- The `send` and `receive` operations on both sides are connected by arrows, indicating the flow of data.
- The `close` operation on the client side sends an **EOF** (End of File) signal to the `receive` operation on the server side.
- The `close` operation on the server side sends a signal back to the client's `close` operation, labeled **Await connection request from next client**.

The diagram illustrates the Sockets Interface, showing the flow of data between a client and a server. It is divided into two main sections: **2. Start client** (Client) and **1. Start server** (Server).

**Client Side (2. Start client):**

- The process starts with `getaddrinfo`, followed by `socket`, and then `connect`.
- A bracket labeled `fd` groups the `getaddrinfo` and `socket` steps.
- The `connect` step sends a **Connection request** to the server's `accept` step.
- After a successful connection, the client enters a loop of `send` and `receive` operations.
- The process ends with `close`.

**Server Side (1. Start server):**

- The process starts with `getaddrinfo`, followed by `socket`, `bind`, and `listen`.
- A bracket labeled `open_listenfd` groups the `socket`, `bind`, and `listen` steps.
- The `accept` step receives the **Connection request** from the client.
- After a successful connection, the server enters a loop of `receive` and `send` operations.
- The process ends with `close`.

**Communication Flow:**

- The client's `send` operation sends data to the server's `receive` operation.
- The server's `send` operation sends data to the client's `receive` operation.
- The client's `close` operation sends an **EOF** (End of File) signal to the server's `receive` operation.
- After the server's `close` operation, a line labeled **Await connection request from next client** loops back to the `accept` step.

The diagram illustrates the Sockets Interface, showing the flow of data and control between a client and a server. It is divided into two main sections: **2. Start client Client** and **1. Start server Server**.

**Client Side (2. Start client Client):**

- The process starts with `getaddrinfo`, followed by `socket`, and then `connect`.
- A bracket labeled `fd` groups the `getaddrinfo` and `socket` steps.
- The `connect` step sends a **Connection request** to the server's `accept` step.
- Once connected, the client enters a loop of `send` and `receive` operations.
- The process ends with `close`.

**Server Side (1. Start server Server):**

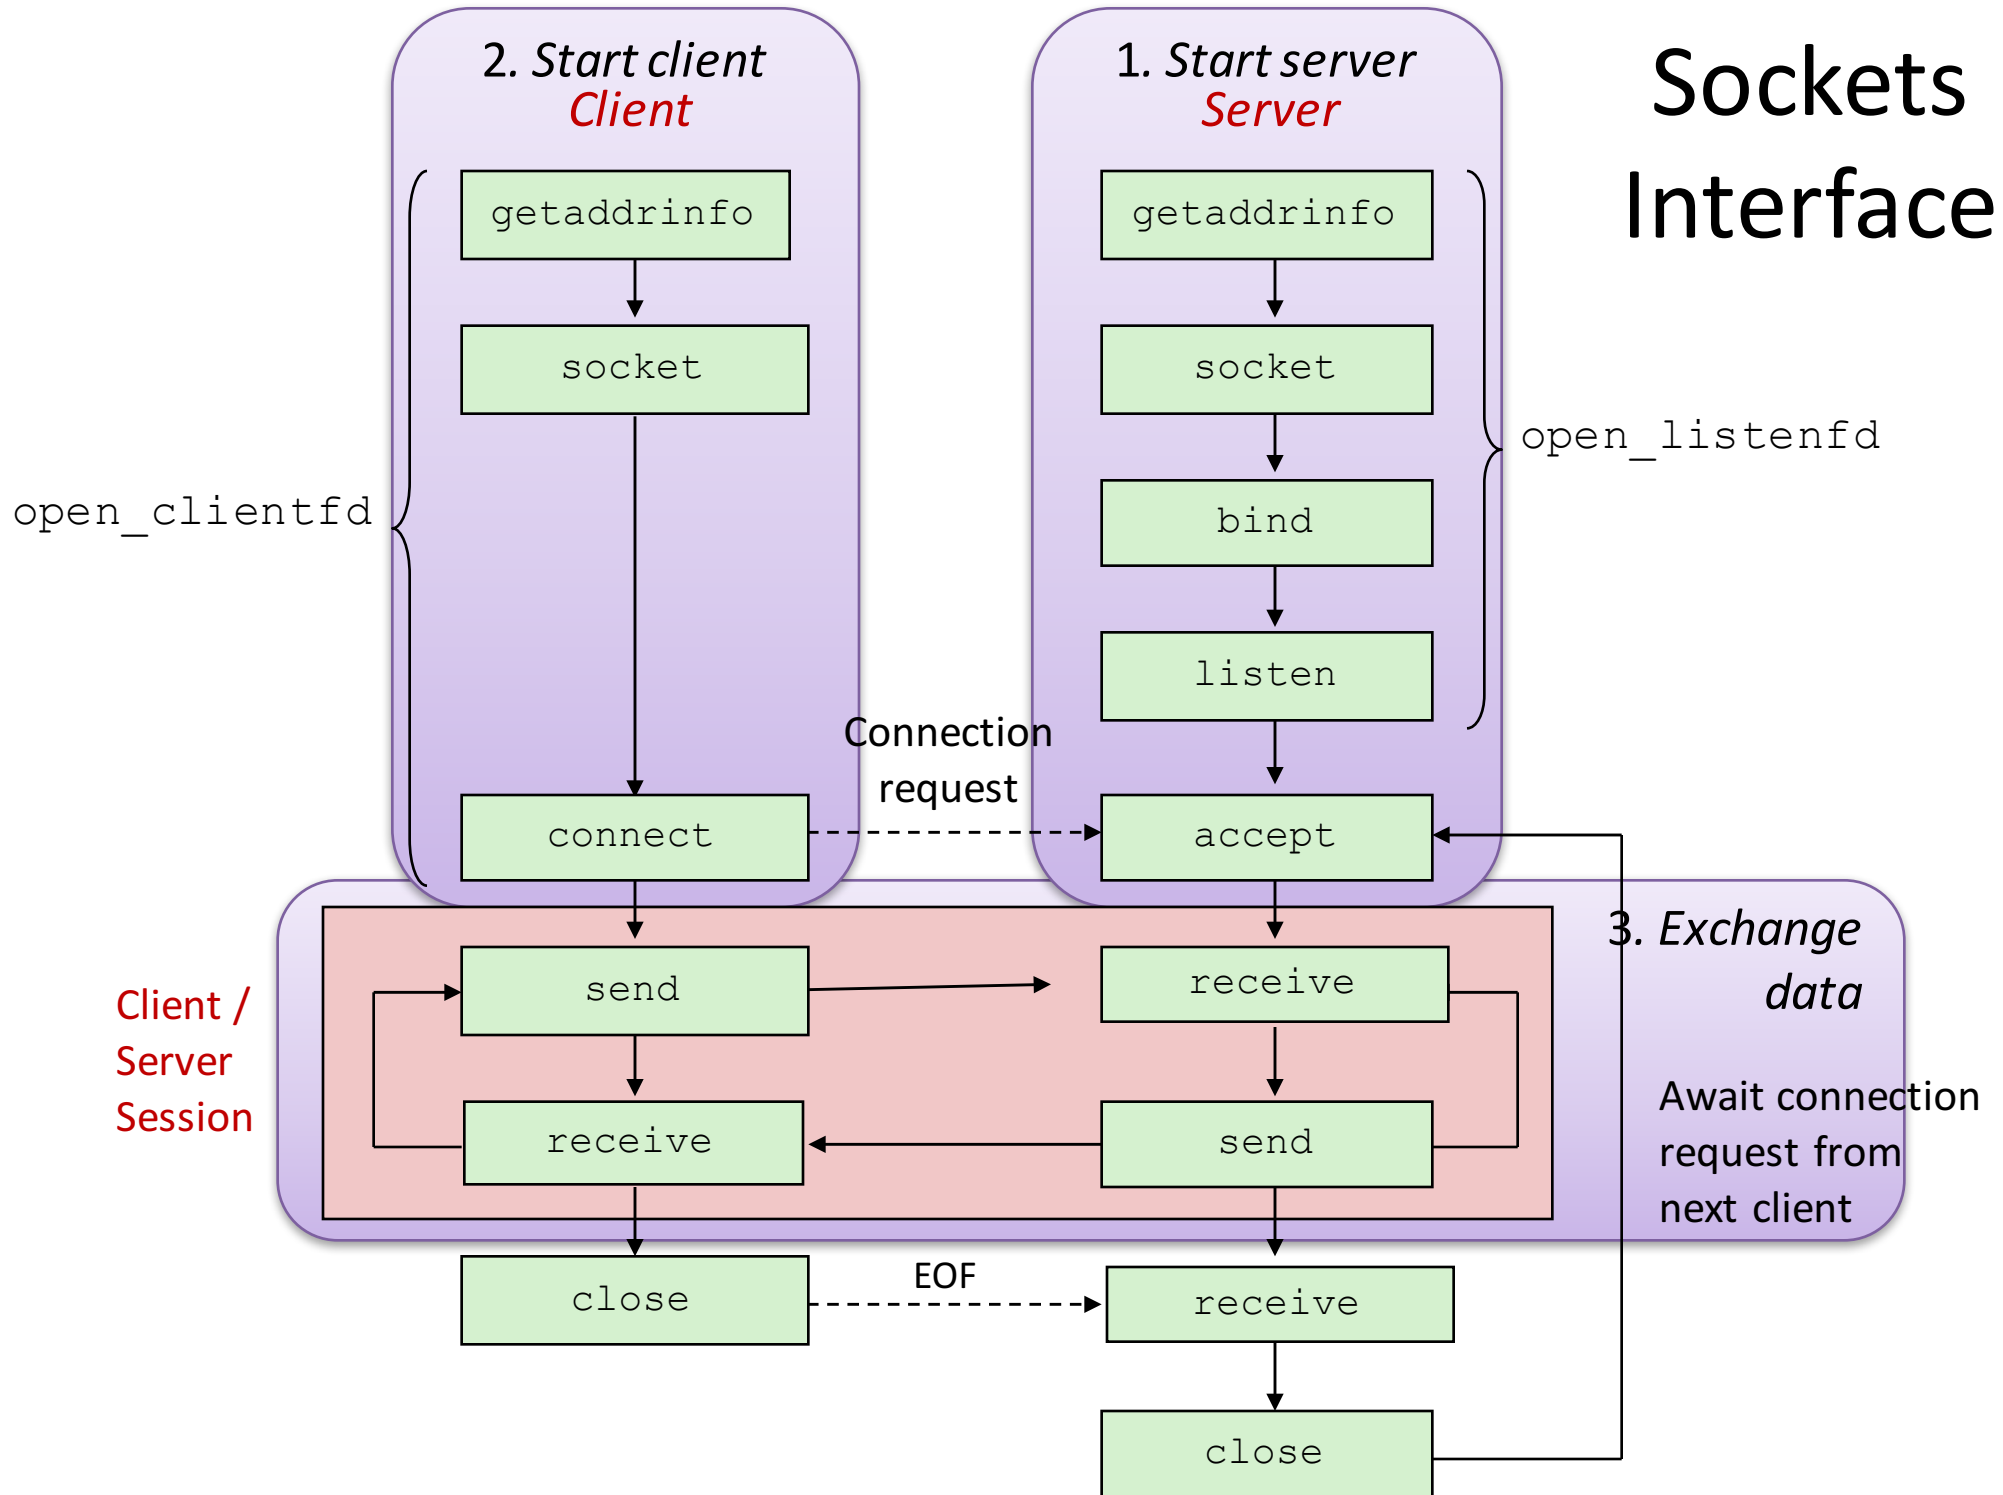
- The process starts with `getaddrinfo`, followed by `socket`, `bind`, and `listen`.
- A bracket labeled `open_listenfd` groups the `socket`, `bind`, and `listen` steps.
- The `accept` step receives the **Connection request** from the client.
- Once connected, the server enters a loop of `receive` and `send` operations.
- The process ends with `close`.

**Data Flow and Control:**

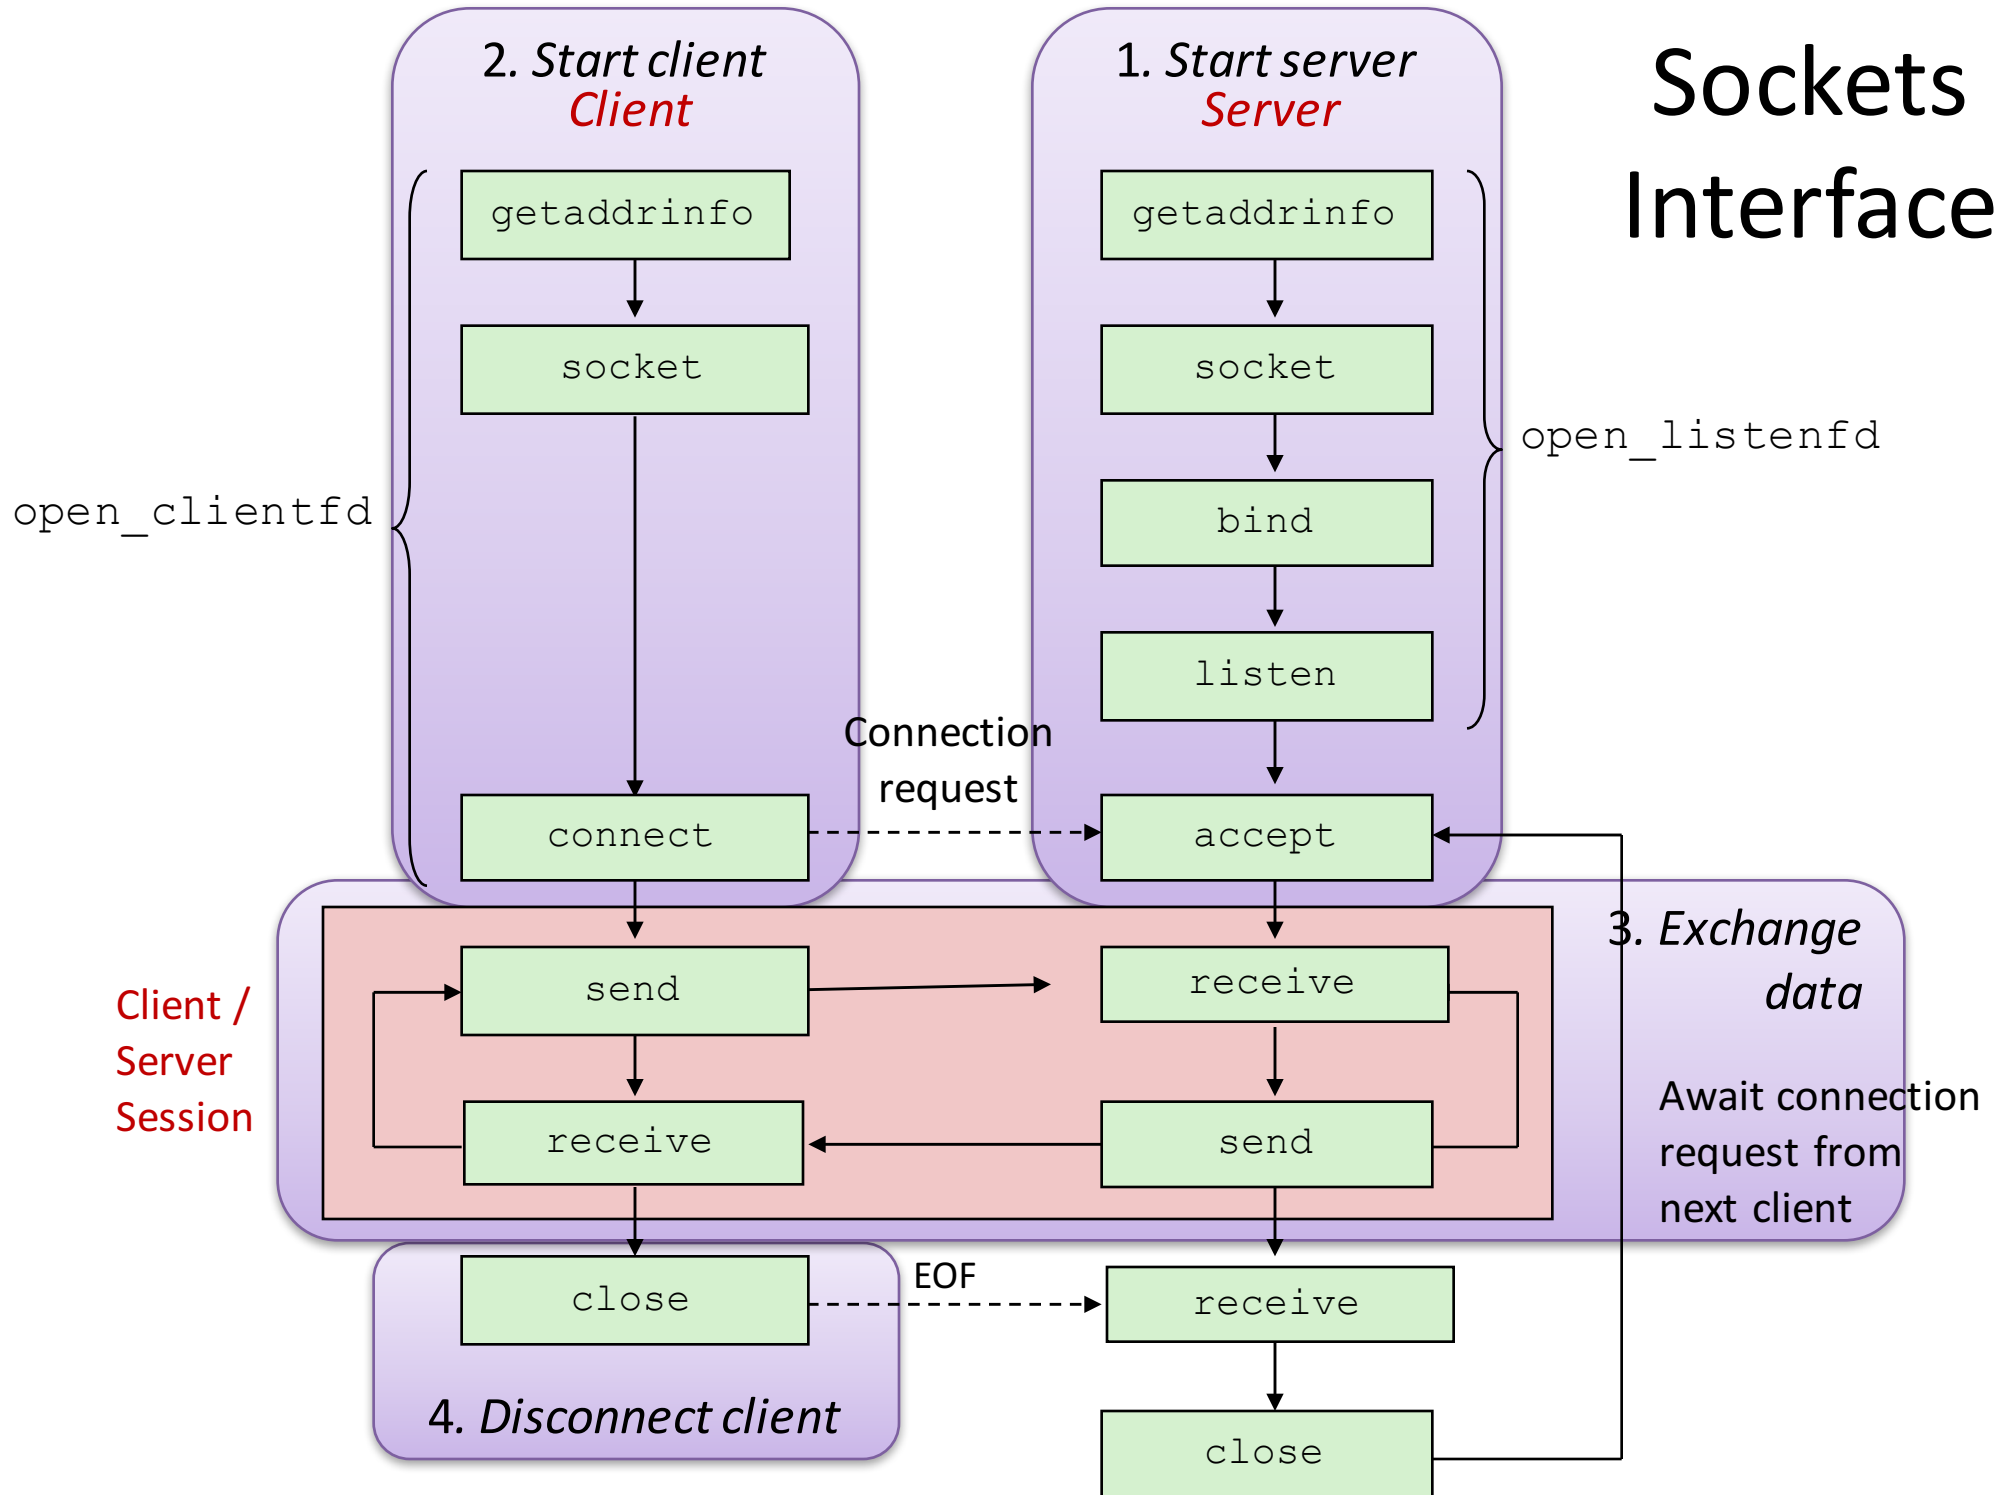
- Connection request:** A dashed arrow from the client's `connect` to the server's `accept`.
- EOF (End of File):** A dashed arrow from the client's `close` to the server's `receive`.
- Await connection request from next client:** A solid arrow from the server's `close` back to the `accept` step.
- Data Exchange:** Solid arrows show data flowing from the client's `send` to the server's `receive`, and from the server's `send` to the client's `receive`.

[illegible]

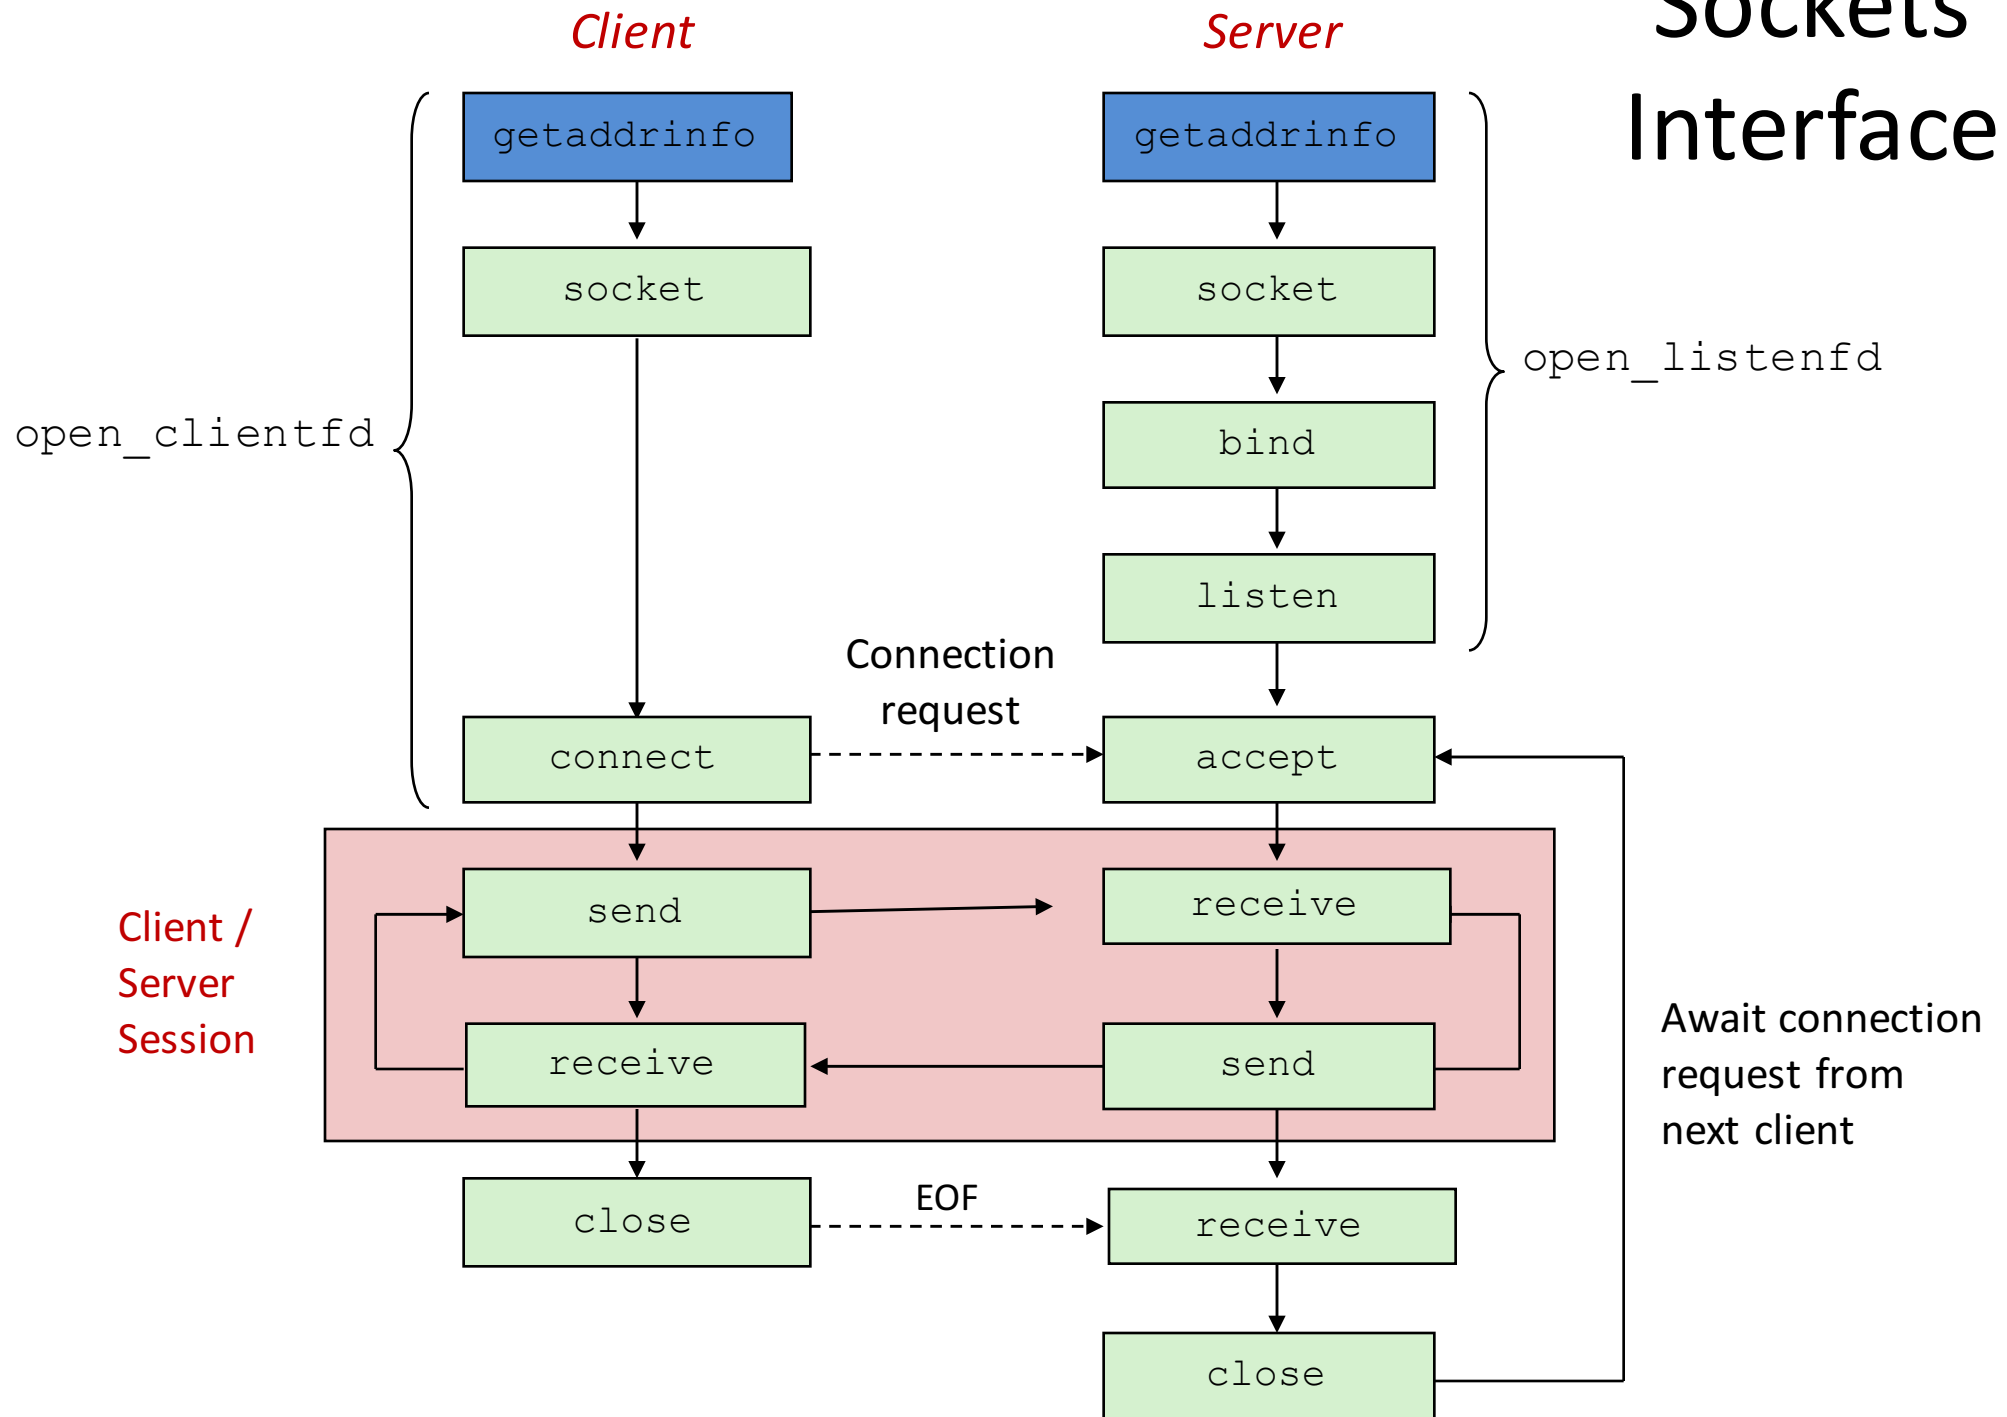
# Sockets Interface



# Sockets Interface



# Sockets Interface



# Host and Service Conversion: getaddrinfo

```
int getaddrinfo(const char *host,          /* Hostname or address */
               const char *service,       /* Port or service name */
               const struct addrinfo *hints, /* Input parameters */
               struct addrinfo **result);  /* Output linked list */
```

- Given `host` and `service`, `getaddrinfo` returns `result` that points to a linked list of `addrinfo` structs, each of which points to a corresponding socket address struct, and which contains arguments for the sockets interface functions.

# Host and Service Conversion: getaddrinfo


```
int getaddrinfo(const char *host,          /* Hostname or address */
               const char *service,       /* Port or service name */
               const struct addrinfo *hints, /* Input parameters */
               struct addrinfo **result);  /* Output linked list */

void freeaddrinfo(struct addrinfo *result); /* Free linked list */

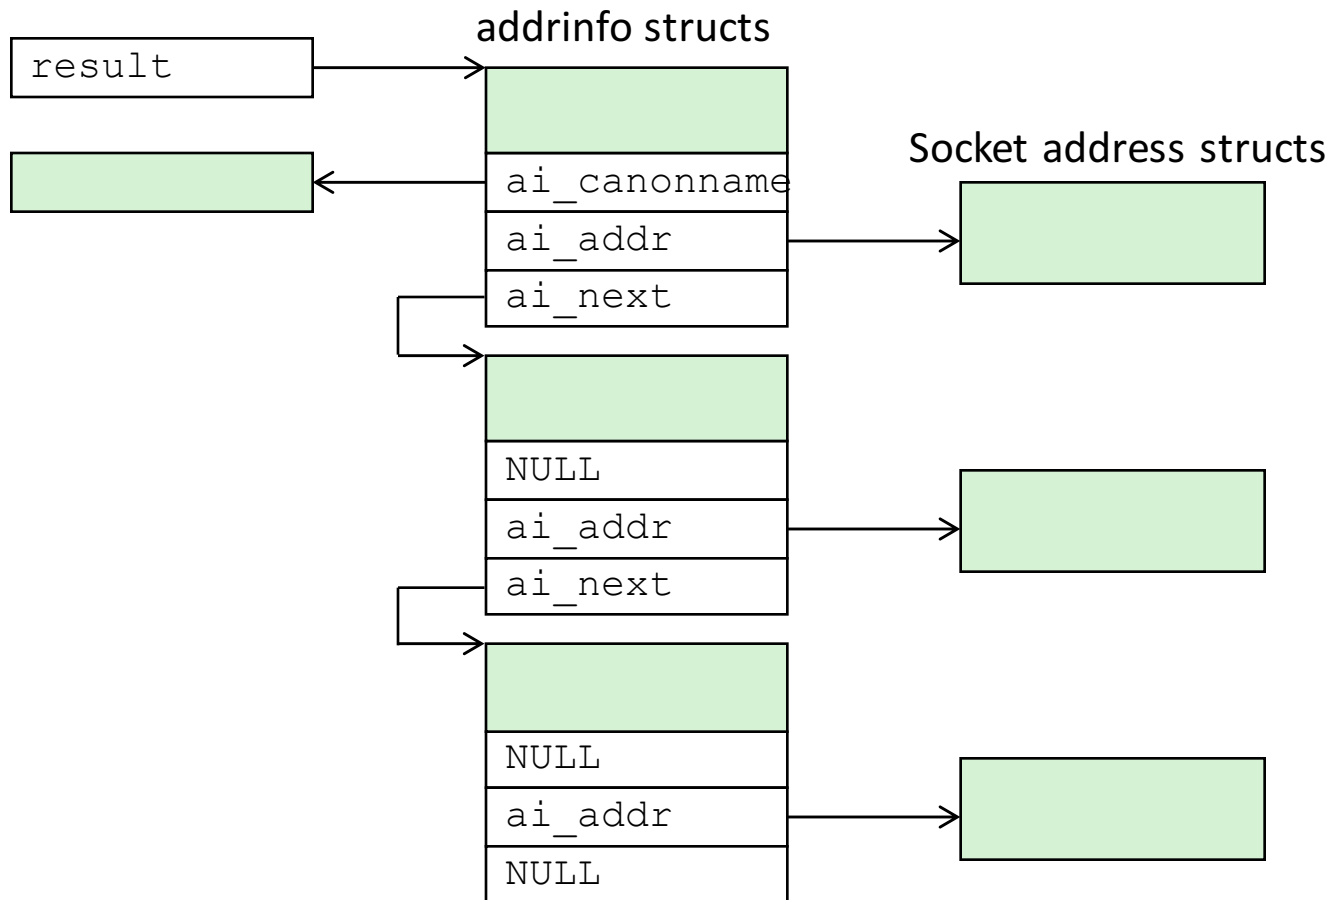
const char *gai_strerror(int errcode);     /* Return error msg */
```

- Given host and service, getaddrinfo returns result that points to a linked list of `addrinfo` structs, each of which points to a corresponding socket address struct, and which contains arguments for the sockets interface functions.
- Helper functions:
  - `freeaddrinfo` frees the entire linked list.
  - `gai_strerror` converts error code to an error message.

# Linked List Returned by `getaddrinfo`



# Linked List Returned by `getaddrinfo`



- Clients: walk this list, trying each socket address in turn, until the calls to `socket` and `connect` succeed.
- Servers: walk the list until calls to `socket` and `bind` succeed.



# addrinfo Struct

```
struct addrinfo {  
    int             ai_flags;        /* Hints argument flags */  
    int             ai_family;      /* First arg to socket function */  
    int             ai_socktype;    /* Second arg to socket function */  
    int             ai_protocol;    /* Third arg to socket function */  
    char            *ai_canonname;   /* Canonical host name */  
    size_t          ai_addrlen;     /* Size of ai_addr struct */  
    struct sockaddr *ai_addr;        /* Ptr to socket address structure */  
    struct addrinfo *ai_next;       /* Ptr to next item in linked list */  
};
```

- Each `addrinfo` struct returned by `getaddrinfo` contains arguments that can be passed directly to `socket` function.
- Also points to a socket address struct that can be passed directly to `connect` and `bind` functions.

# Host and Service Conversion: `getnameinfo`

- `getnameinfo` is the inverse of `getaddrinfo`, converting a socket address to the corresponding host and service.
  - Replaces obsolete `gethostbyaddr` and `getservbyport` funcs.
  - Reentrant and protocol independent.

```
int getnameinfo(const SA *sa, socklen_t salen, /* In: socket addr */
                char *host, size_t hostlen, /* Out: host */
                char *serv, size_t servlen, /* Out: service */
                int flags); /* optional flags */
```

# Conversion Example

```
#include "csapp.h"

int main(int argc, char **argv)
{
    struct addrinfo *p, *listp, hints;
    char buf[MAXLINE];
    int rc, flags;

    /* Get a list of addrinfo records */
    memset(&hints, 0, sizeof(struct addrinfo));
    hints.ai_family = AF_INET;          /* IPv4 only */
    hints.ai_socktype = SOCK_STREAM;    /* Connections only */
    if ((rc = getaddrinfo(argv[1], NULL, &hints, &listp)) != 0) {
        fprintf(stderr, "getaddrinfo error: %s\n", gai_strerror(rc));
        exit(1);
    }
}
```

hostinfo.c

# Conversion Example (cont)

```
/* Walk the list and display each IP address */
flags = NI_NUMERICHOST; /* Display address instead of name */
for (p = listp; p; p = p->ai_next) {
    Getnameinfo(p->ai_addr, p->ai_addrlen,
                buf, MAXLINE, NULL, 0, flags);
    printf("%s\n", buf);
}

/* Clean up */
Freeaddrinfo(listp);

exit(0);
}
```

hostinfo.c

# Running hostinfo

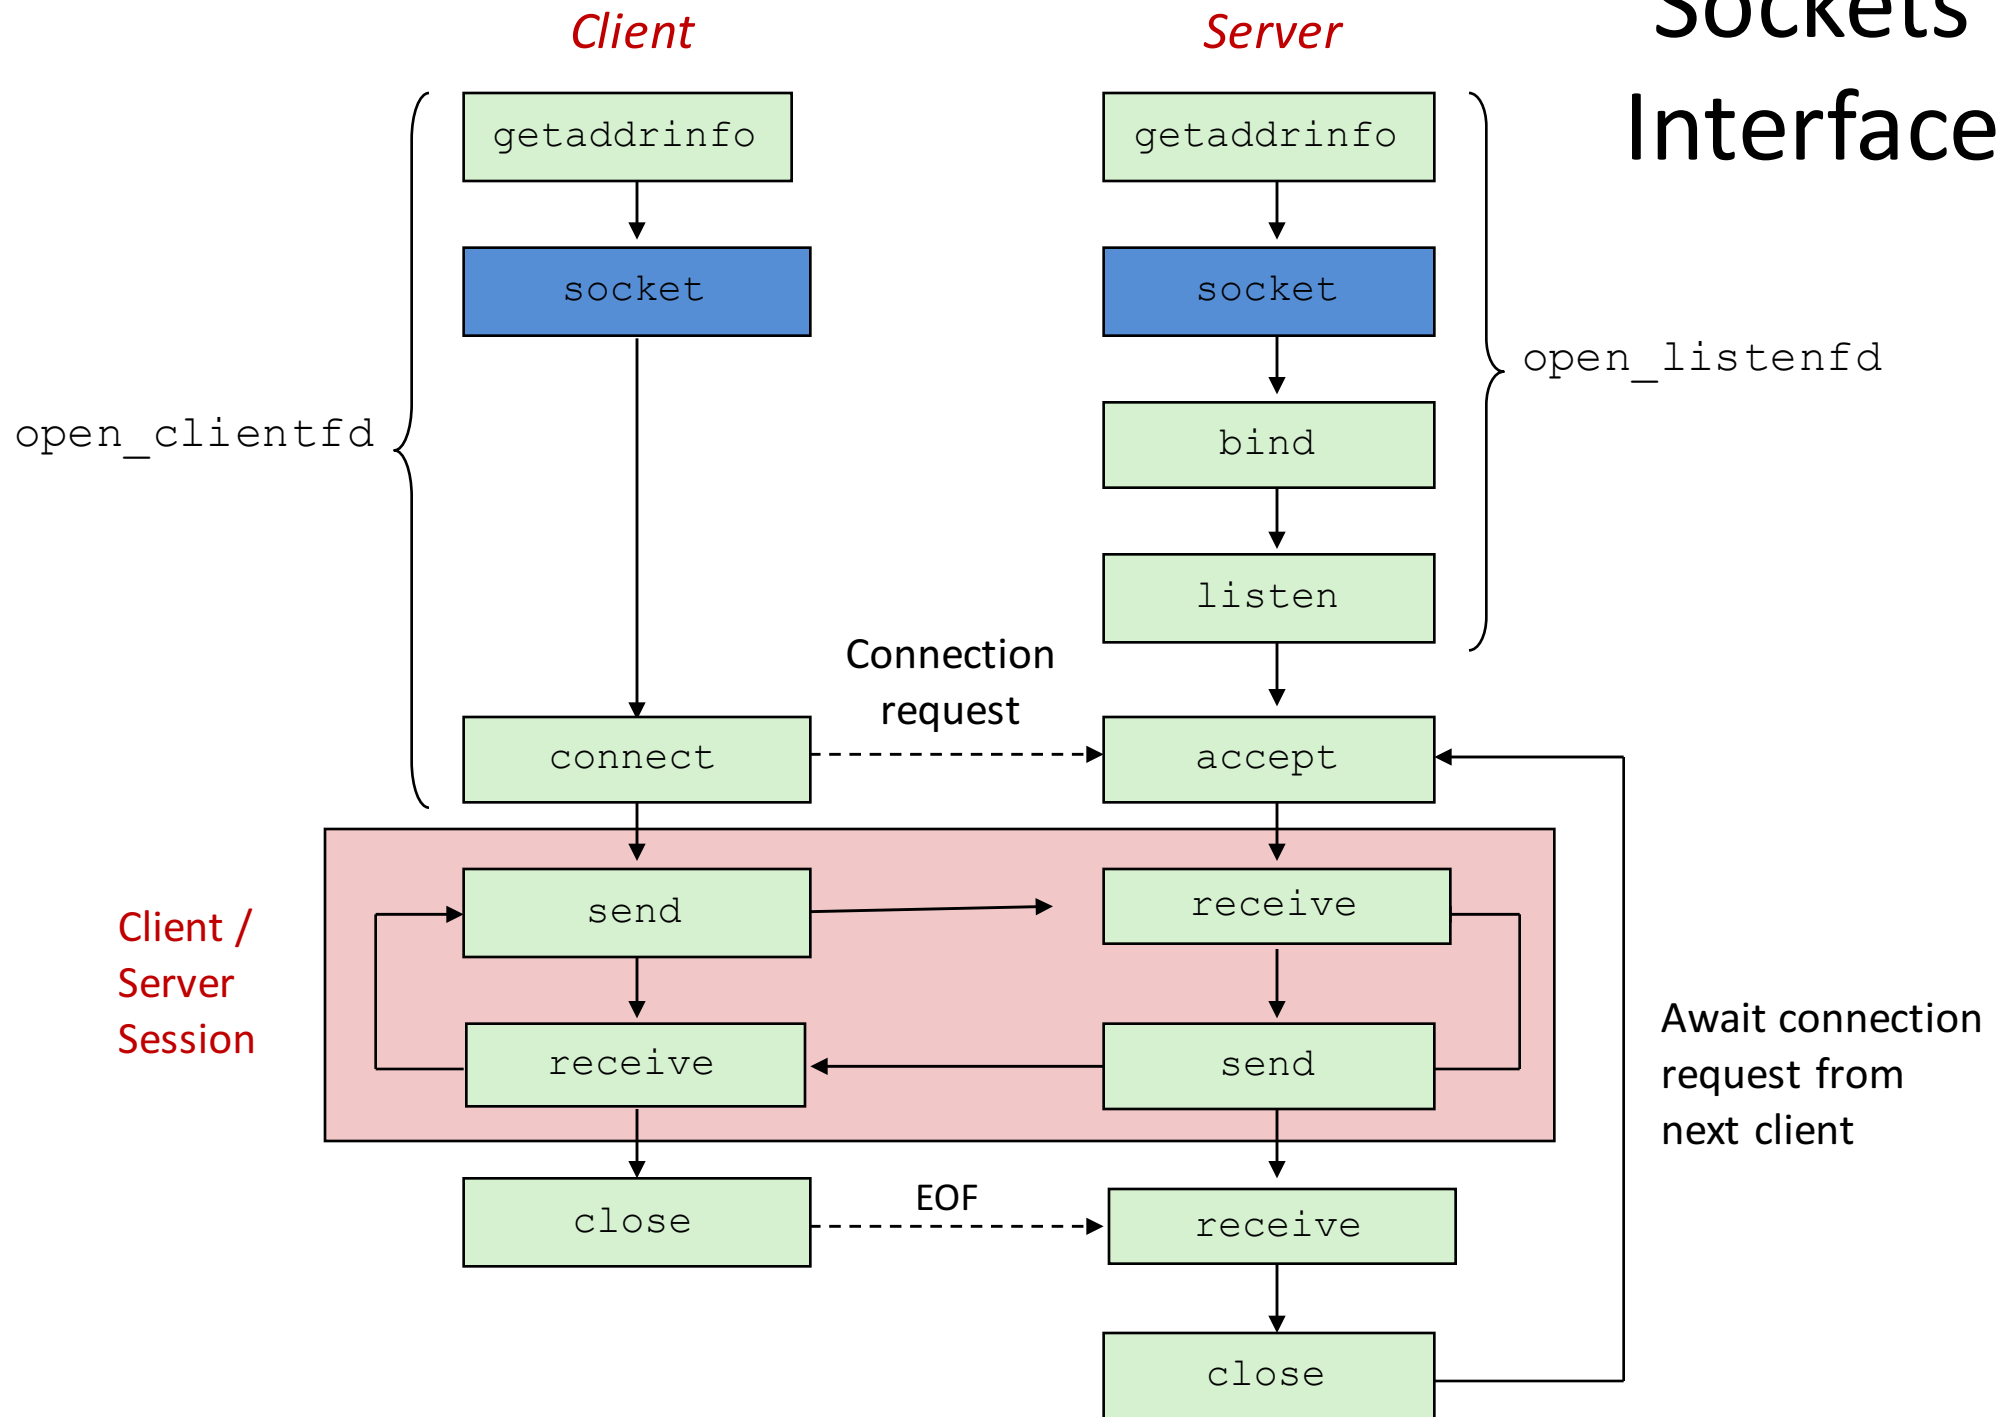
```
student@osboxes:~/Code/lec24$ ./hostinfo amazon.com  
54.239.25.192  
54.239.25.208  
54.239.25.200  
54.239.17.6  
54.239.26.128  
54.239.17.7
```

```
student@osboxes:~/Code/lec24$ ./hostinfo google.com  
216.58.219.206
```

```
student@osboxes:~/Code/lec24$ ./hostinfo localhost  
127.0.0.1
```

```
student@osboxes:~/Code/lec24$ ./hostinfo umass.edu  
128.119.103.148
```

# Sockets Interface



# Sockets Interface: `socket`

- Clients and servers use the `socket` function to create a *socket descriptor*:
- Example:

```
int socket(int domain, int type, int protocol)
```

Indicate protocol number,  
normally 0

```
int clientfd = Socket(AF_INET, SOCK_STREAM, 0);
```

Indicates that we are using  
32-bit IPV4 addresses

Indicates socket type

# Sockets Interface: `socket`

- Clients and servers use the `socket` function to create a *socket descriptor*.
- Example:

Protocol specific! Best practice is to use `getaddrinfo` to generate the parameters automatically, so that code is protocol independent.

```
int socket(int domain, int type, int protocol)
```

Indicate protocol number,  
normally 0

```
int clientfd = Socket(AF_INET, SOCK_STREAM, 0);
```

Indicates that we are using  
32-bit IPV4 addresses

Indicates socket type



# addrinfo Struct

```
struct addrinfo {  
    int             ai_flags;        /* Hints argument flags */  
    int             ai_family;       /* First arg to socket function */  
    int             ai_socktype;     /* Second arg to socket function */  
    int             ai_protocol;     /* Third arg to socket function */  
    char            *ai_canonname;    /* Canonical host name */  
    size_t          ai_addrlen;      /* Size of ai_addr struct */  
    struct sockaddr *ai_addr;         /* Ptr to socket address structure */  
    struct addrinfo *ai_next;         /* Ptr to next item in linked list */  
};
```

- Each addrinfo struct returned by `getaddrinfo` contains arguments that can be passed directly to `socket` function.
- Also points to a socket address struct that can be passed directly to `connect` and `bind` functions.

# Sockets Interface: `socket`

- Clients and servers use the `socket` function to create a *socket descriptor*.
- Example:

Protocol specific! Best practice is to use `getaddrinfo` to generate the parameters automatically, so that code is protocol independent.

```
int socket(int domain, int type, int protocol)
```

Indicate protocol number,  
normally 0

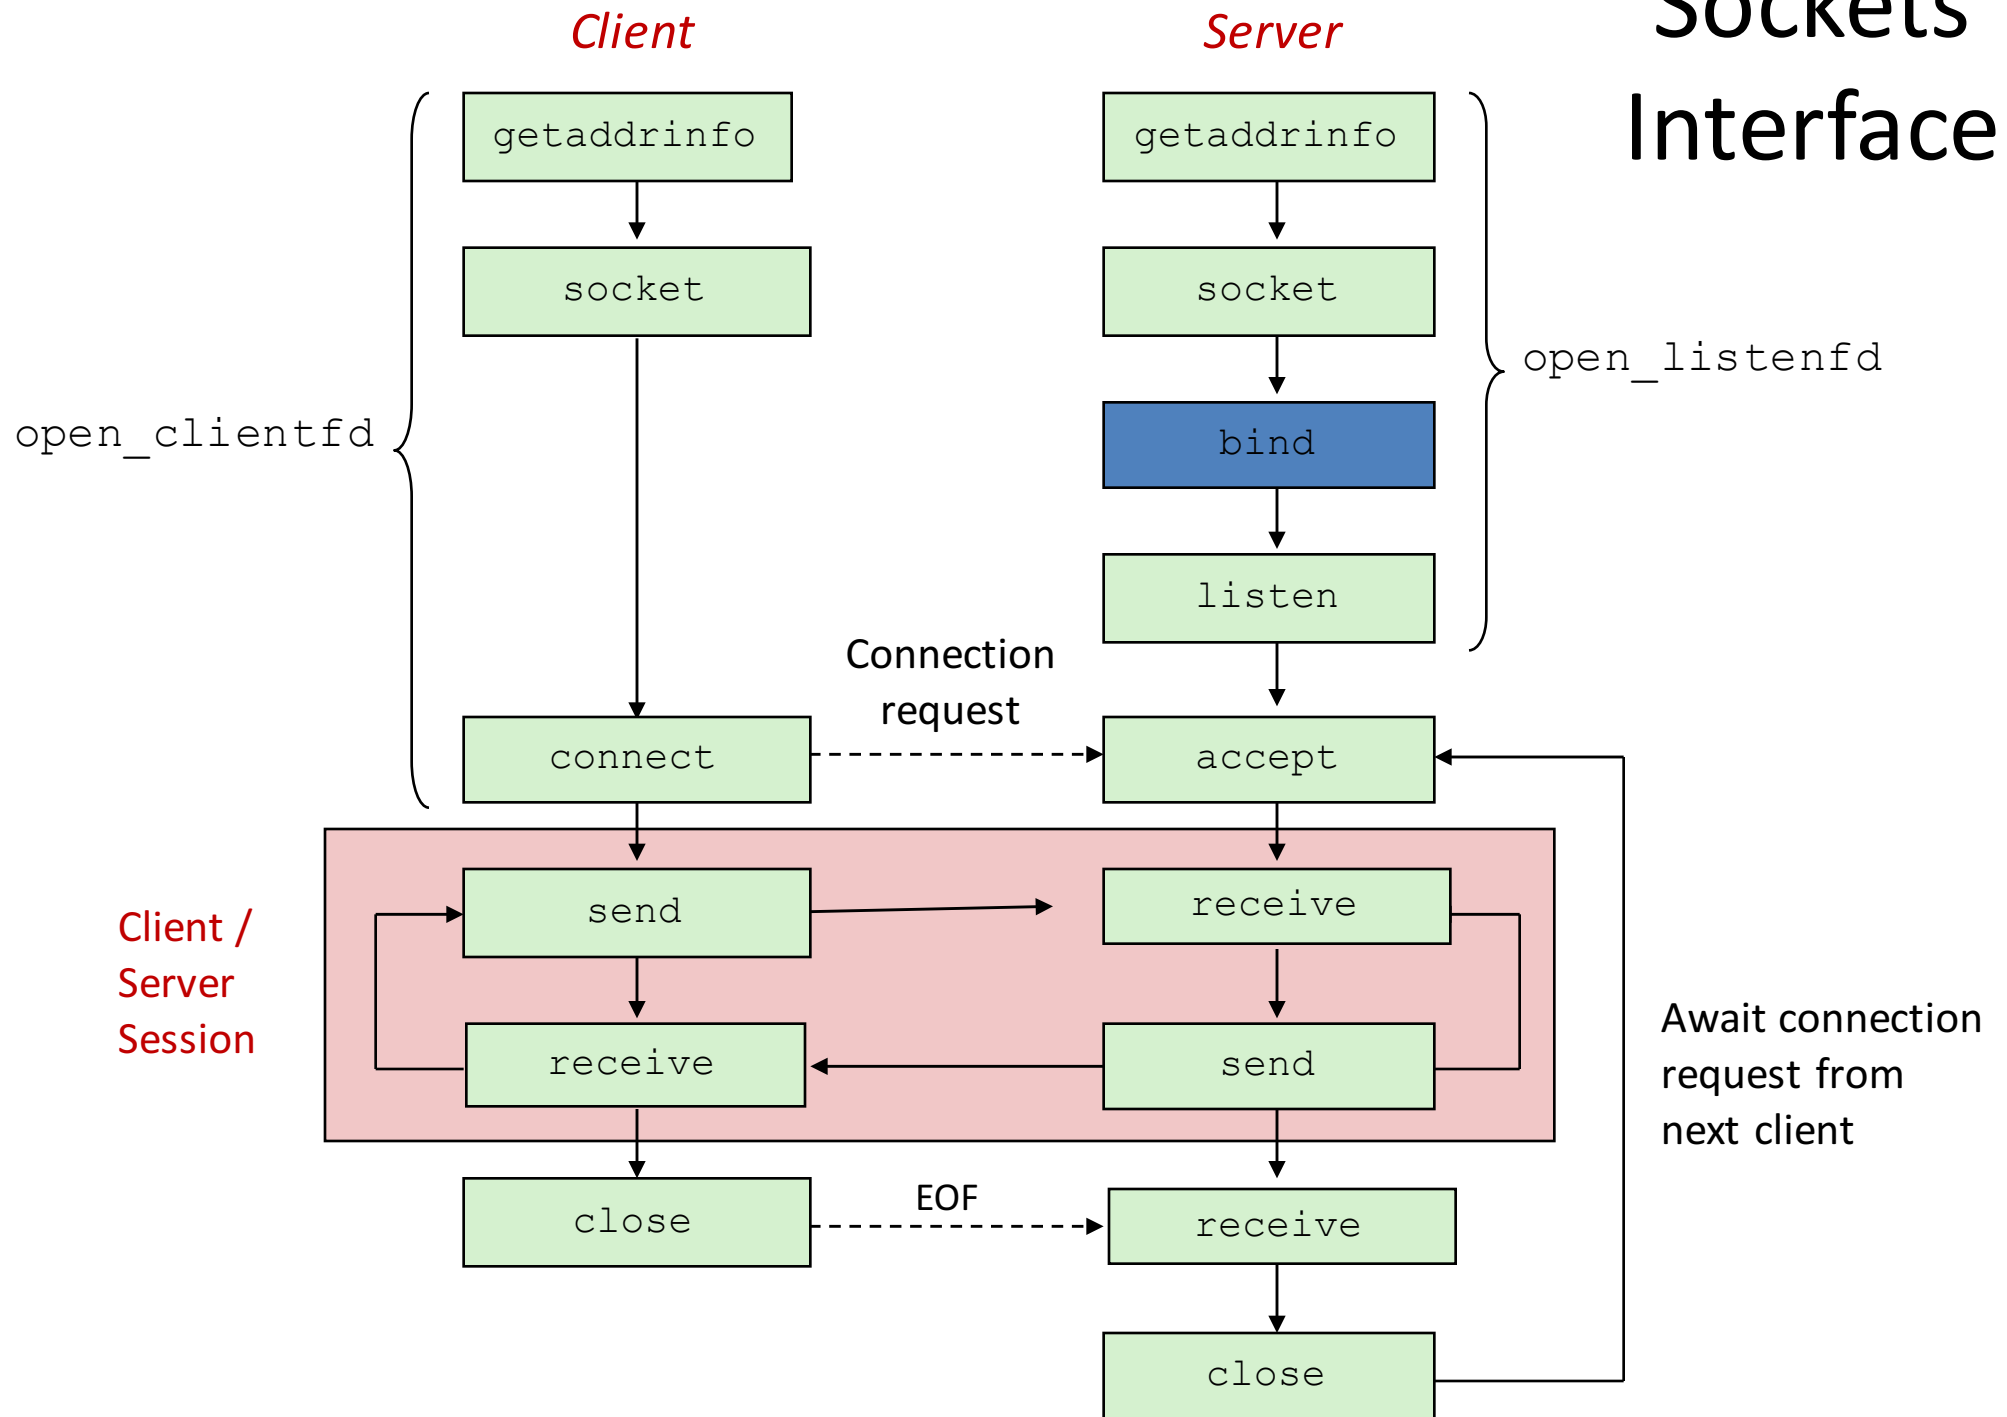
```
int clientfd = Socket(AF_INET, SOCK_STREAM, 0);
```

A socket descriptor

Indicates that we are using  
32-bit IPV4 addresses

Indicates socket type

# Sockets Interface




# Sockets Interface: `bind`

- A server uses `bind` to ask the kernel to associate the server's socket address with a socket descriptor:
- The process can read bytes that arrive on the connection whose endpoint is `addr` by reading from descriptor `sockfd`.
- Similarly, writes to `sockfd` are transferred along connection whose endpoint is `addr`.

# Sockets Interface: `bind`

- A server uses `bind` to ask the kernel to associate the server's socket address with a socket descriptor:

```
int bind(int sockfd, SA *addr, socklen_t addrlen);
```




# Sockets Interface: `bind`

- A server uses `bind` to ask the kernel to associate the server's socket address with a socket descriptor:

Best practice is to use `getaddrinfo` to supply the arguments `addr` and `addrlen`.

```
int bind(int sockfd, SA *addr, socklen_t addrlen);
```

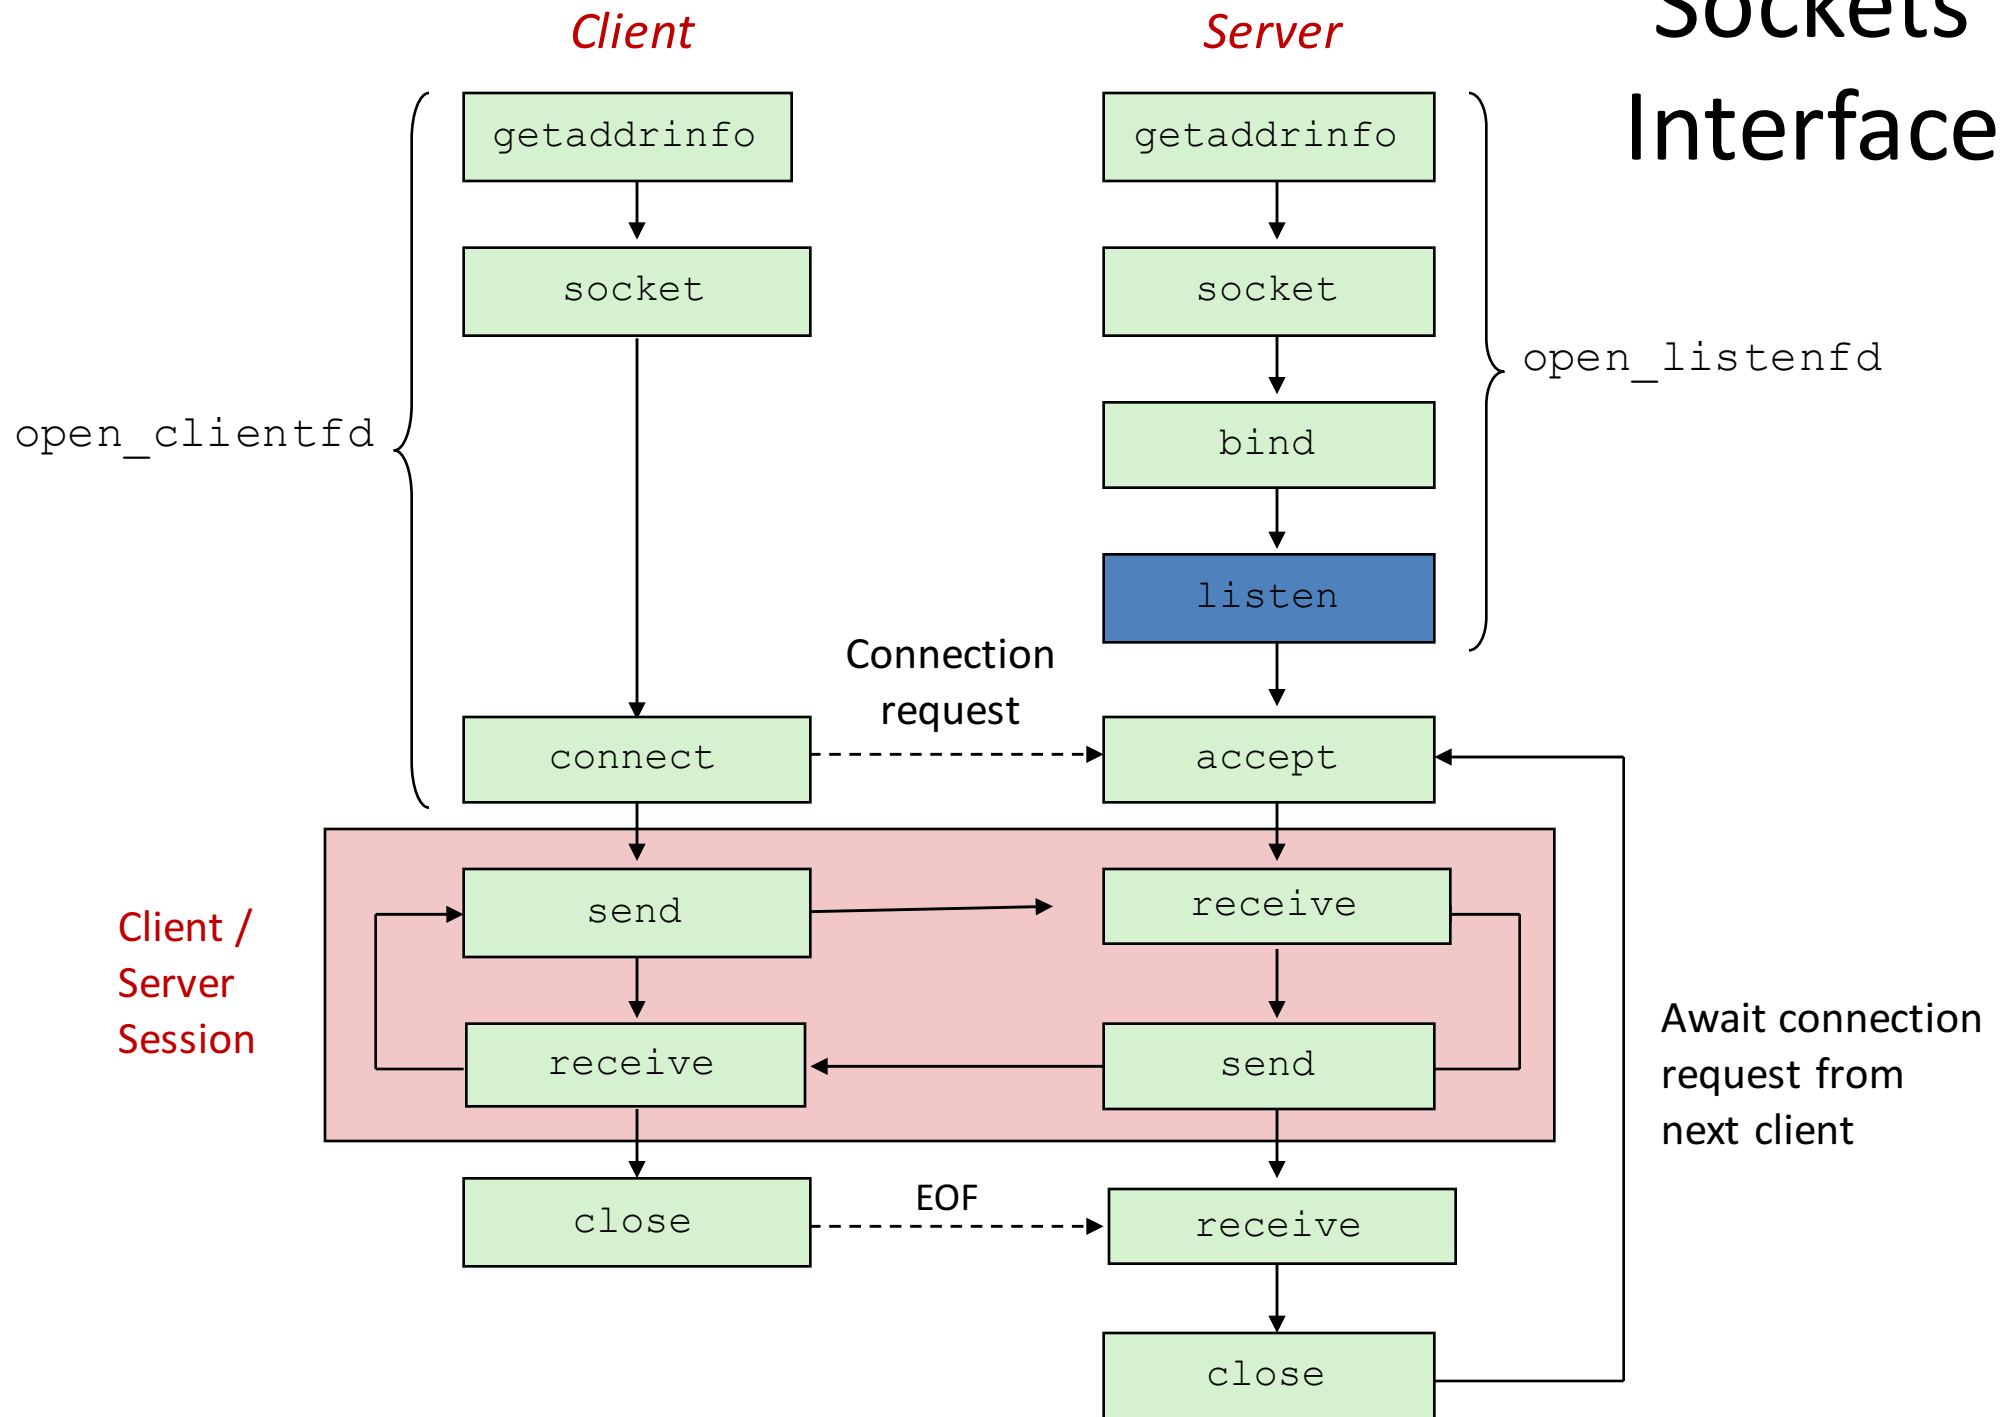


# addrinfo Struct

```
struct addrinfo {  
    int             ai_flags;        /* Hints argument flags */  
    int             ai_family;      /* First arg to socket function */  
    int             ai_socktype;    /* Second arg to socket function */  
    int             ai_protocol;    /* Third arg to socket function */  
    char            *ai_canonname;   /* Canonical host name */  
    size_t          ai_addrlen;     /* Size of ai_addr struct */  
    struct sockaddr *ai_addr;        /* Ptr to socket address structure */  
    struct addrinfo *ai_next;       /* Ptr to next item in linked list */  
};
```

- Each `addrinfo` struct returned by `getaddrinfo` contains arguments that can be passed directly to `socket` function.
- Also points to a socket address struct that can be passed directly to `connect` and `bind` functions.

# Sockets Interface





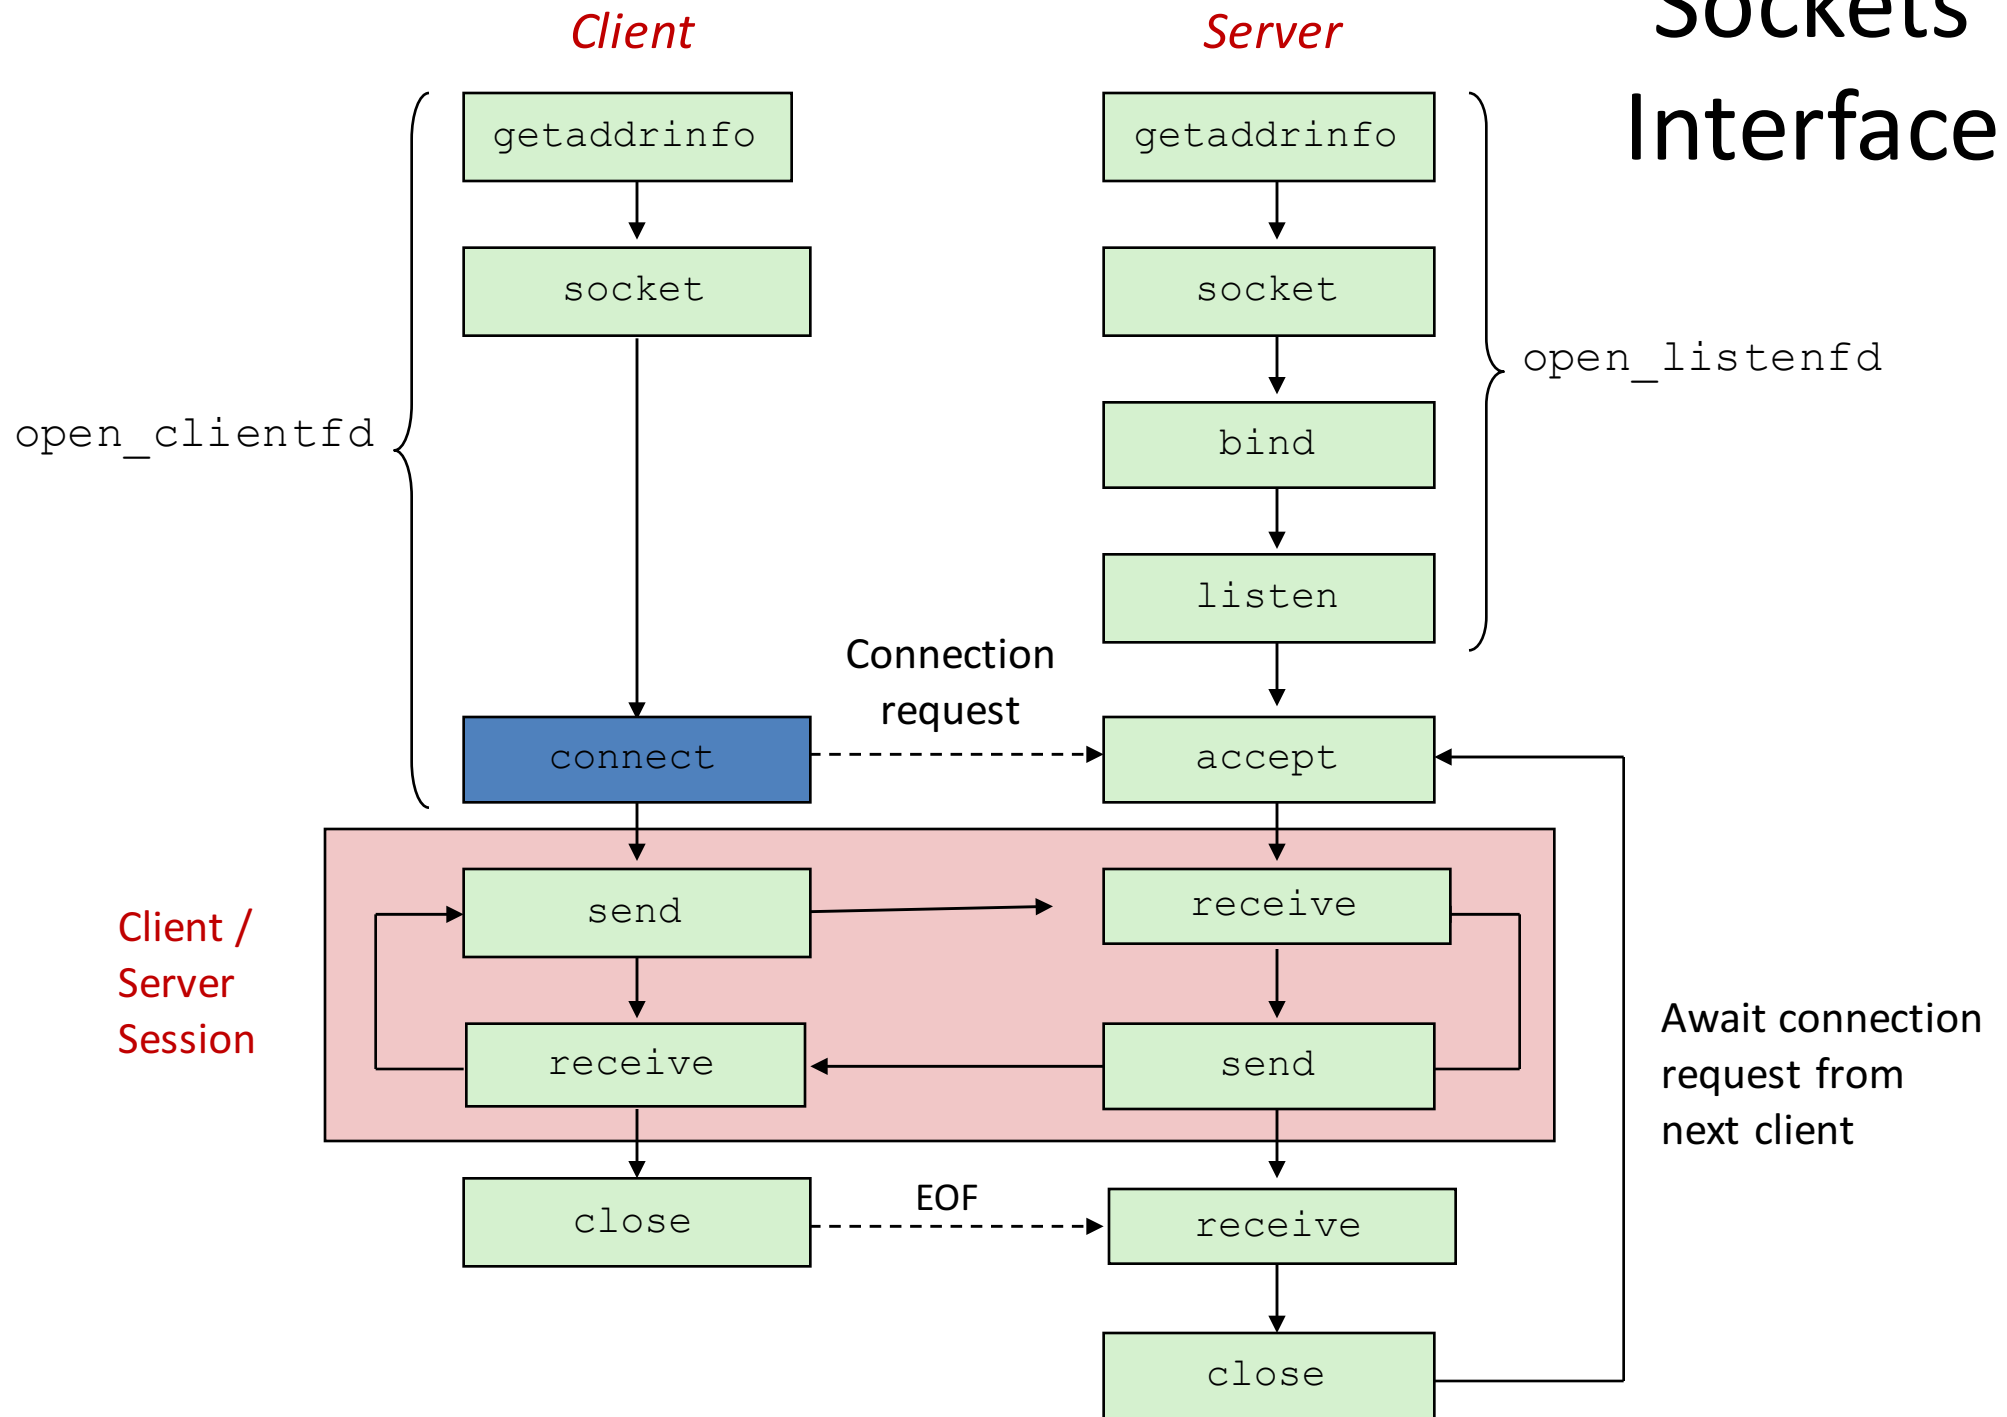
# Sockets Interface: `listen`

- By default, kernel assumes that descriptor from `socket` function is an *active socket* that will be on the client end of a connection.
- A server calls the `listen` function to tell the kernel that a descriptor will be used by a server rather than a client:

```
int listen(int sockfd, int backlog);
```

- Converts `sockfd` from an active socket to a *listening socket* that can accept connection requests from clients.
- `backlog` is a hint about the number of outstanding connection requests that the kernel should queue up before starting to refuse requests.

# Sockets Interface



# Sockets Interface: connect

- A client establishes a connection with a server by calling connect:
- Attempts to establish a connection with server at socket address `addr`
  - If successful, then `clientfd` is now ready for reading and writing.
  - Resulting connection is characterized by socket pair  
(`x:y`, `addr.sin_addr:addr.sin_port`)
    - `x` is client address
    - `y` is ephemeral port that uniquely identifies client process on client host

Best practice is to use `getaddrinfo` to supply the arguments `addr` and `addrlen`.

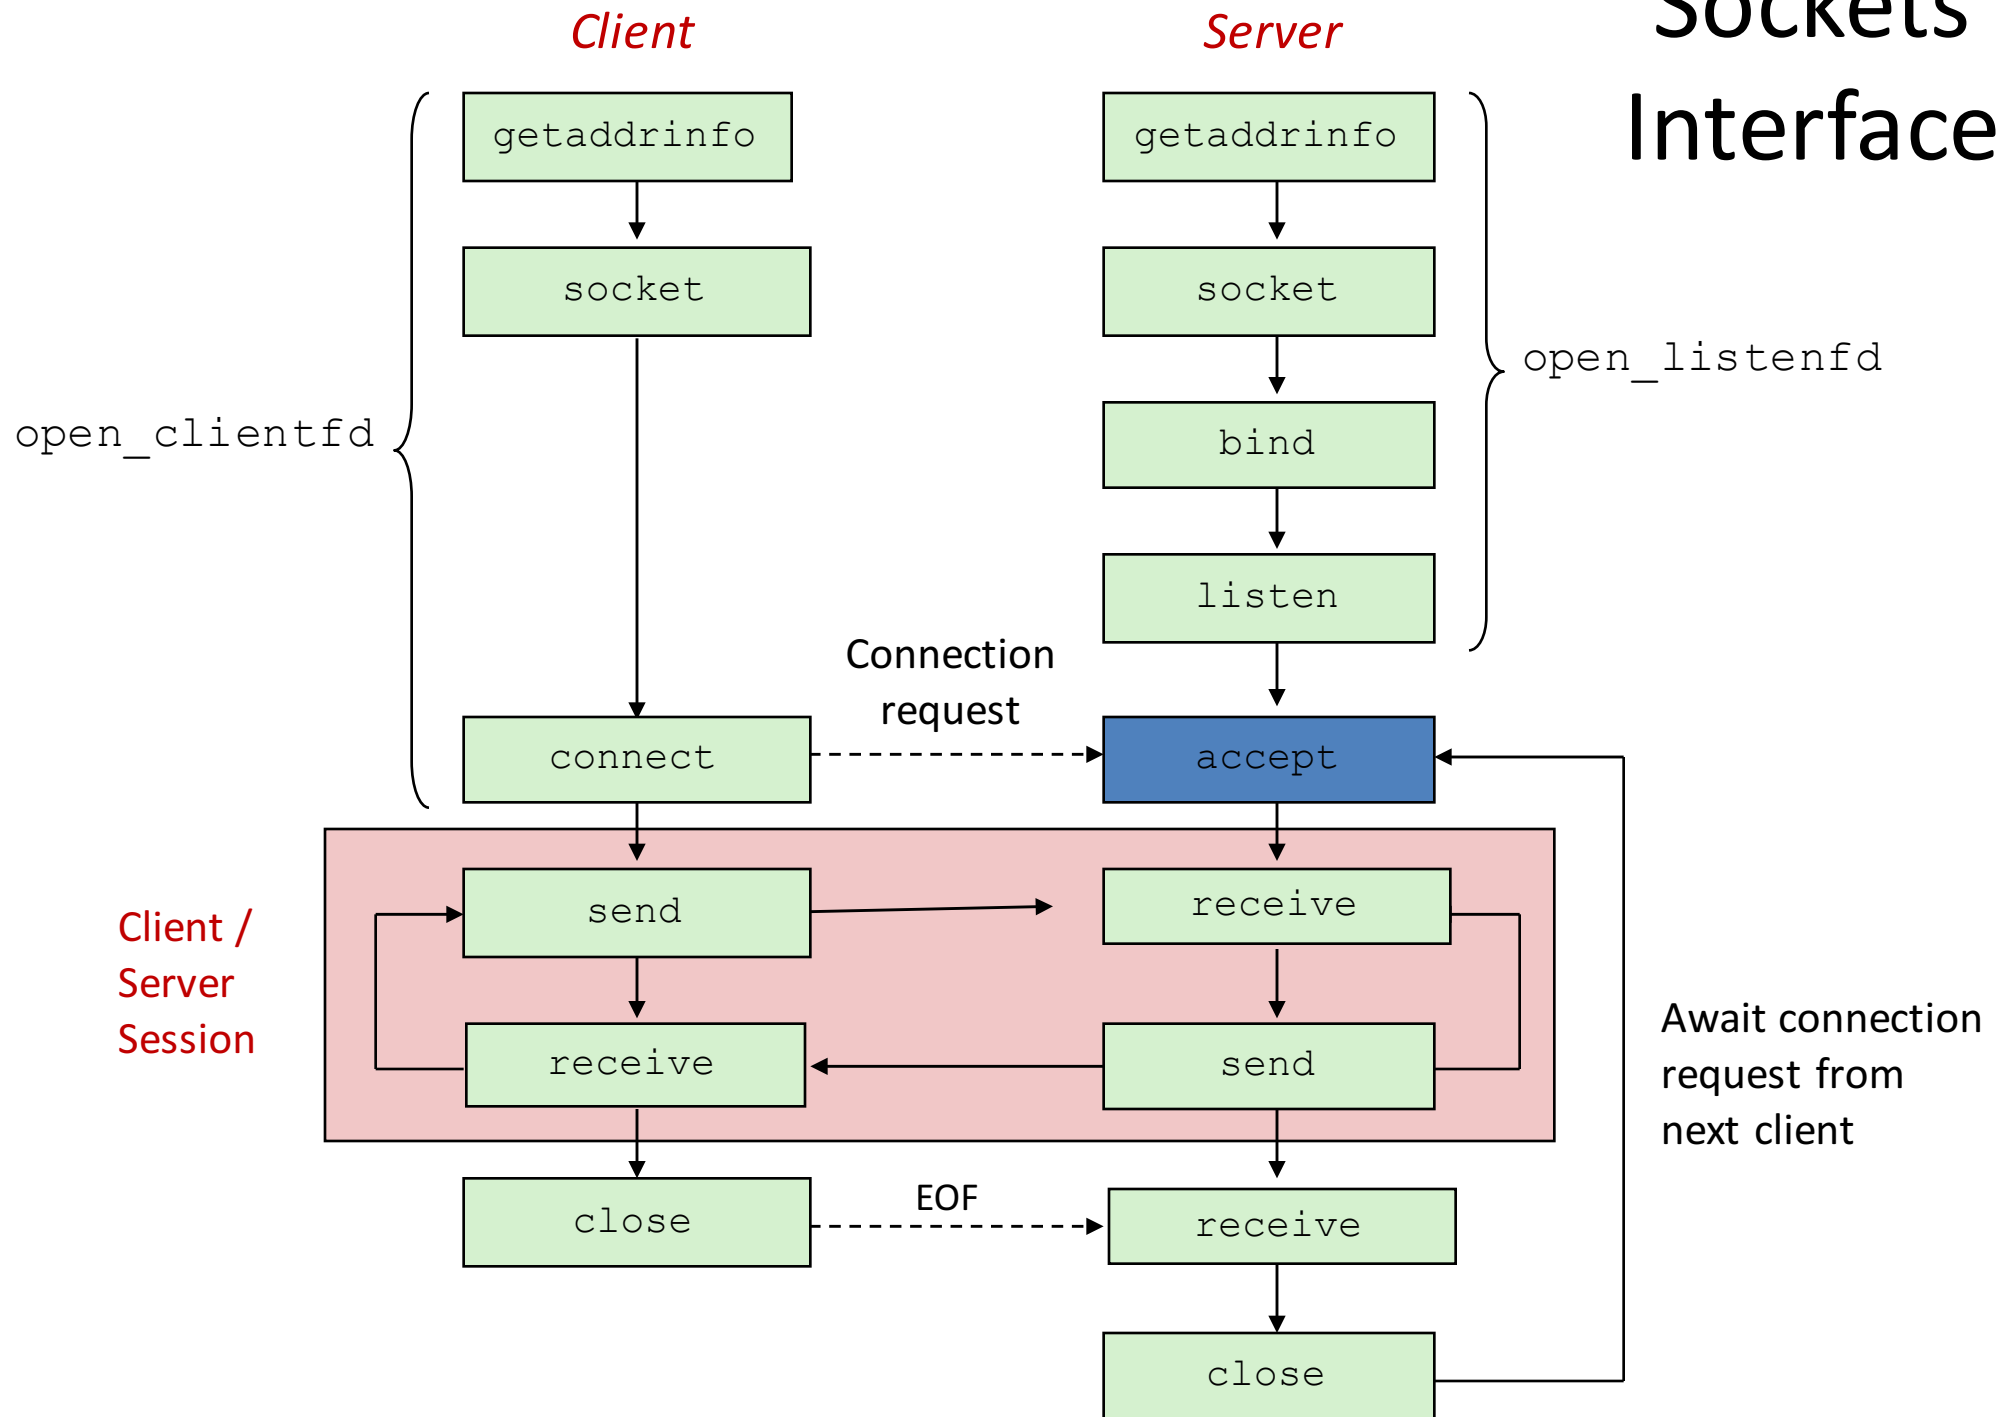
```
int connect(int clientfd, SA *addr, socklen_t addrlen);
```

# addrinfo Struct

```
struct addrinfo {  
    int             ai_flags;        /* Hints argument flags */  
    int             ai_family;       /* First arg to socket function */  
    int             ai_socktype;     /* Second arg to socket function */  
    int             ai_protocol;     /* Third arg to socket function */  
    char            *ai_canonname;    /* Canonical host name */  
    size_t          ai_addrlen;      /* Size of ai_addr struct */  
    struct sockaddr *ai_addr;         /* Ptr to socket address structure */  
    struct addrinfo *ai_next;        /* Ptr to next item in linked list */  
};
```

- Each `addrinfo` struct returned by `getaddrinfo` contains arguments that can be passed directly to `socket` function.
- Also points to a socket address struct that can be passed directly to `connect` and `bind` functions.

# Sockets Interface



# Sockets Interface: `accept`

- Servers wait for connection requests from clients by calling `accept`:

```
int accept(int listenfd, SA *addr, int *addrlen);
```

- Waits for connection request to arrive on the connection bound to `listenfd`, then fills in client's socket address in `addr` and size of the socket address in `addrlen`.
- Returns a *connected descriptor* that can be used to communicate with the client via Unix I/O routines.

# accept Illustrated




*1. Server blocks in `accept`,  
waiting for connection request  
on listening descriptor  
`listenfd`*

# accept Illustrated



*1. Server blocks in `accept`, waiting for connection request on listening descriptor `listenfd`*




*2. Client makes connection request by calling and blocking in `connect`*



# accept Illustrated



1. Server blocks in *accept*, waiting for connection request on listening descriptor *listenfd*



2. Client makes connection request by calling and blocking in *connect*

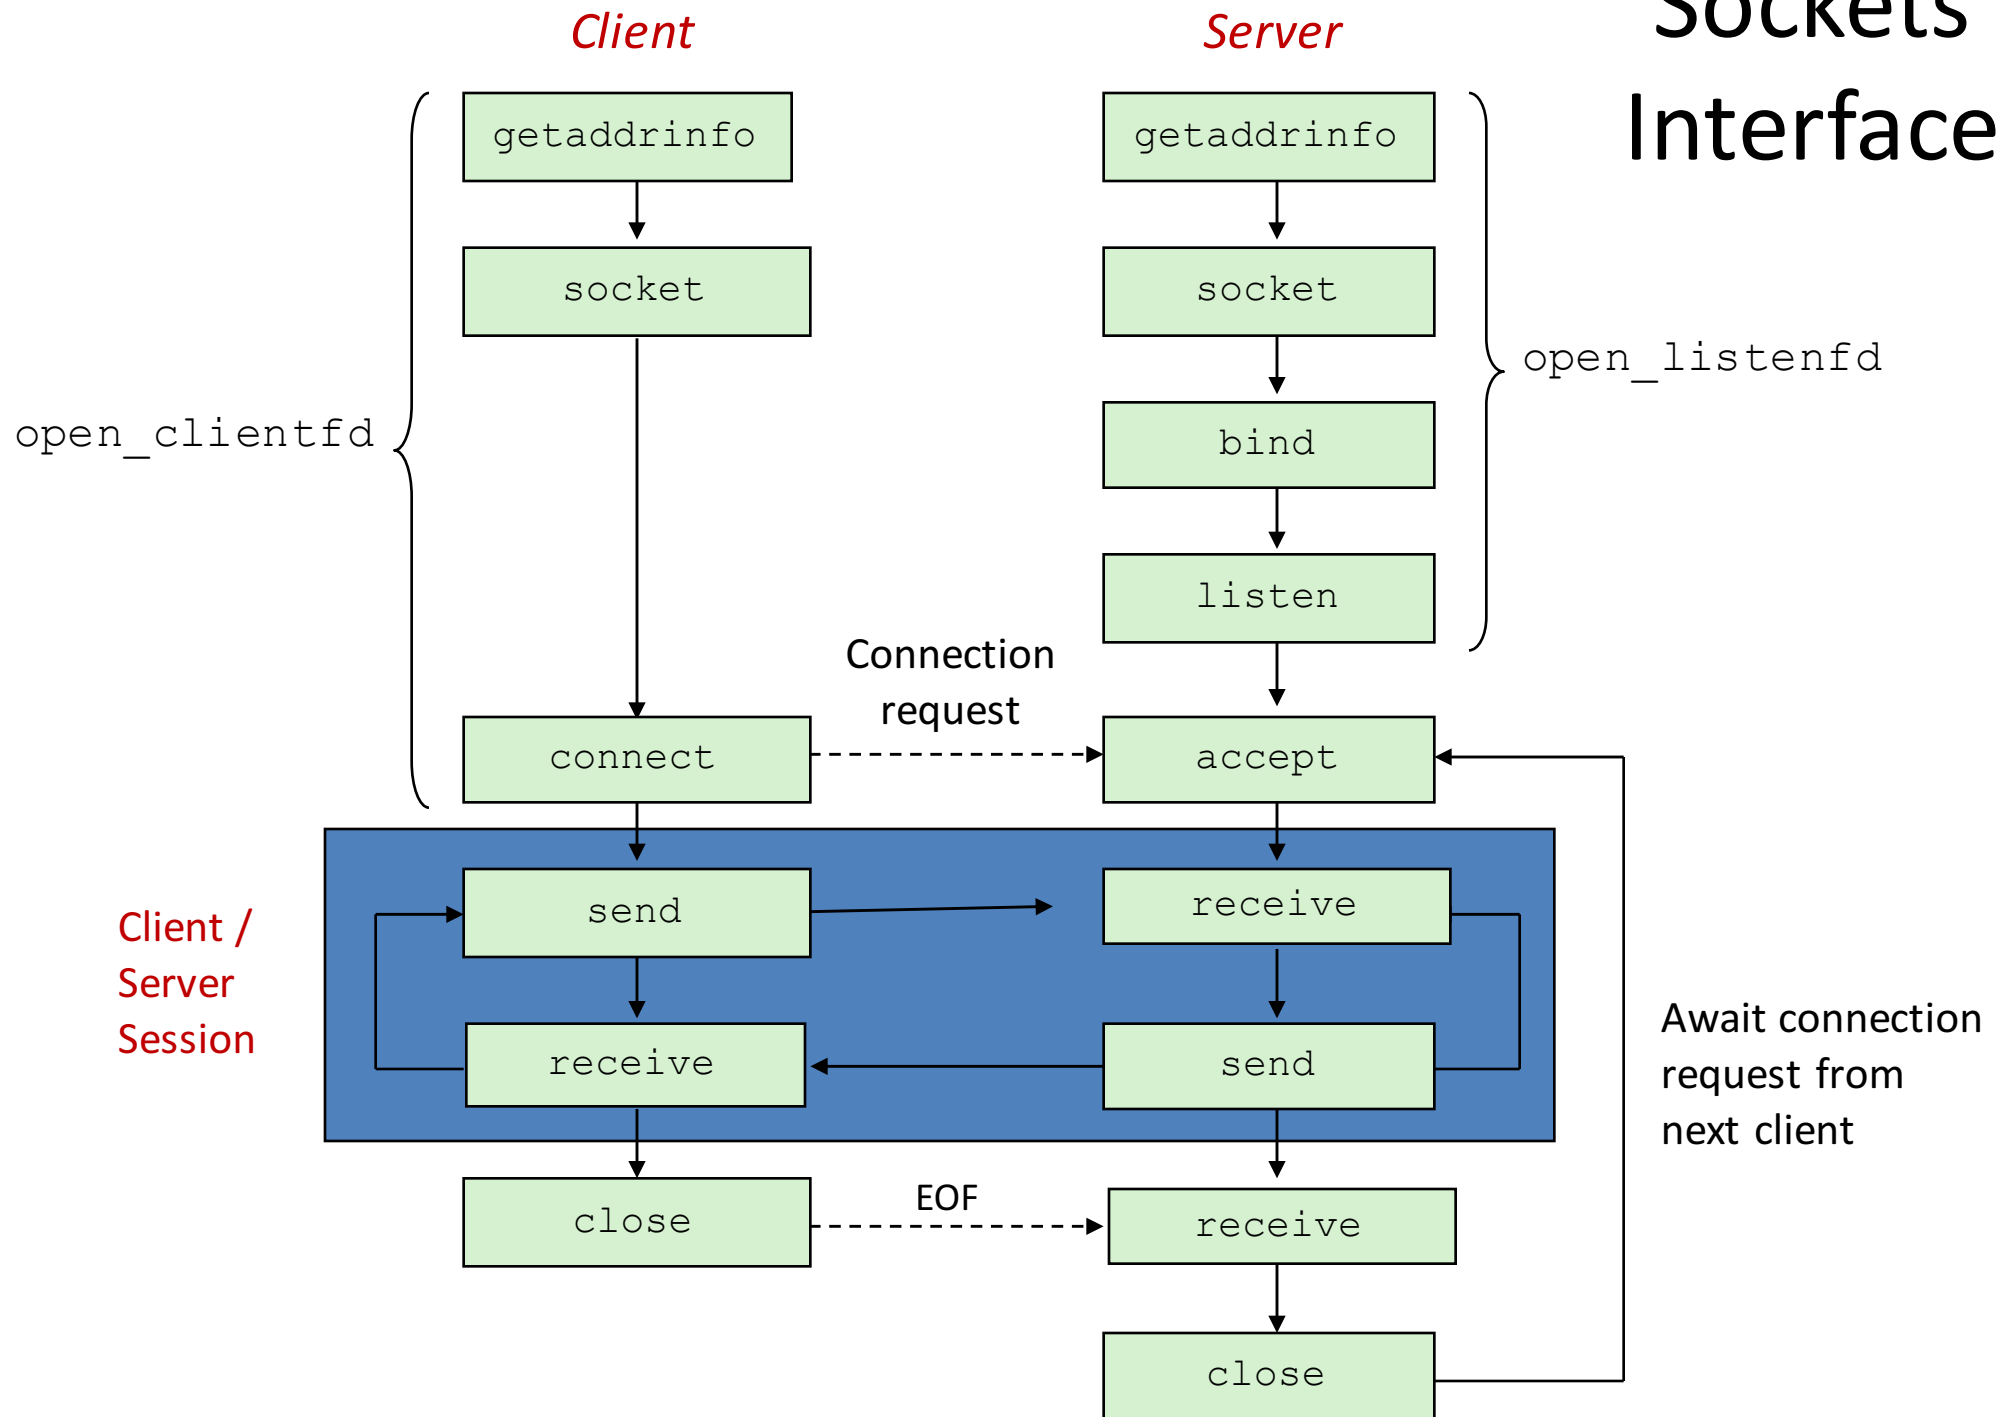


3. Server returns *connfd* from *accept*. Client returns from *connect*. Connection is now established between *clientfd* and *connfd*

# Connected vs. Listening Descriptors

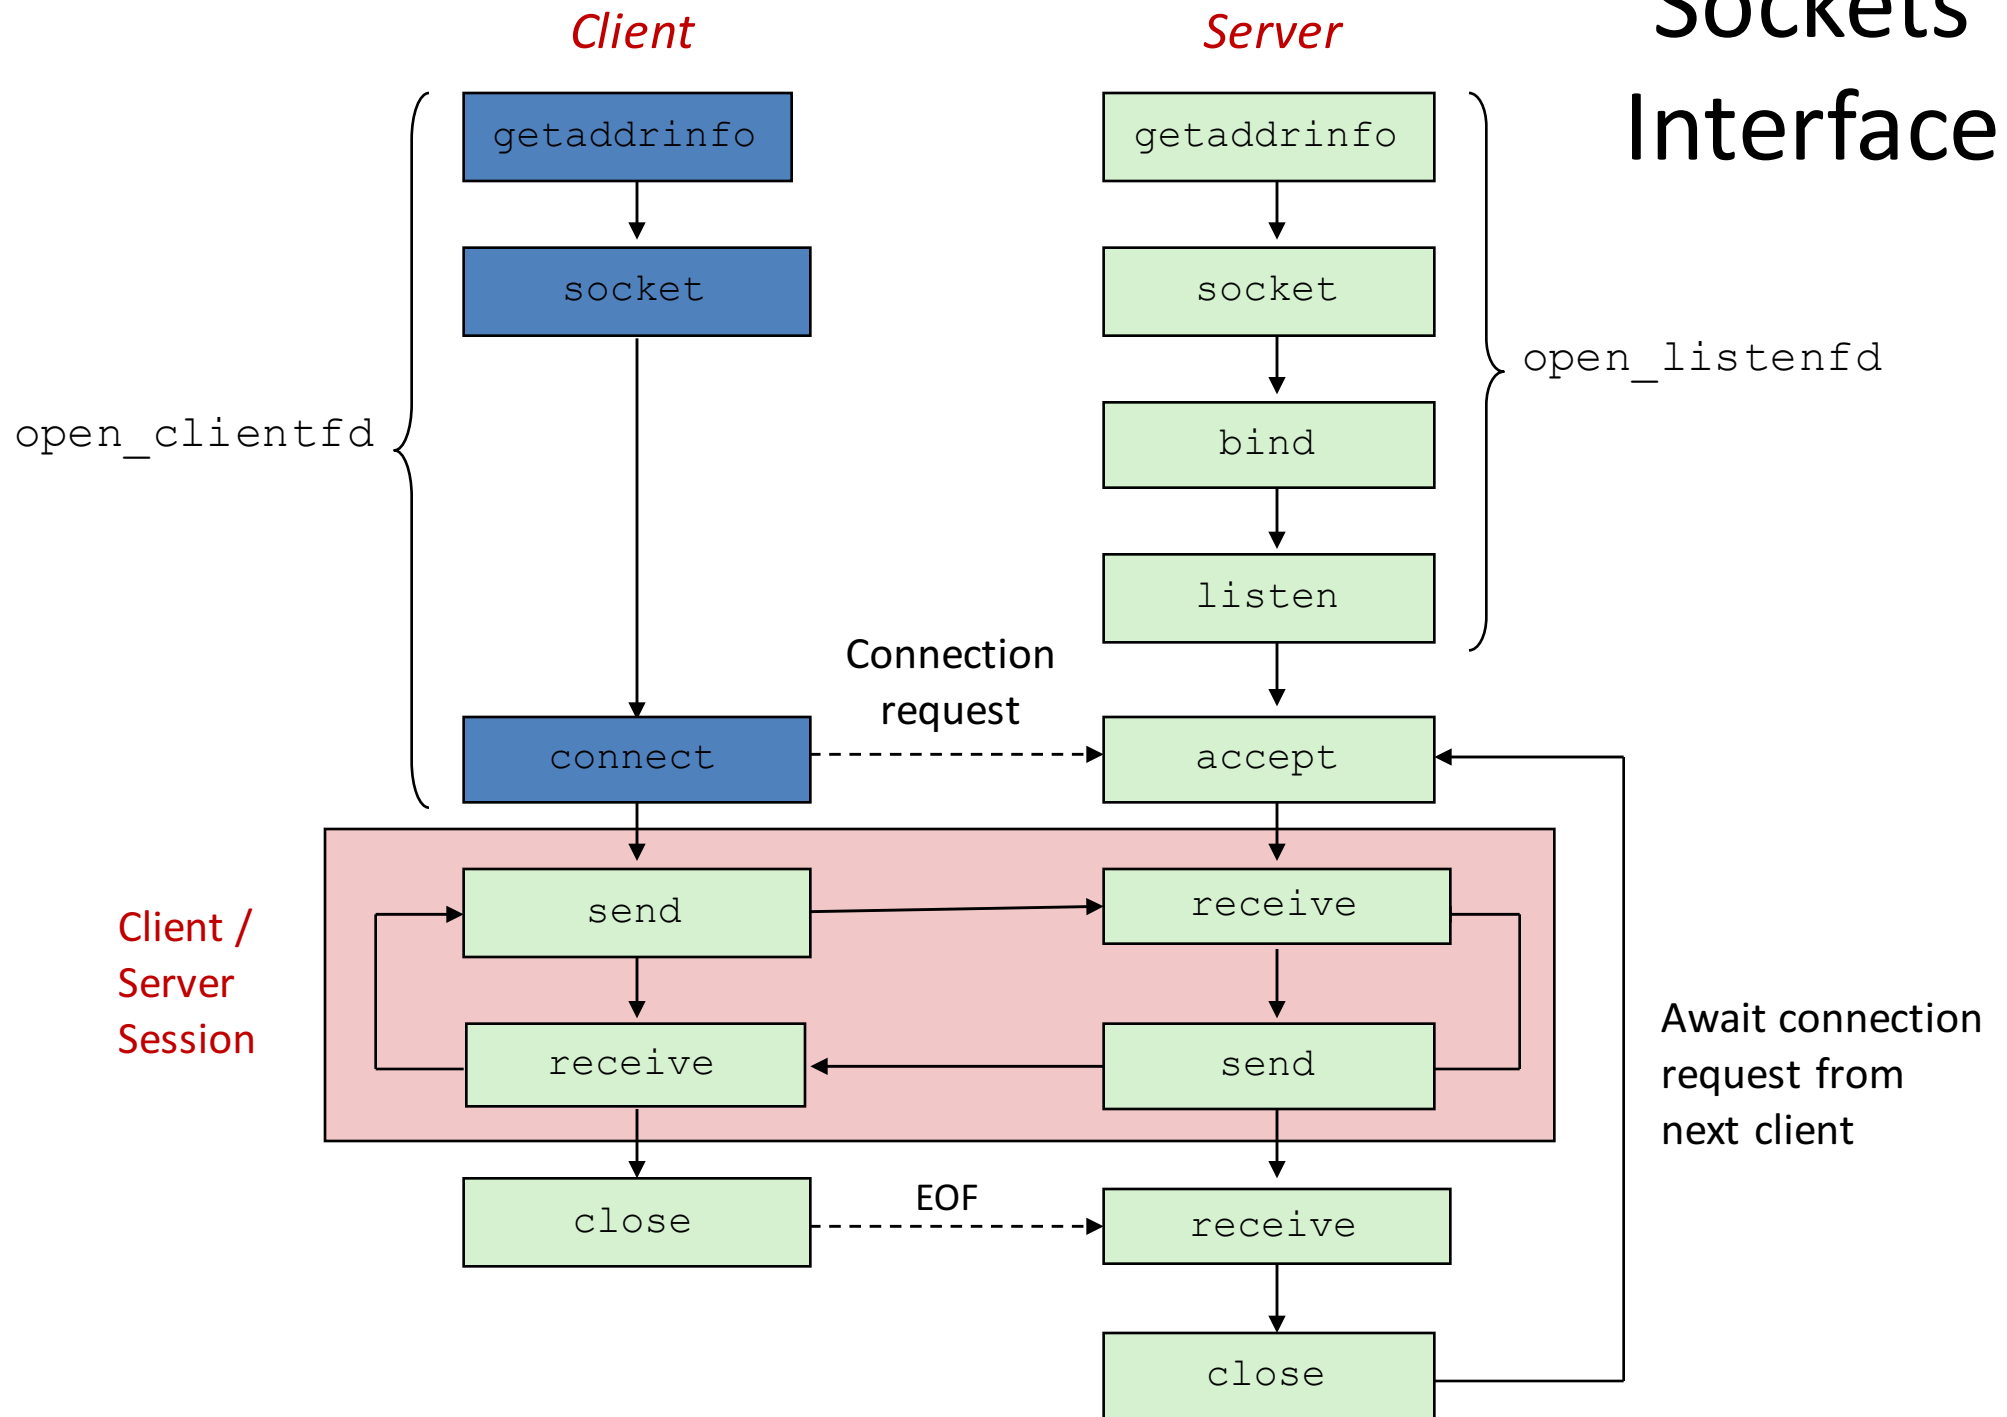
- Listening descriptor
  - End point for client connection requests
  - Created once and exists for lifetime of the server
- Connected descriptor
  - End point of the connection between client and server
  - A new descriptor is created each time the server accepts a connection request from a client
  - Exists only as long as it takes to service client
- Why the distinction?
  - Allows for concurrent servers that can communicate over many client connections simultaneously
    - E.g., Each time we receive a new request, we fork a child to handle the request

# Sockets Interface



**ECHO SERVER**

# Sockets Interface



# Sockets Helper: `open_clientfd`

- Establish a connection with a server

```
int open_clientfd(char *hostname, char *port) {  
    ...  
}
```

# Sockets Helper: open\_clientfd

- Establish a connection with a server

```
int open_clientfd(char *hostname, char *port) {
    int clientfd;
    struct addrinfo hints, *listp, *p;

    /* Get a list of potential server addresses */
    memset(&hints, 0, sizeof(struct addrinfo));
    hints.ai_socktype = SOCK_STREAM; /* Open a connection */
    hints.ai_flags = AI_NUMERICSERV; /* ...using numeric port arg. */
    hints.ai_flags |= AI_ADDRCONFIG; /* Recommended for connections */
    Getaddrinfo(hostname, port, &hints, &listp);
```

csapp.c

# Sockets Helper: open\_clientfd (cont)

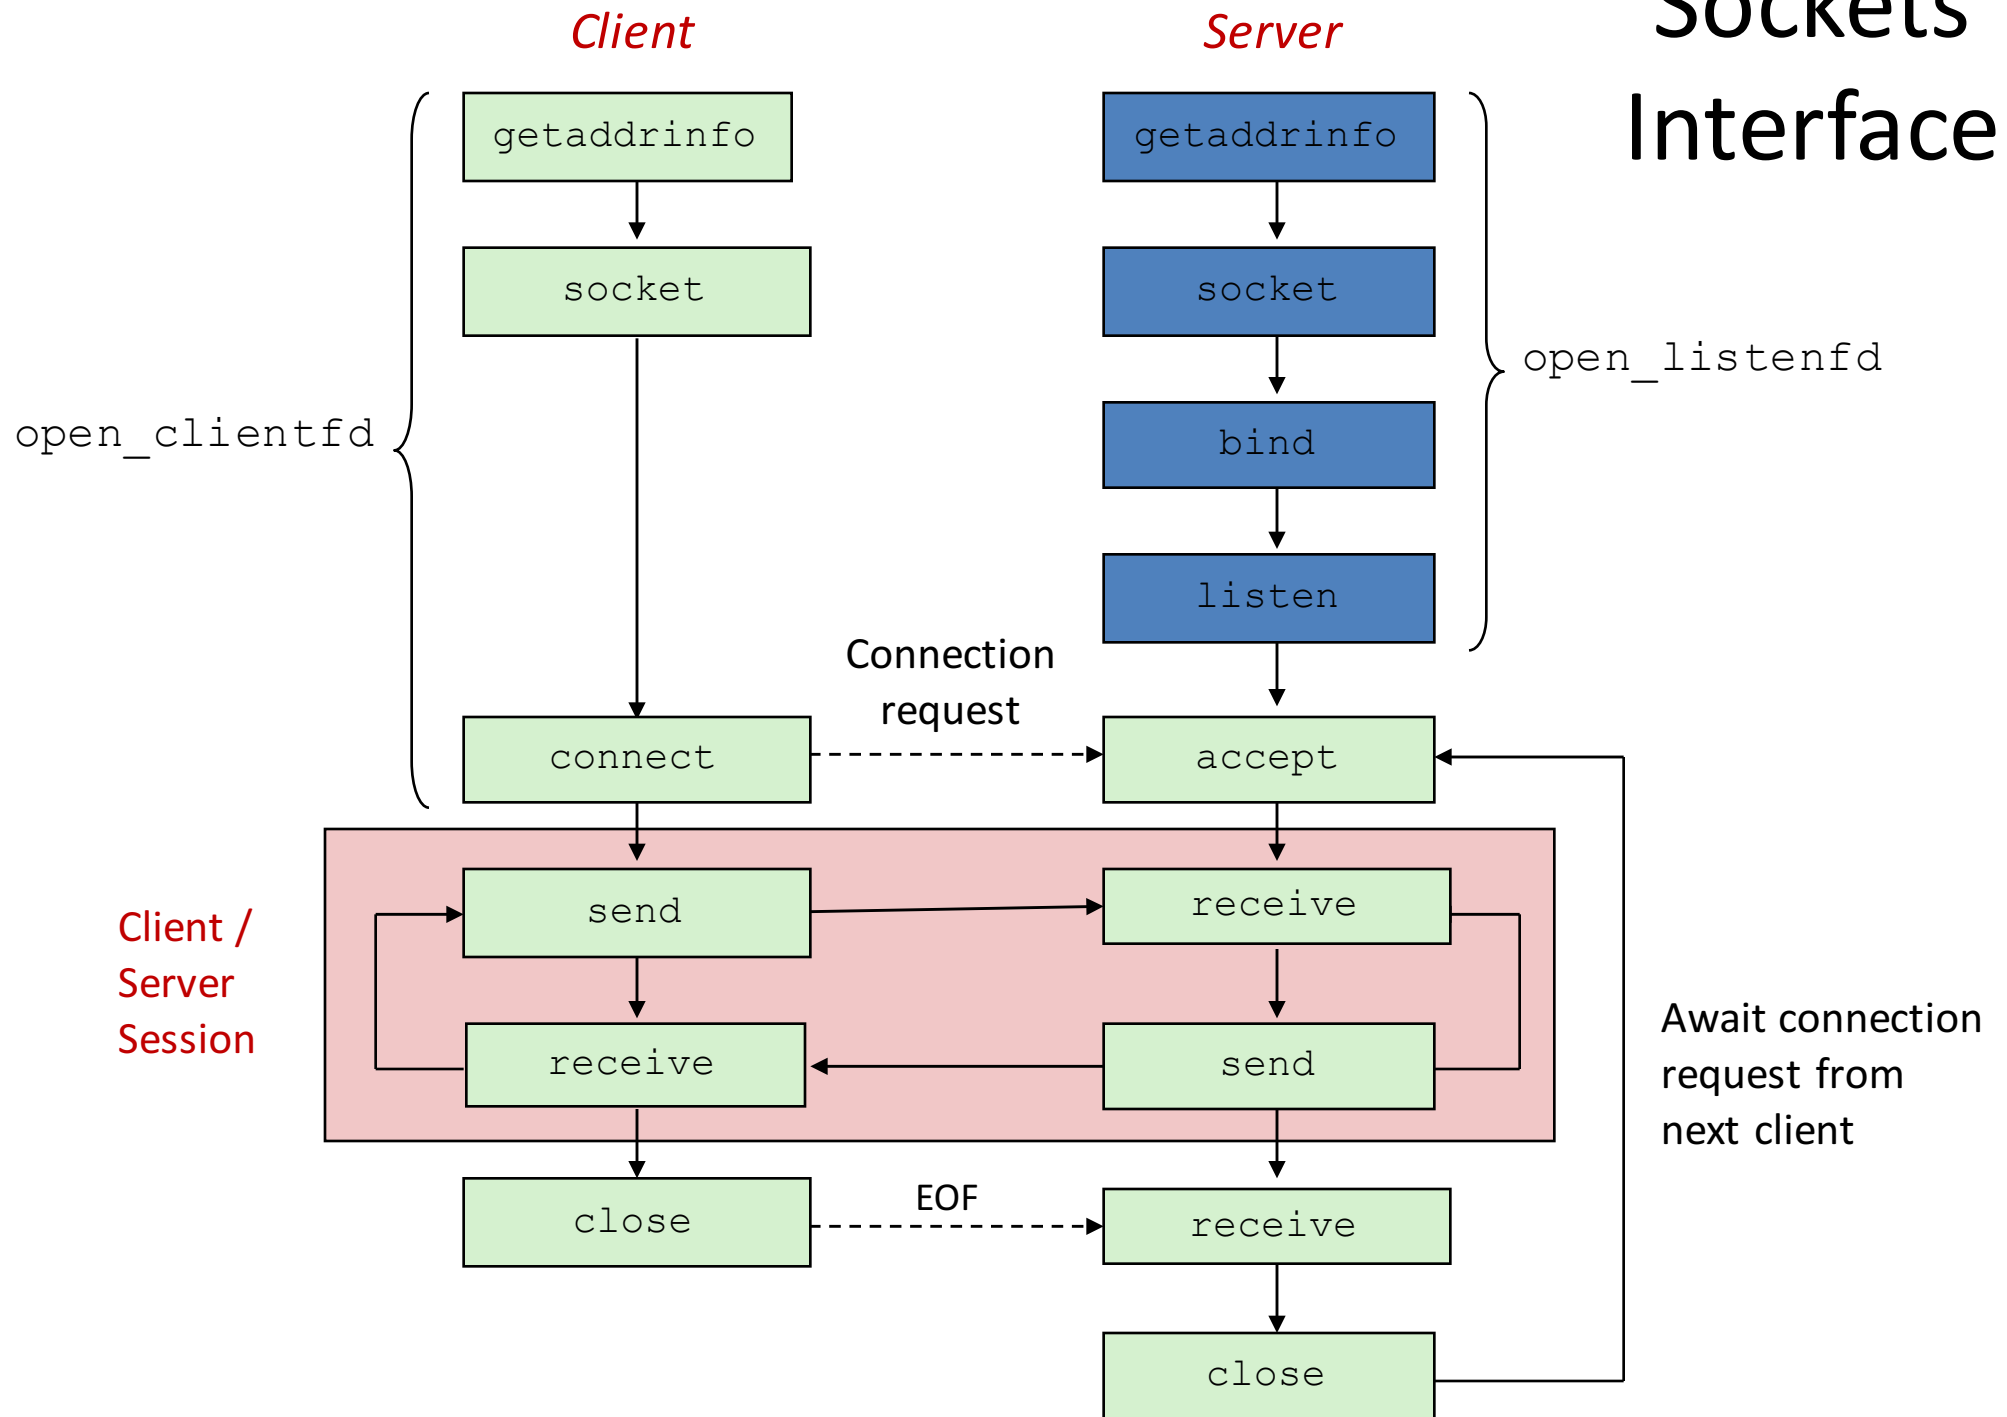
```
/* Walk the list for one that we can successfully connect to */
for (p = listp; p; p = p->ai_next) {
    /* Create a socket descriptor */
    if ((clientfd = socket(p->ai_family, p->ai_socktype,
                          p->ai_protocol)) < 0)
        continue; /* Socket failed, try the next */

    /* Connect to the server */
    if (connect(clientfd, p->ai_addr, p->ai_addrlen) != -1)
        break; /* Success */
    Close(clientfd); /* Connect failed, try another */
}

/* Clean up */
Freeaddrinfo(listp);
if (!p) /* All connects failed */
    return -1;
else /* The last connect succeeded */
    return clientfd;
}
```



# Sockets Interface



# Sockets Helper: `open_listenfd`

- Create a listening descriptor that can be used to accept connection requests from clients.

```
int open_listenfd(char *port)
{
    struct addrinfo hints, *listp, *p;
    int listenfd, optval=1;

    /* Get a list of potential server addresses */
    memset(&hints, 0, sizeof(struct addrinfo));
    hints.ai_socktype = SOCK_STREAM; /* Accept connect. */
    hints.ai_flags = AI_PASSIVE | AI_ADDRCONFIG; /* ...on any IP addr */
    hints.ai_flags |= AI_NUMERICSERV; /* ...using port no. */
    Getaddrinfo(NULL, port, &hints, &listp);
```

# Sockets Helper: open\_listenfd (cont)

```
/* Walk the list for one that we can bind to */
for (p = listp; p; p = p->ai_next) {
    /* Create a socket descriptor */
    if ((listenfd = socket(p->ai_family, p->ai_socktype,
                          p->ai_protocol)) < 0)
        continue; /* Socket failed, try the next */

    Setsockopt(listenfd, SOL_SOCKET, SO_REUSEADDR,
               (const void *)&optval , sizeof(int));

    /* Bind the descriptor to the address */
    if (bind(listenfd, p->ai_addr, p->ai_addrlen) == 0)
        break; /* Success */
    Close(listenfd); /* Bind failed, try the next */
}
```

csapp.c

# Sockets Helper: open\_listenfd (cont)

```
/* Clean up */
Freeaddrinfo(listp);
if (!p) /* No address worked */
    return -1;

/* Make it a listening socket ready to accept conn. requests */
if (listen(listenfd, LISTENQ) < 0) {
    Close(listenfd);
    return -1;
}
return listenfd;
}
```

csapp.c

- **Key point:** open\_clientfd and open\_listenfd are both independent of any particular version of IP.

# Echo Client: Main Routine

```
#include "csapp.h"

int main(int argc, char **argv)
{
    ...
    clientfd = open_clientfd(host, port);

    printf("type: ");
    fflush(stdout);

    while (fgets(buf, MAXLINE, stdin) != NULL) {
        send(clientfd, buf, strlen(buf), 0);
        recv(clientfd, buf, MAXLINE, 0);

        // Display and read again:
        printf("echo: ");
        fputs(buf, stdout);
        printf("type: ");
        fflush(stdout);
    }
    ...
}
```

echoclient.c

# Iterative Echo Server: Main Routine

```
#include "csapp.h"
void echo(int connfd);

int main(int argc, char **argv)
{
    ...
    port = atoi(argv[1]);

    listenfd = open_listenfd(port);
    while (1) {
        clientlen = sizeof(clientaddr);
        connfd = accept(listenfd, (SA *)&clientaddr, &clientlen);

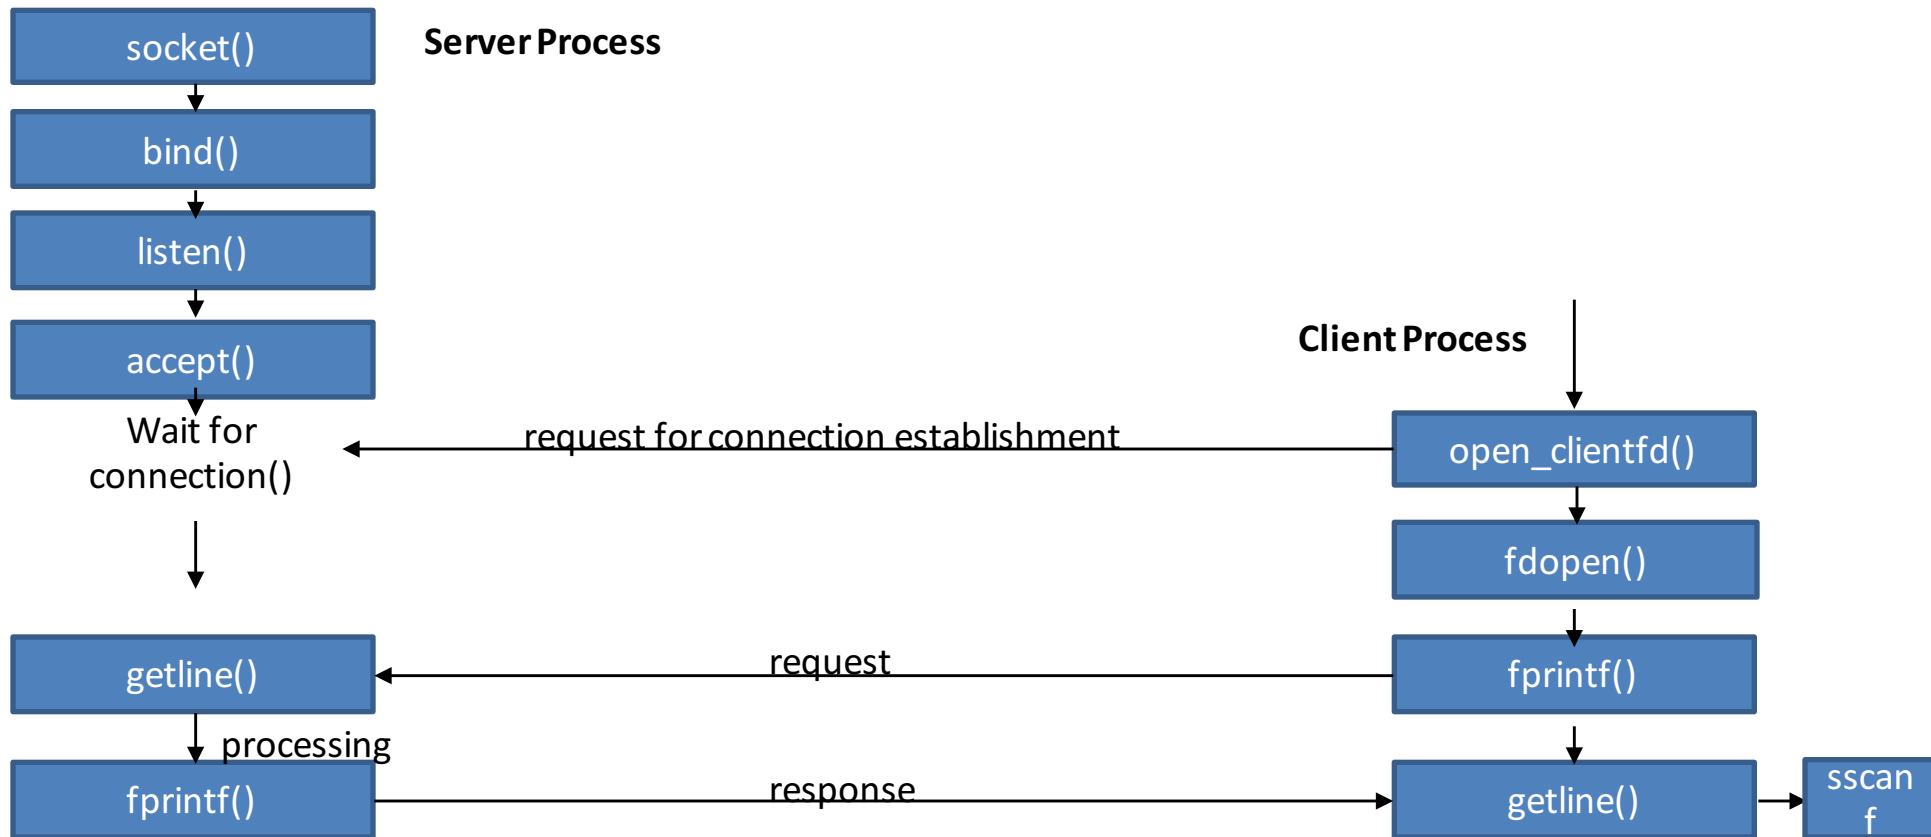
        /* determine the domain name and IP address of the client */
        hp = gethostbyaddr((const char *)&clientaddr.sin_addr.s_addr,
                           sizeof(clientaddr.sin_addr.s_addr), AF_INET);
        haddrp = inet_ntoa(clientaddr.sin_addr);
        client_port = ntohs(clientaddr.sin_port);
        printf("server connected to %s (%s), port %u\n",
               hp->h_name, haddrp, client_port);
        echo(connfd);
        printf("Connection closed\n");
        close(connfd);
    }
}
```

echoserver.c

# Echo Server: echo function

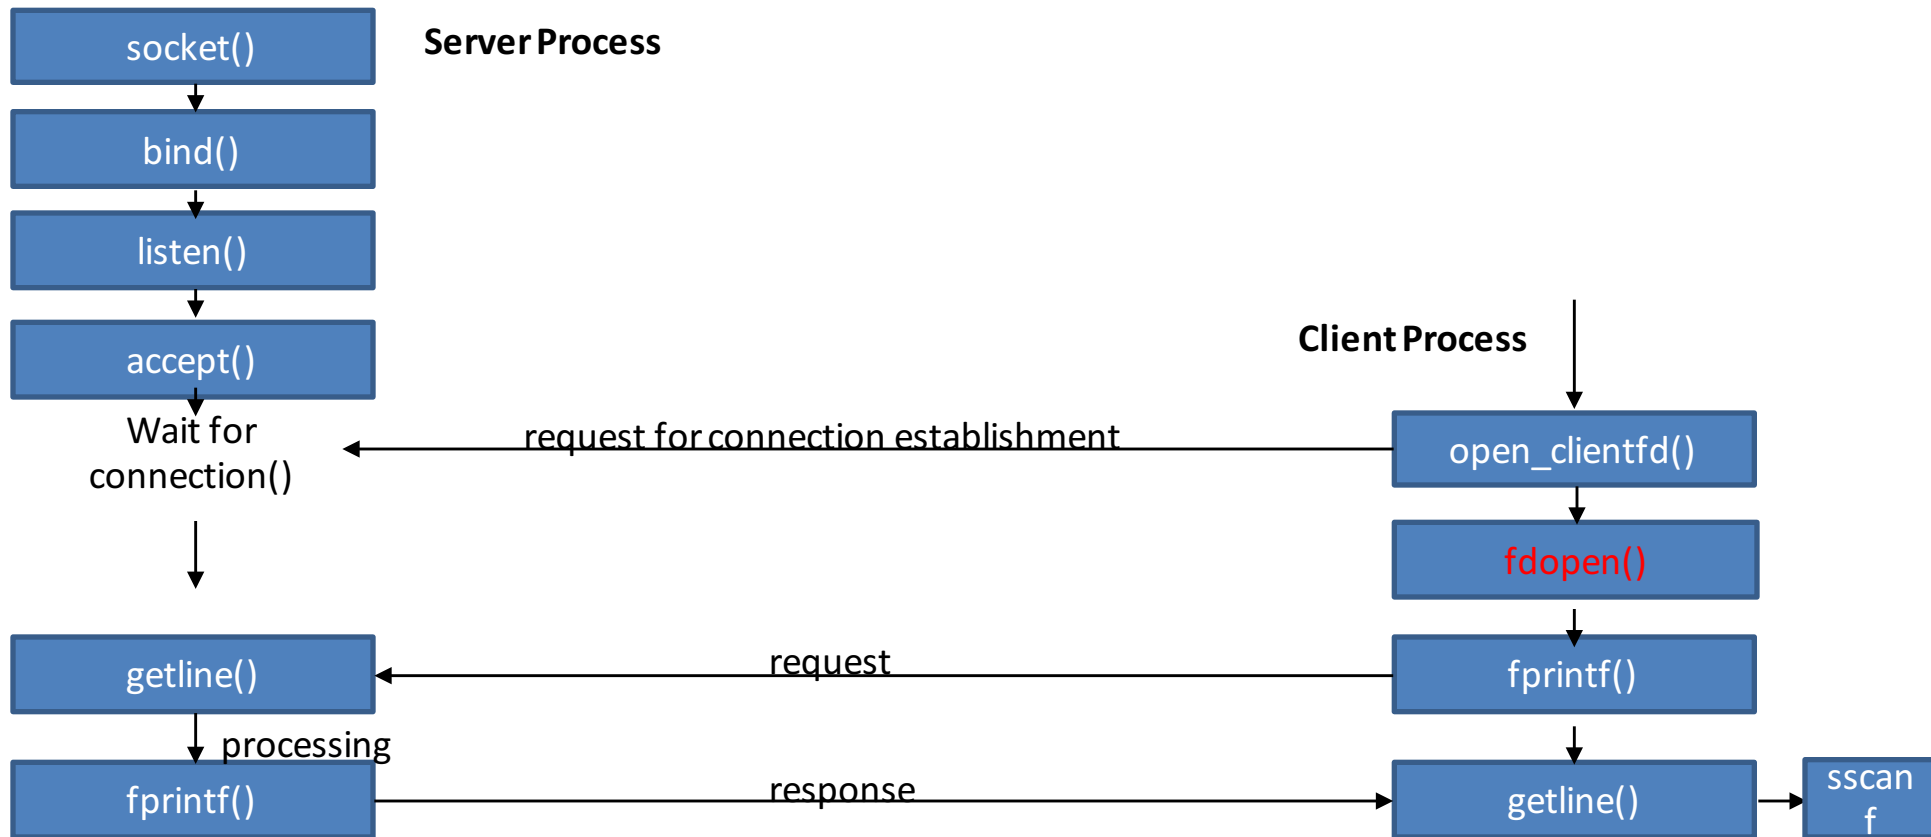
```
void echo (int connfd) {  
    // Local variable declarations:  
    size_t n;  
    char buf[MAXLINE];  
  
    // Keep reading lines until client closes connection:  
    while((n = recv(connfd, buf, MAXLINE, 0)) != 0) {  
        printf("server received %d bytes\n", (int) n);  
        upper_case(buf);  
        send(connfd, buf, n, 0);  
    }  
}
```

# HW9

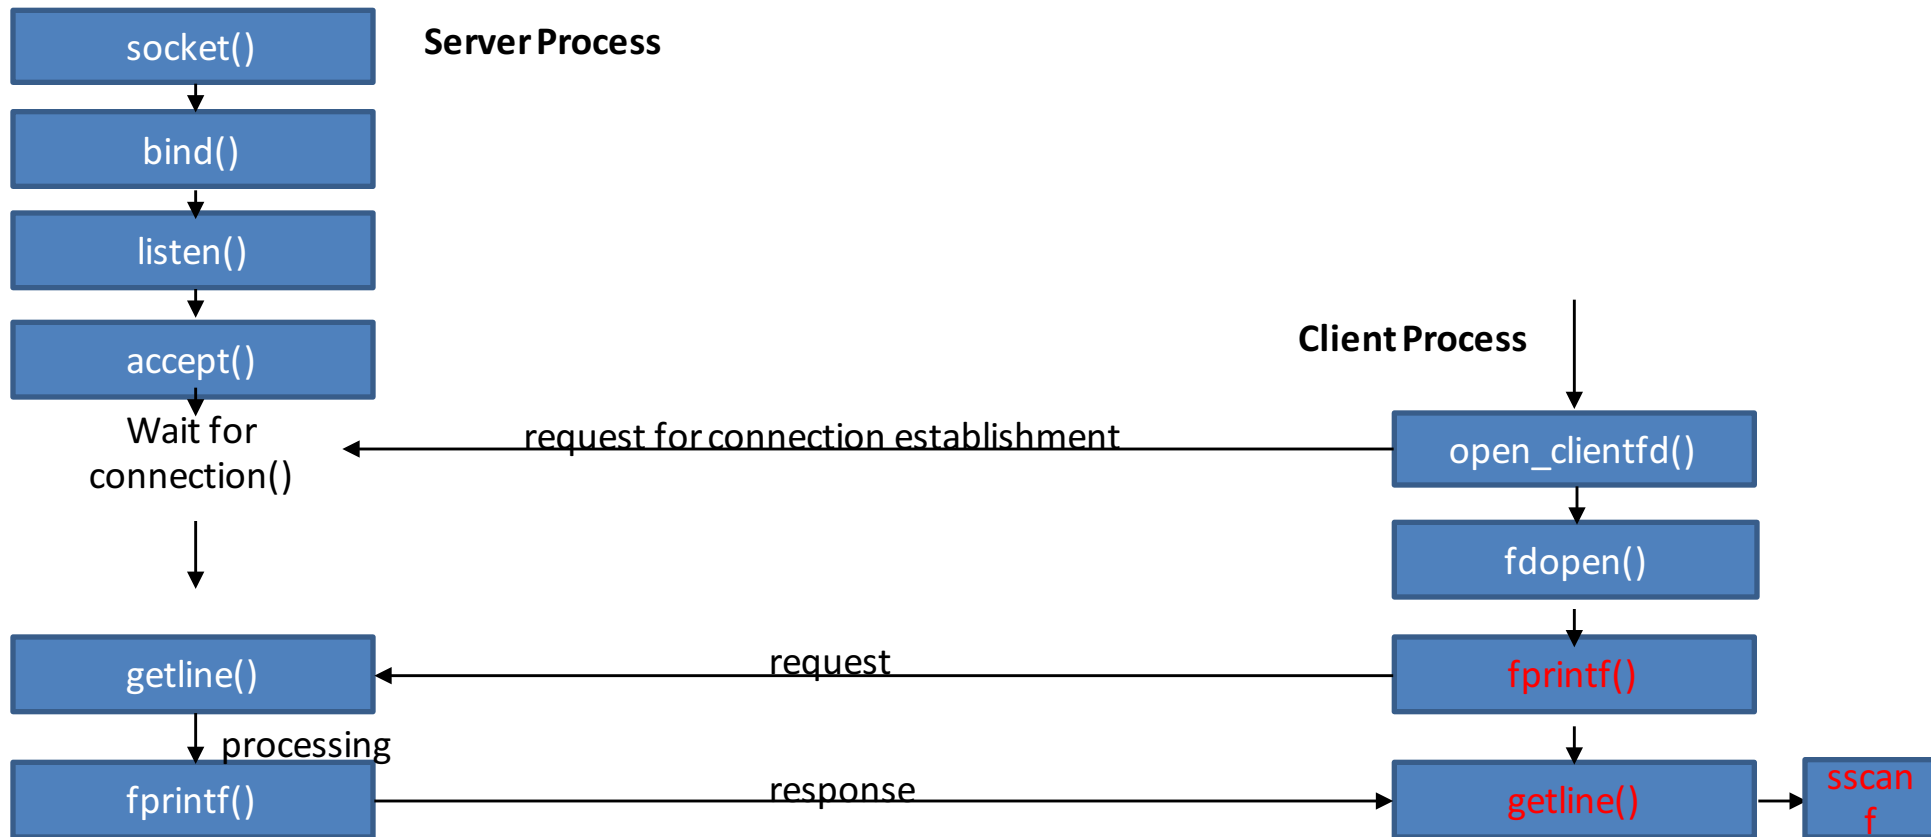




# TCP connection



# TCP connection



# i-clicker question

Which of the following option is wrong?

- A. If `open_clientfd()` returns -1, it may indicate that your internet is not working.
- B. client can send request to the server using `getline()` and read response from the server using `fprintf()`
- C. We use `getline` to get the whole line, and `sscanf` to break the input into its pieces
- D. The *`fdopen()`* function shall associate a stream with the socket descriptor.

# i-clicker question

Which of the following option is wrong?

- A. If `open_clientfd()` returns -1, it may indicate that your internet is not working.
- B. client can send request to the server using `getline()` and read response from the server using `fprintf()`**
- C. We use `getline` to get the whole line, and `sscanf` to break the input into its pieces
- D. The *`fdopen()`* function shall associate a stream with the socket descriptor.